

Lecture 7:

Geometry Processing

Computer Graphics
CMU 15-462/15-662, Fall 2015

Last time: Curves, Surfaces & Meshes

■ Mathematical description of geometry

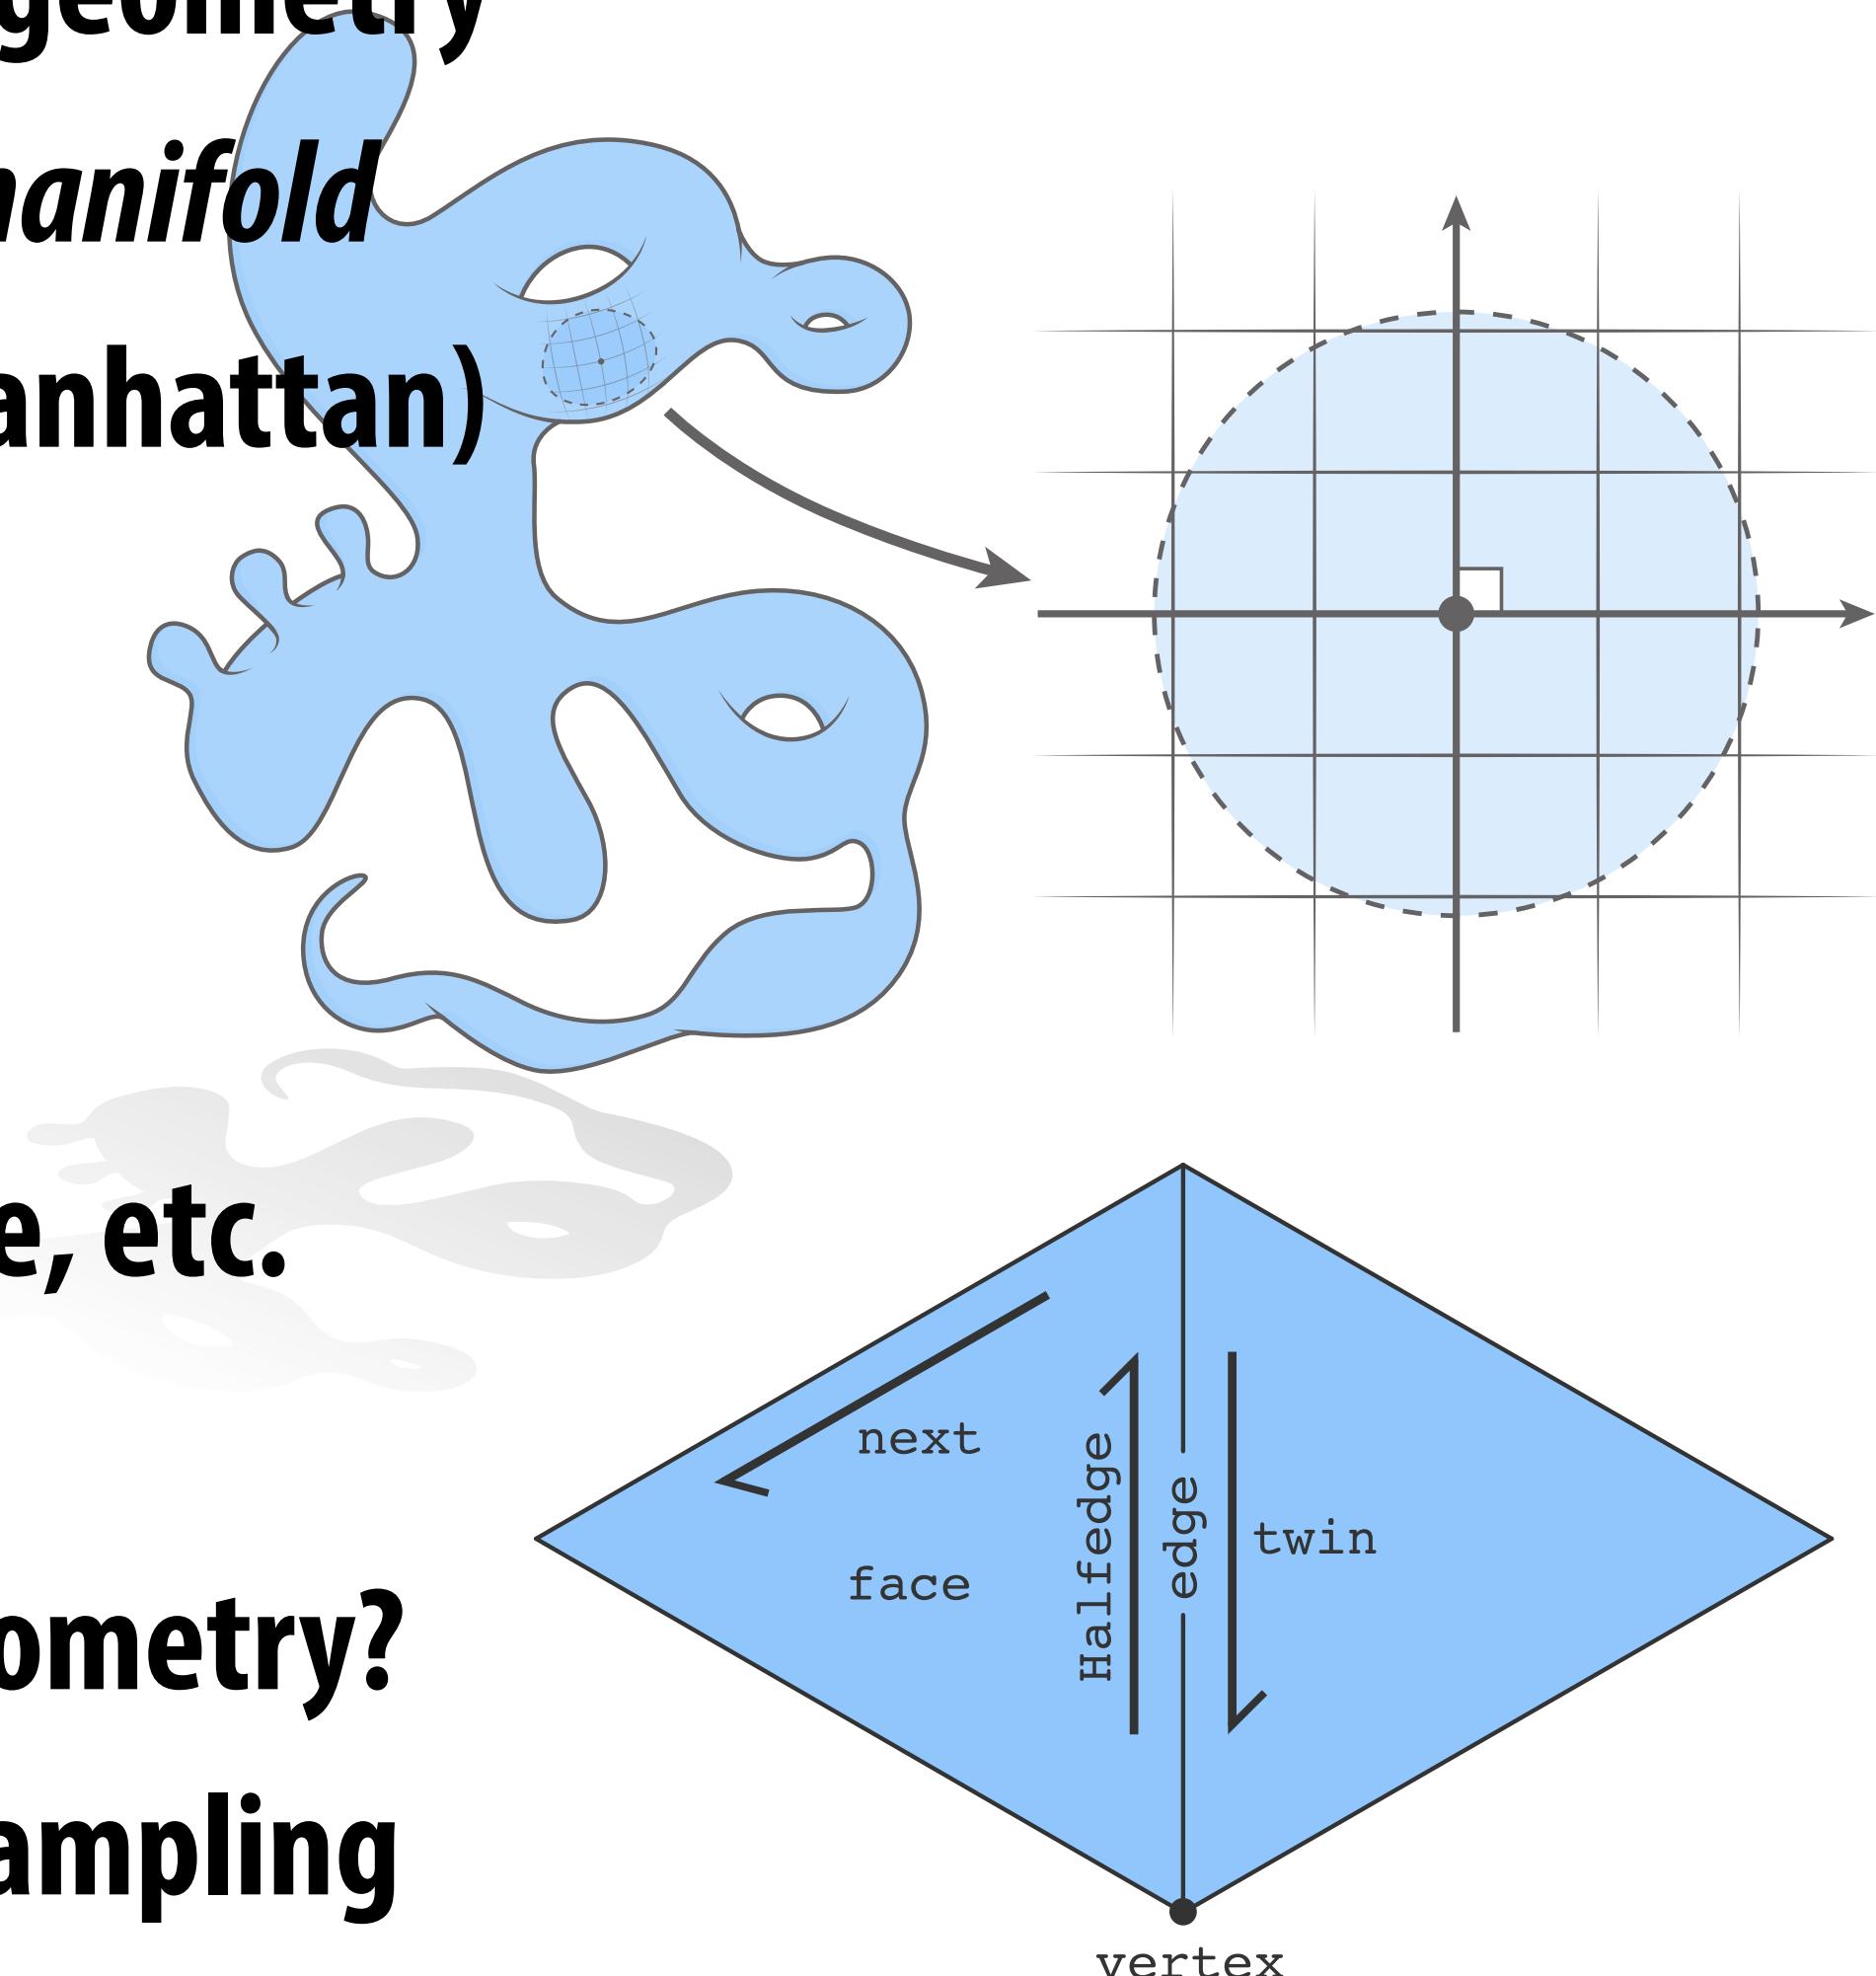
- simplifying assumption: *manifold*
- gives local coordinates (Manhattan)

■ Data structures for surfaces

- polygon soup
- halfedge mesh
- storage cost vs. access time, etc.

■ Today:

- how do we manipulate geometry?
- geometry processing / resampling

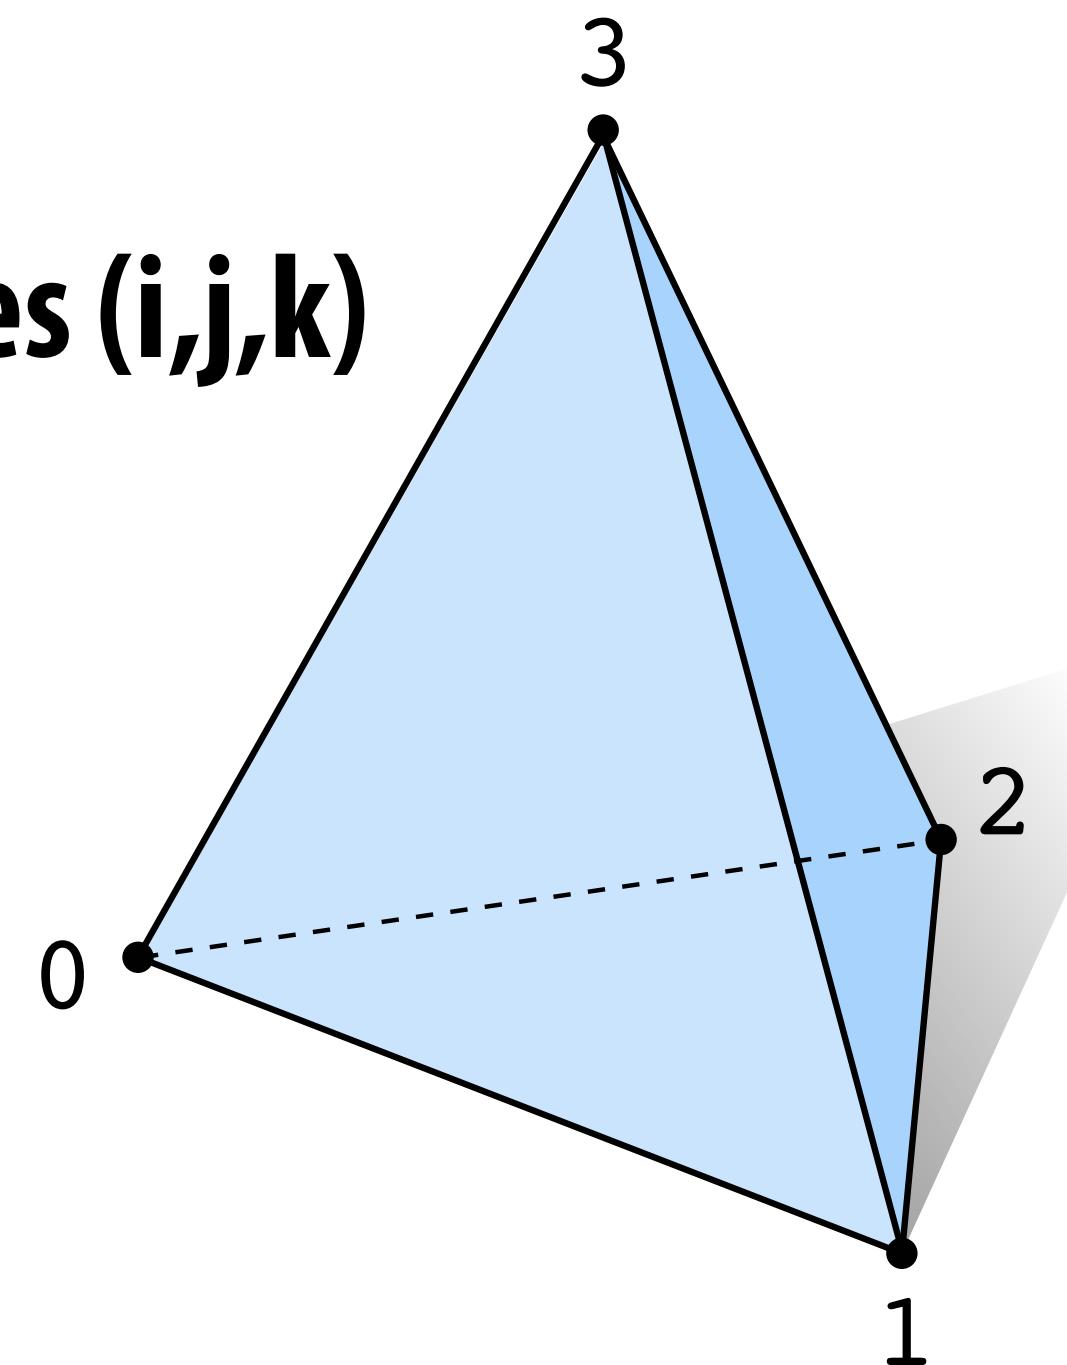


Refresher: Polygon Soup

- Store triples of coordinates (x,y,z) and indices (i,j,k)
- E.g., tetrahedron:

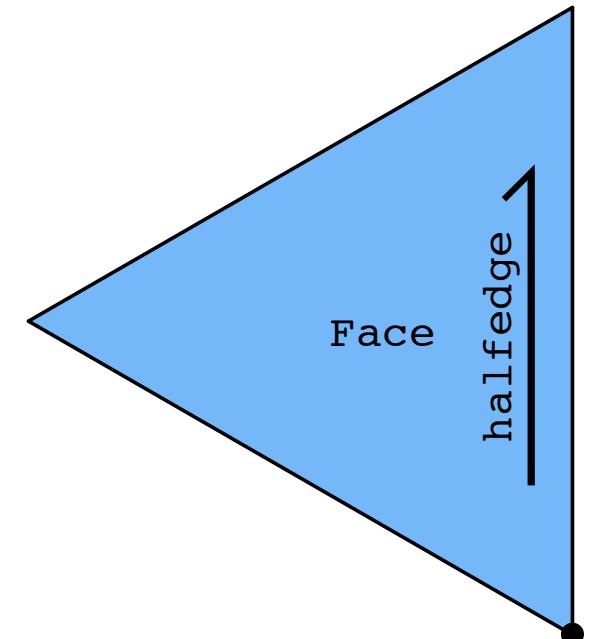
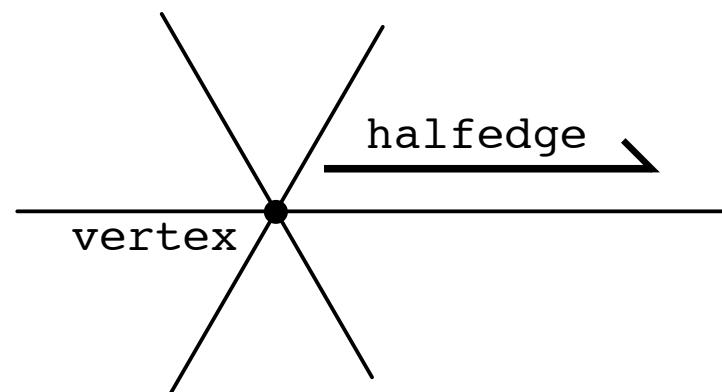
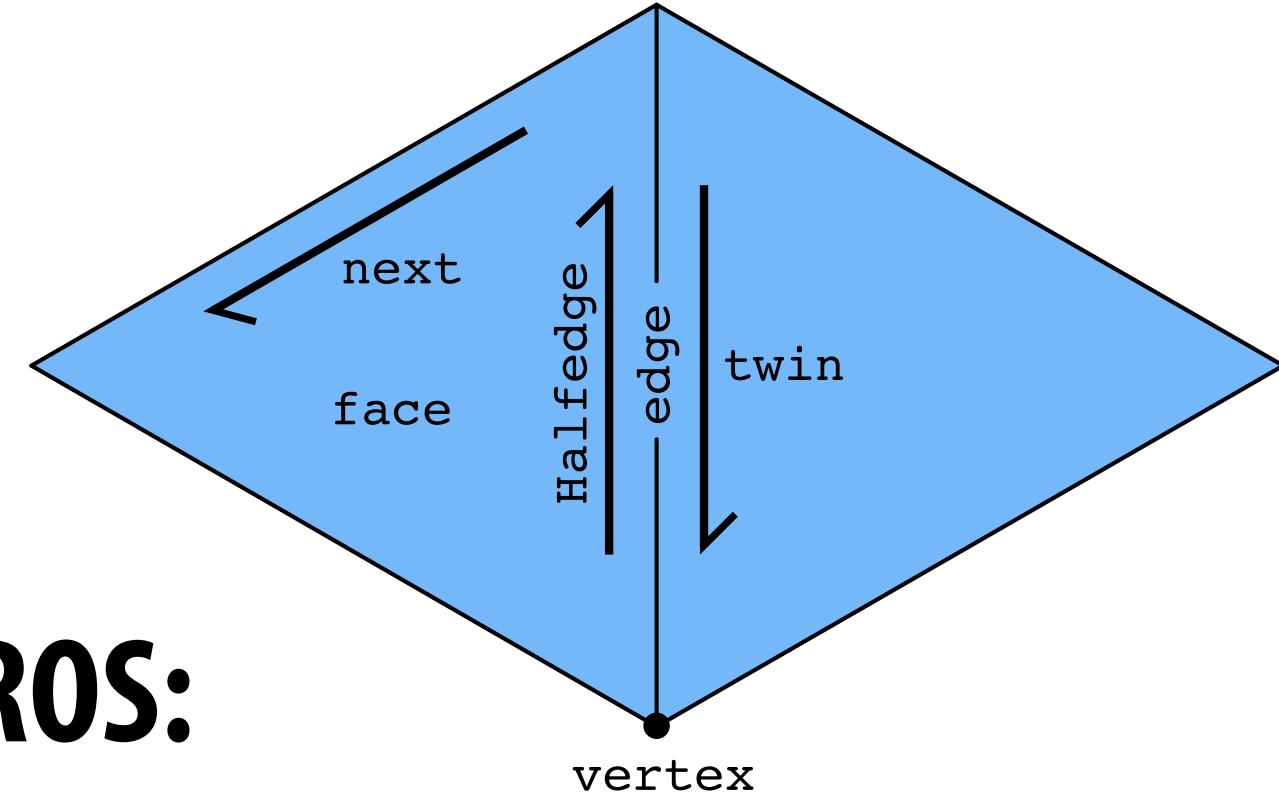
	VERTICES			TRIANGLES		
	x	y	z	i	j	k
0:	-1	-1	-1	0	2	1
1:	1	-1	1	0	3	2
2:	1	1	-1	3	0	1
3:	-1	1	1	3	1	2

- PROS:
 - very low storage cost (how much?)
 - coherent memory access
 - can easily represent nonmanifold geometry
- CONS:
 - very expensive to access neighbors
 - very difficult to change connectivity (how?)



Refresher: Halfedge Mesh

■ Connectivity encoded via *halfedges*



■ PROS:

- easy to access neighbors
- easy to change connectivity

■ CONS:

- greater storage cost
- less coherent memory access
- can only represent manifold, orientable surfaces

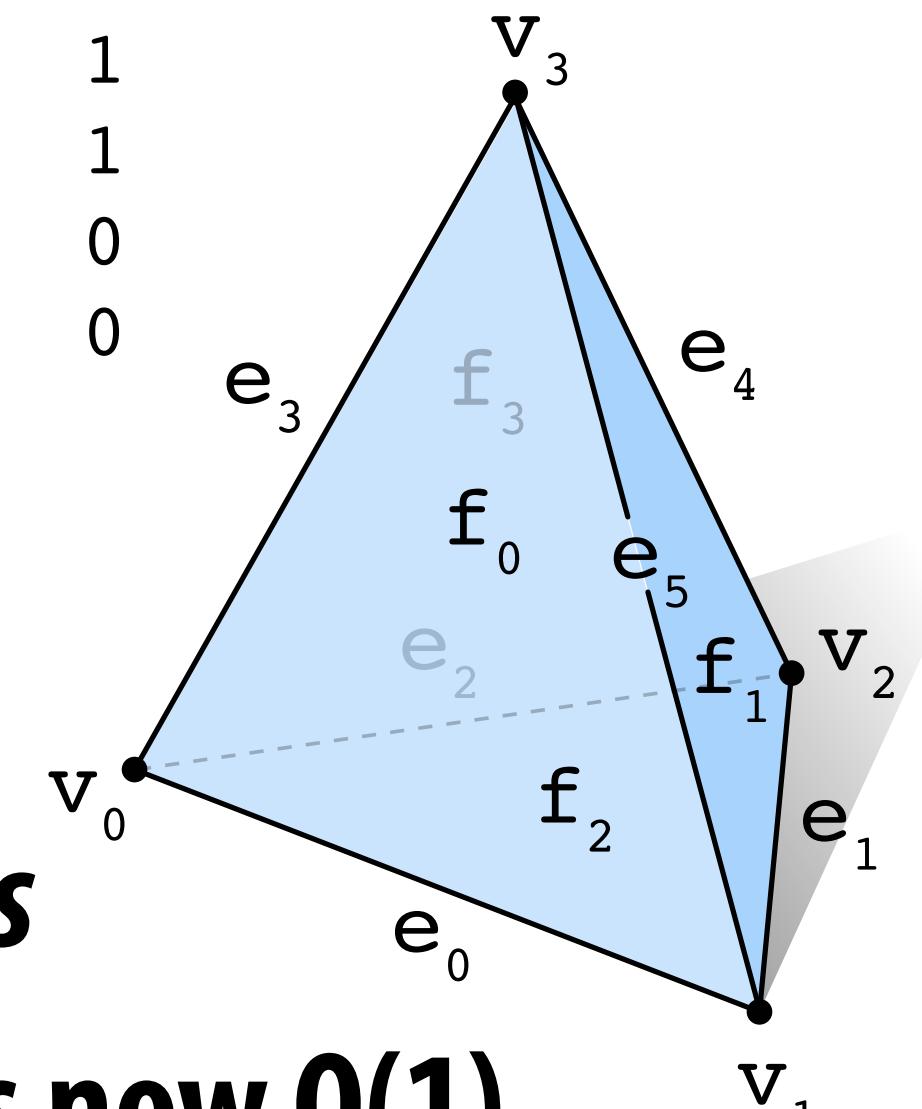
Food for thought: Incidence Matrices

- If we want to answer neighborhood queries, why not simply store a list of neighbors?
- Can encode all neighbor information via *incidence matrices*
- E.g., tetrahedron:

VERTEX	EDGE	FACE												
v ₀	v ₁	v ₂	v ₃	e ₀	e ₁	e ₂	e ₃	e ₄	e ₅	f ₀	f ₁	f ₂	f ₃	
e ₀	1	1	0	0	f ₀	1	0	0	1	0	1	0	1	1
e ₁	0	1	1	0	f ₁	0	1	0	0	1	1	1	1	1
e ₂	1	0	1	0	f ₂	1	1	1	0	0	0	0	0	0
e ₃	1	0	0	1	f ₃	0	0	1	1	1	1	0	0	0
e ₄	0	0	1	1										
e ₅	0	1	0	1										

	v ₀	v ₁	v ₂	v ₃		e ₀	e ₁	e ₂	e ₃	e ₄	e ₅		f ₀	f ₁	f ₂	f ₃
e ₀	1	1	0	0		f ₀	1	0	0	1	0	1				
e ₁	0	1	1	0		f ₁	0	1	0	0	1	1				
e ₂	1	0	1	0		f ₂	1	1	1	0	0	0				
e ₃	1	0	0	1		f ₃	0	0	1	1	1	0				
e ₄	0	0	1	1												
e ₅	0	1	0	1												

- 1 means “touches”; 0 means “does not touch”
- Instead of storing lots of 0's, use *sparse matrices*
- Still large storage cost, but finding neighbors is now O(1)
- Hard to change connectivity, since we used fixed indices
- Bonus feature: mesh does not have to be manifold



Comparison of Polygon Mesh Data Structures

	Polygon Soup	Incidence Matrices	Halfedge Mesh
storage cost*	$\sim 3 \times \#vertices$	$\sim 33 \times \#vertices$	$\sim 36 \times \#vertices$
constant-time neighborhood access?	NO	YES	YES
easy to add/remove mesh elements?	NO	NO	YES
nonmanifold geometry?	YES	YES	NO

Conclusion: pick the right data structure for the job!

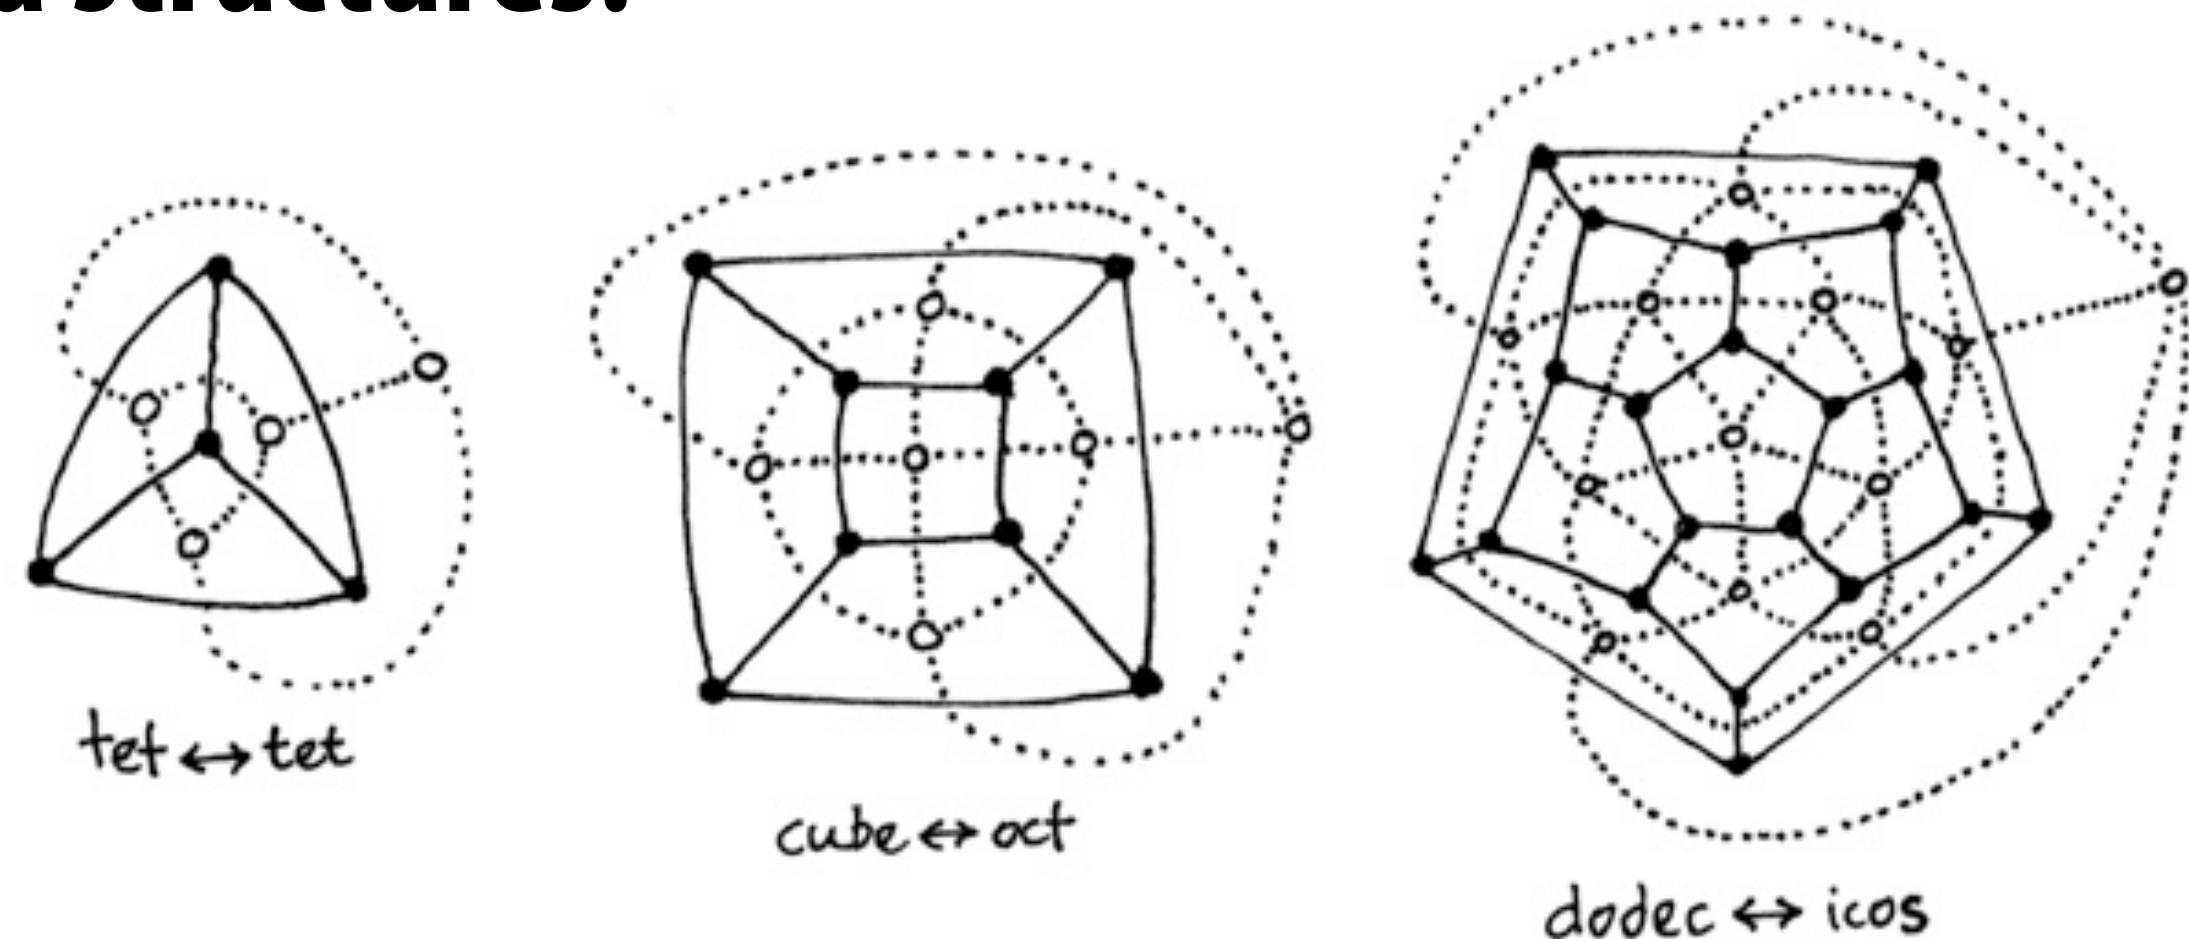
*number of integer values and/or pointers required to encode *connectivity*
(all data structures require same amount of storage for vertex positions)

Even more alternatives to Halfedge

■ Many very similar data structures:

- winged edge
- corner table
- quadedge
- ...

Paul Heckbert (former CMU prof.)
quadedge code - <http://bit.ly/1QZLHos>



■ Each stores local neighborhood information

■ Similar tradeoffs relative to simple polygon list:

- **CONS**: additional storage, incoherent memory access
 - **PROS**: better access time for individual elements, intuitive traversal of local neighborhoods
- (Food for thought: can you design a halfedge-like data structure with reasonably coherent data storage?)

So, what's the job?

Digital Geometry Processing

- Extend traditional digital signal processing (audio, video, etc.) to deal with *geometric* signals:
 - upsampling / downsampling / resampling / filtering ...
 - aliasing (reconstructed surface gives “false impression”)
- Also some new challenges (very recent field!):
 - over which domain is a geometric signal expressed?
 - no terrific sampling theory, no fast Fourier transform, ...
- Often need new data structures & new algorithms

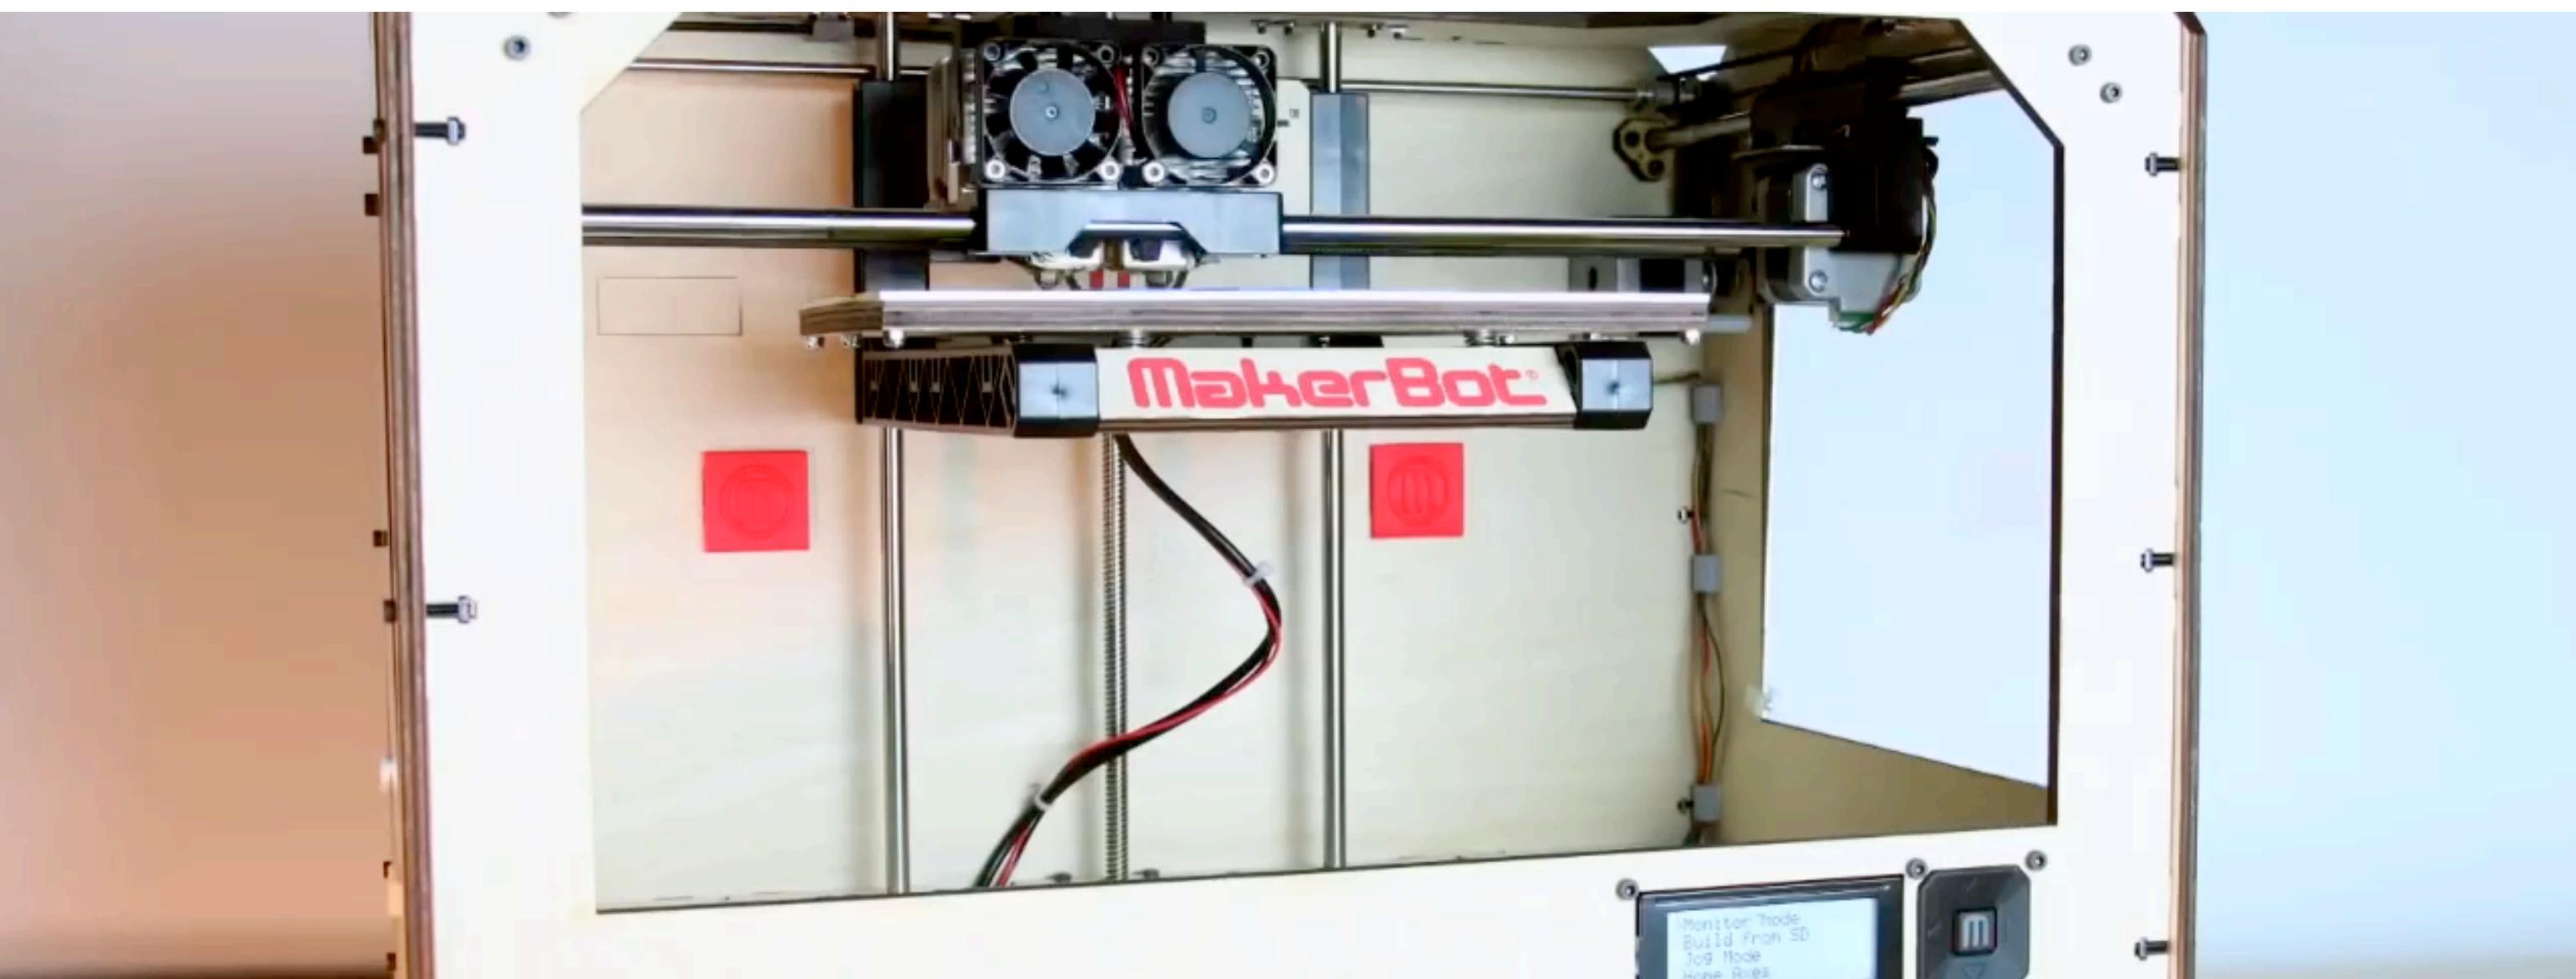


Digital Geometry Processing: Motivation

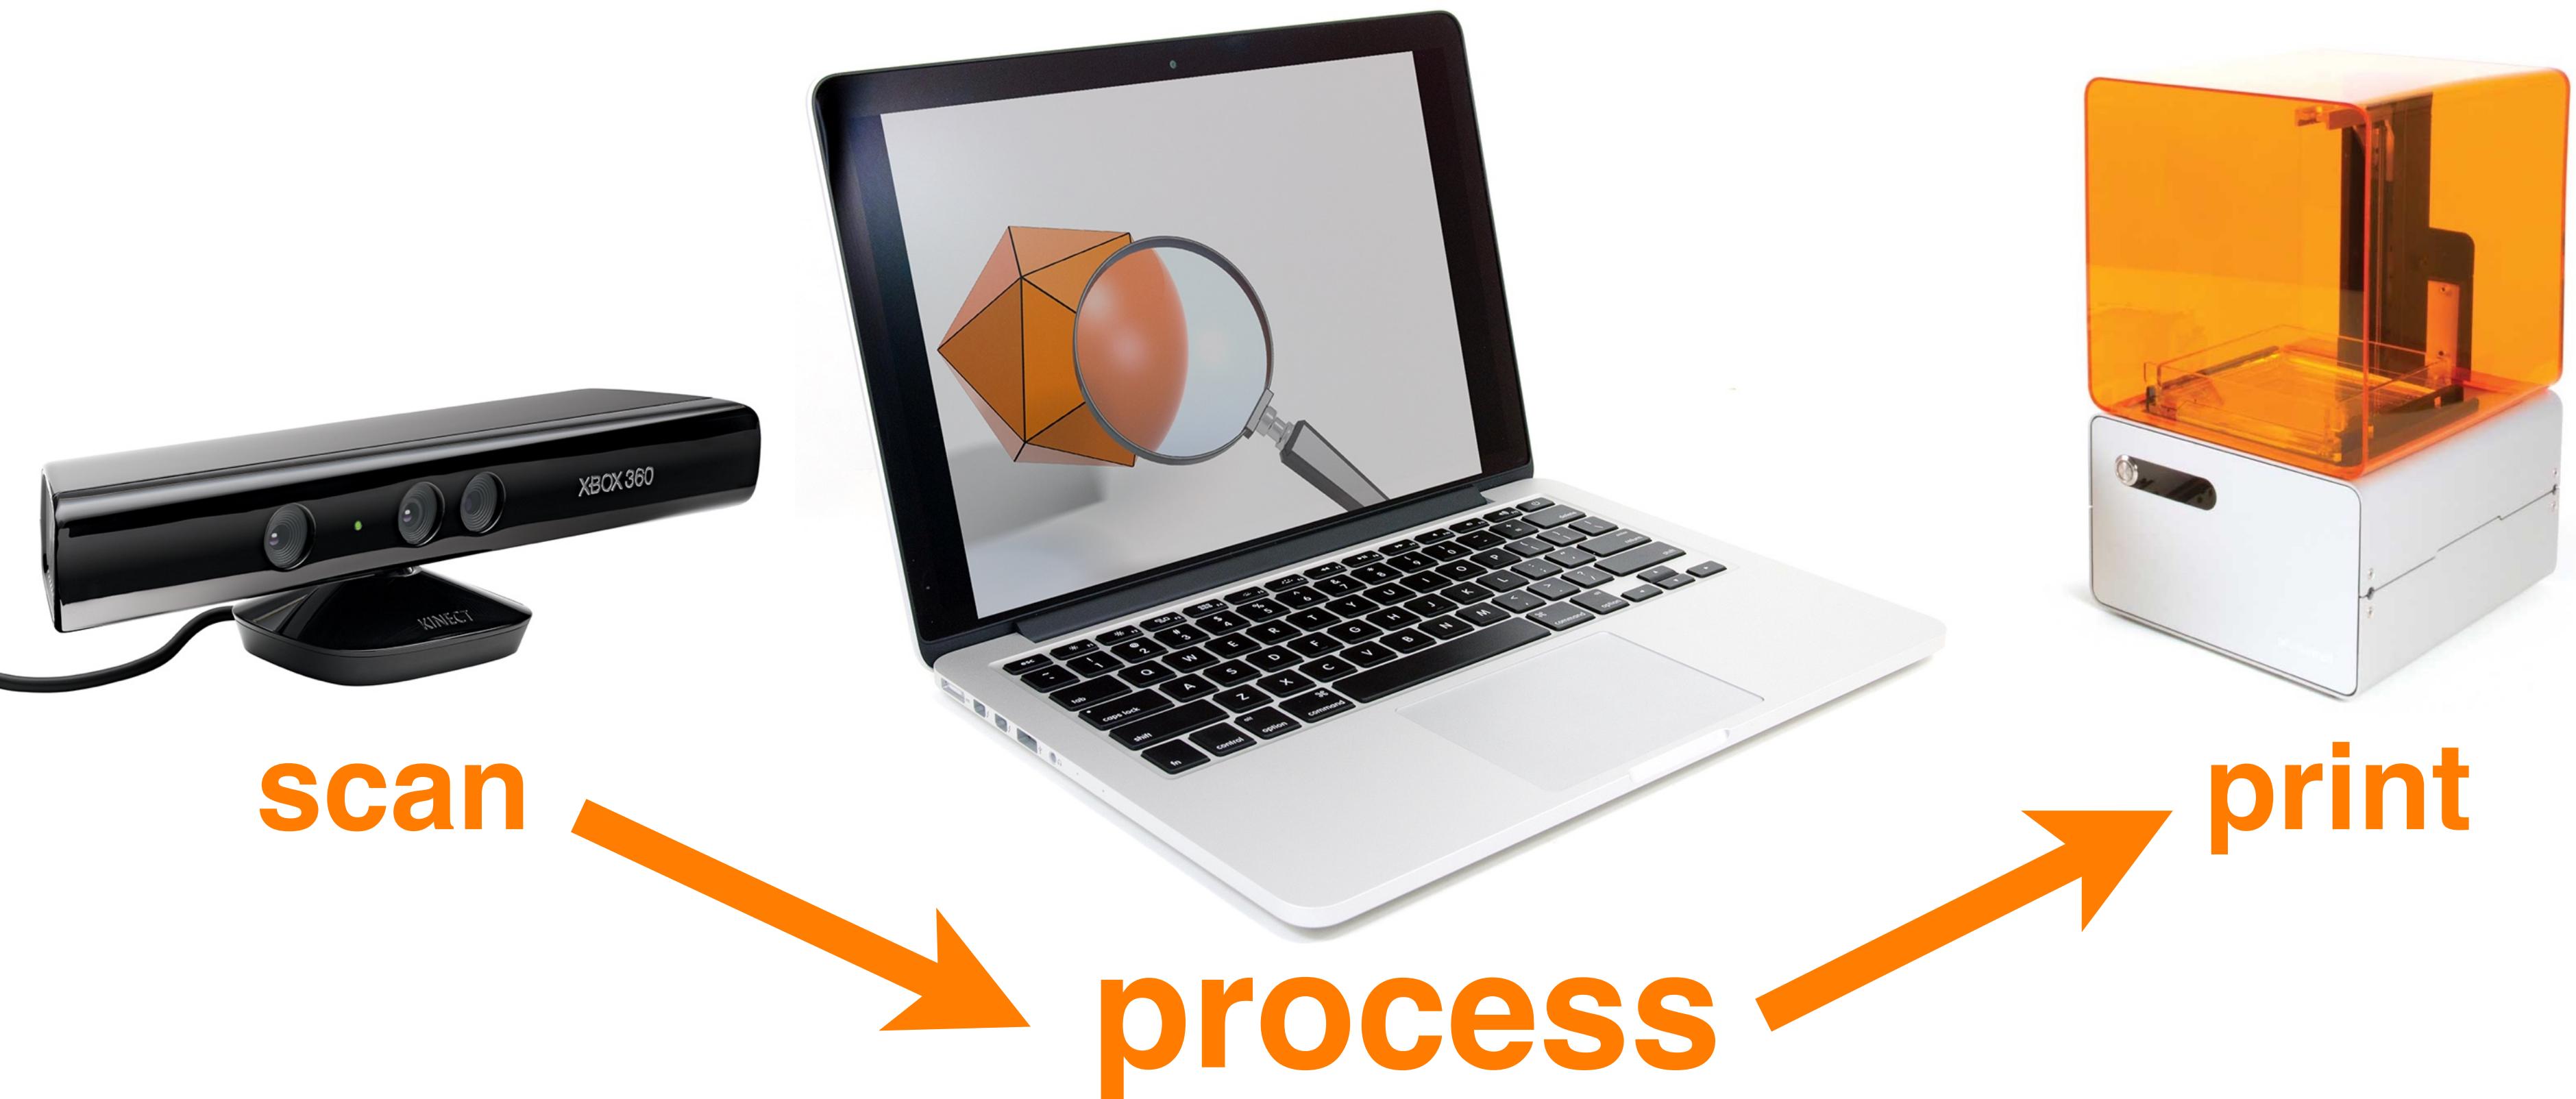
3D Scanning



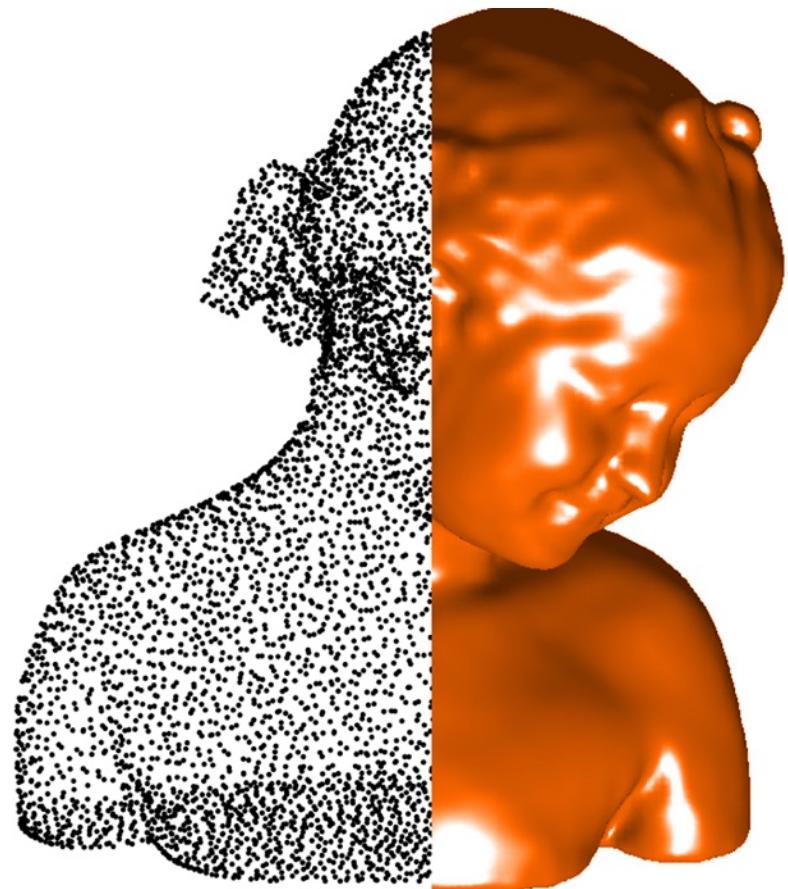
3D Printing



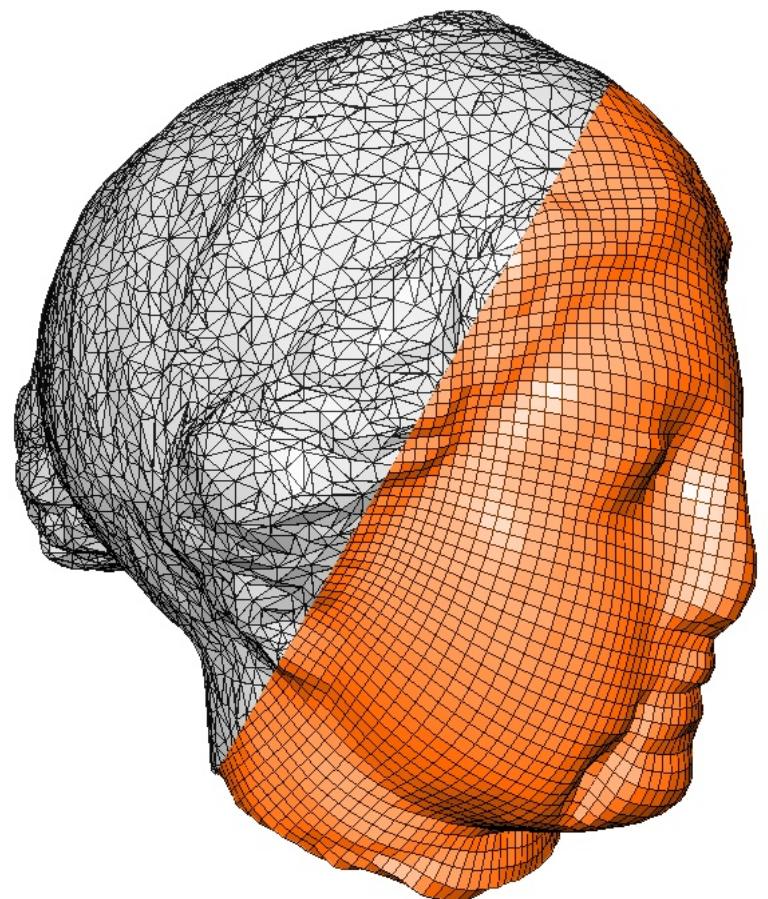
Geometry Processing Pipeline



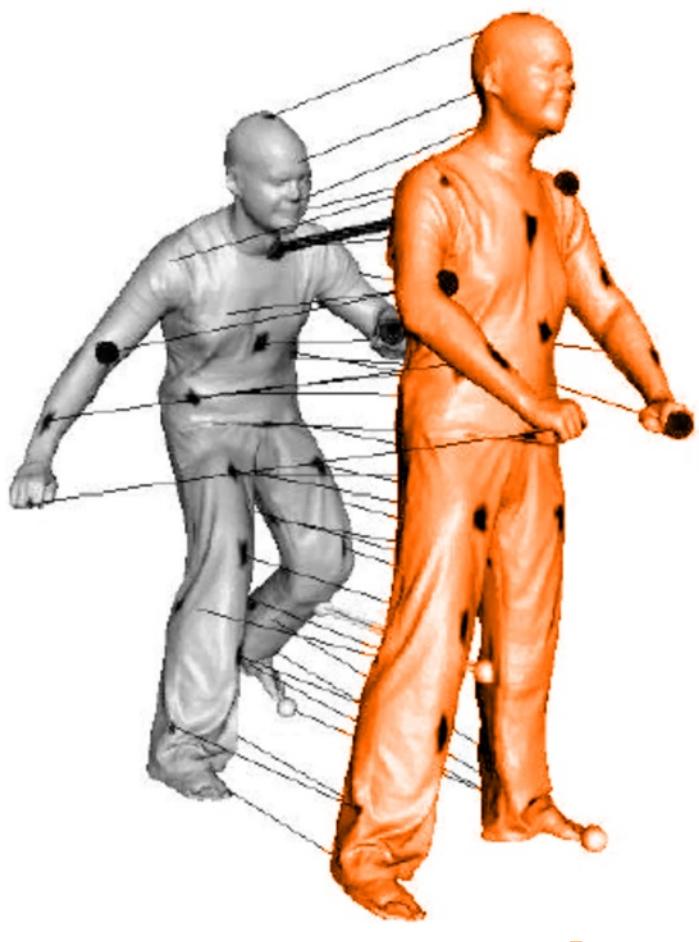
Geometry Processing Tasks



reconstruction



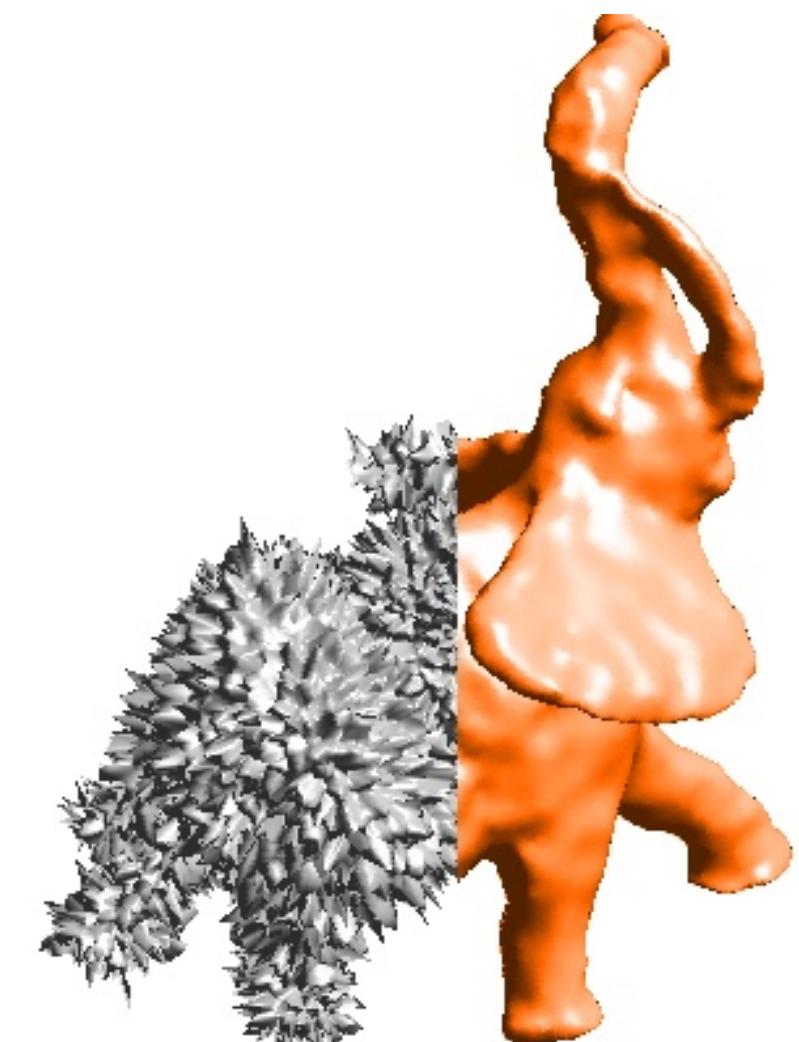
remeshing



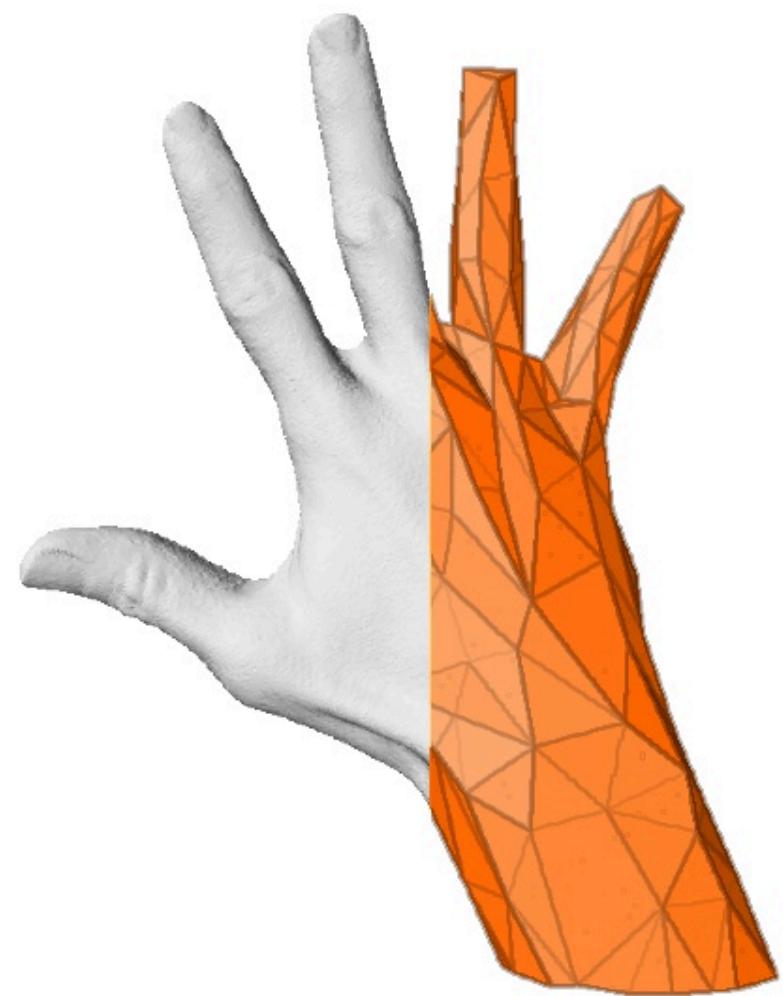
shape analysis



parameterization



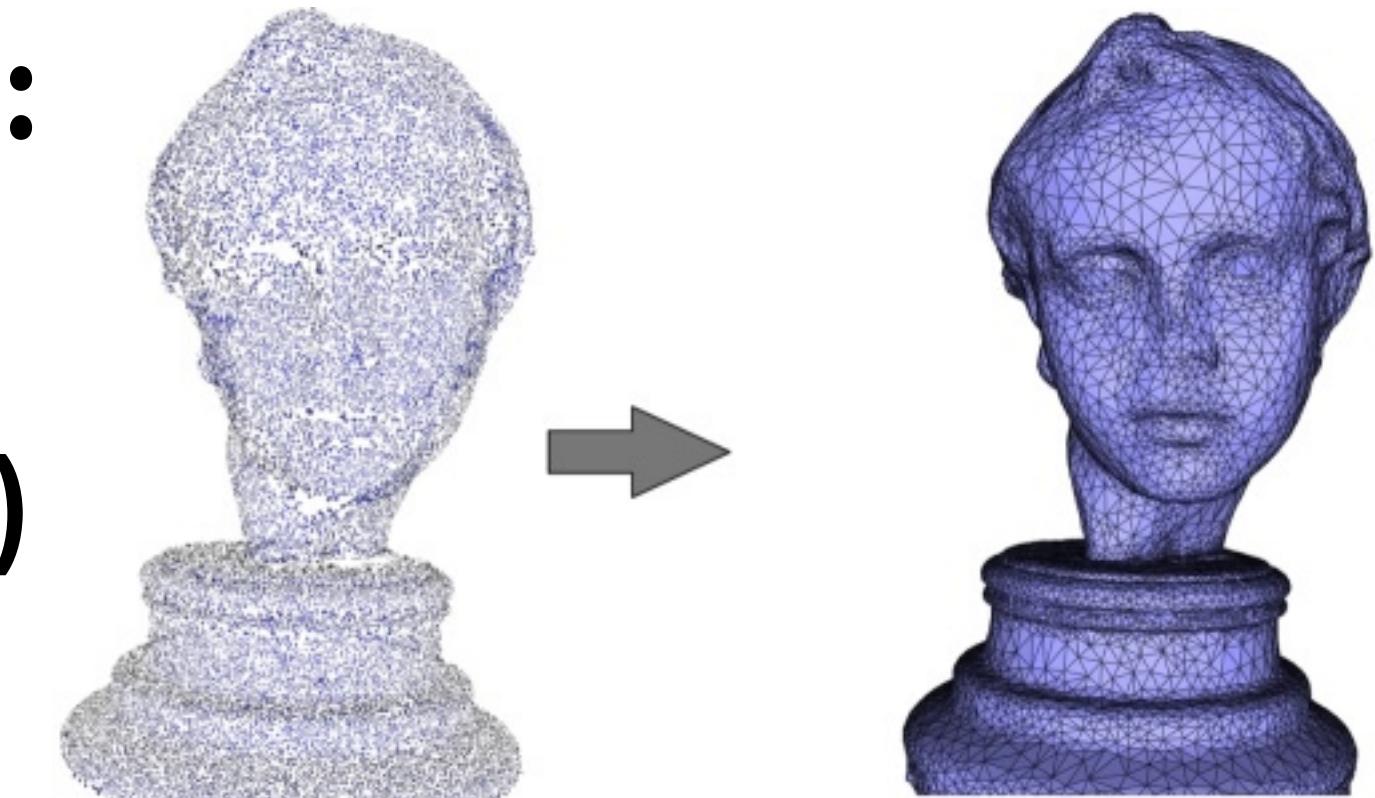
filtering



compression

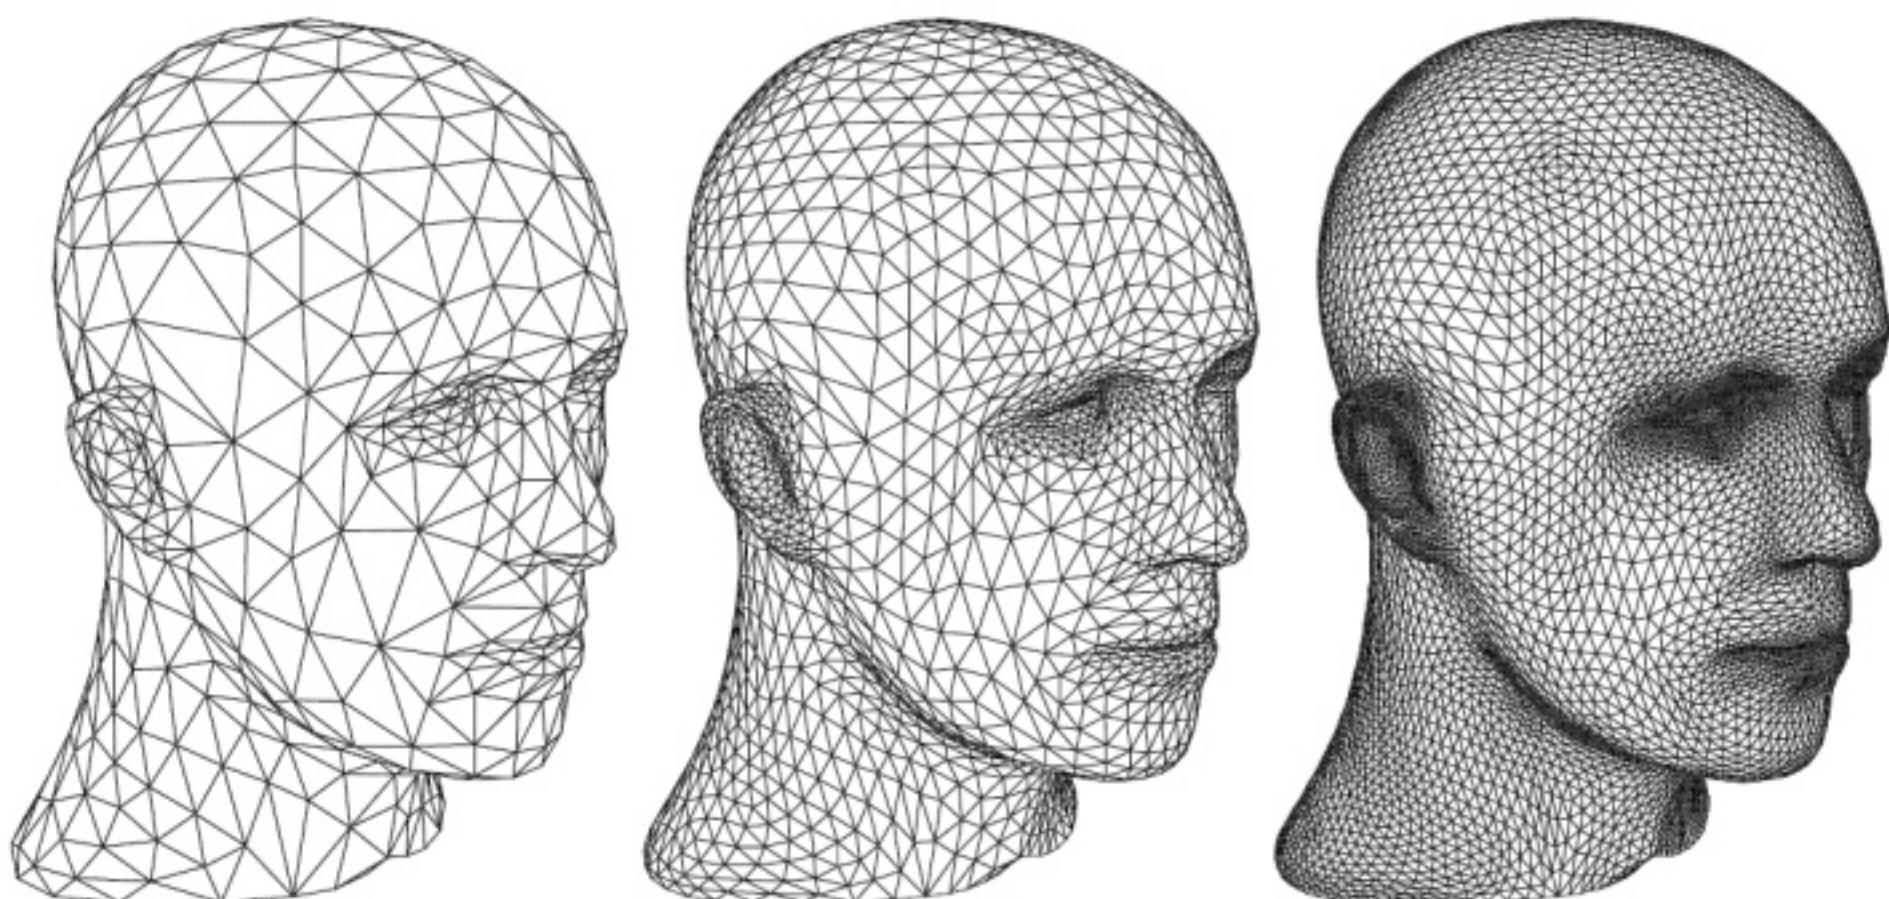
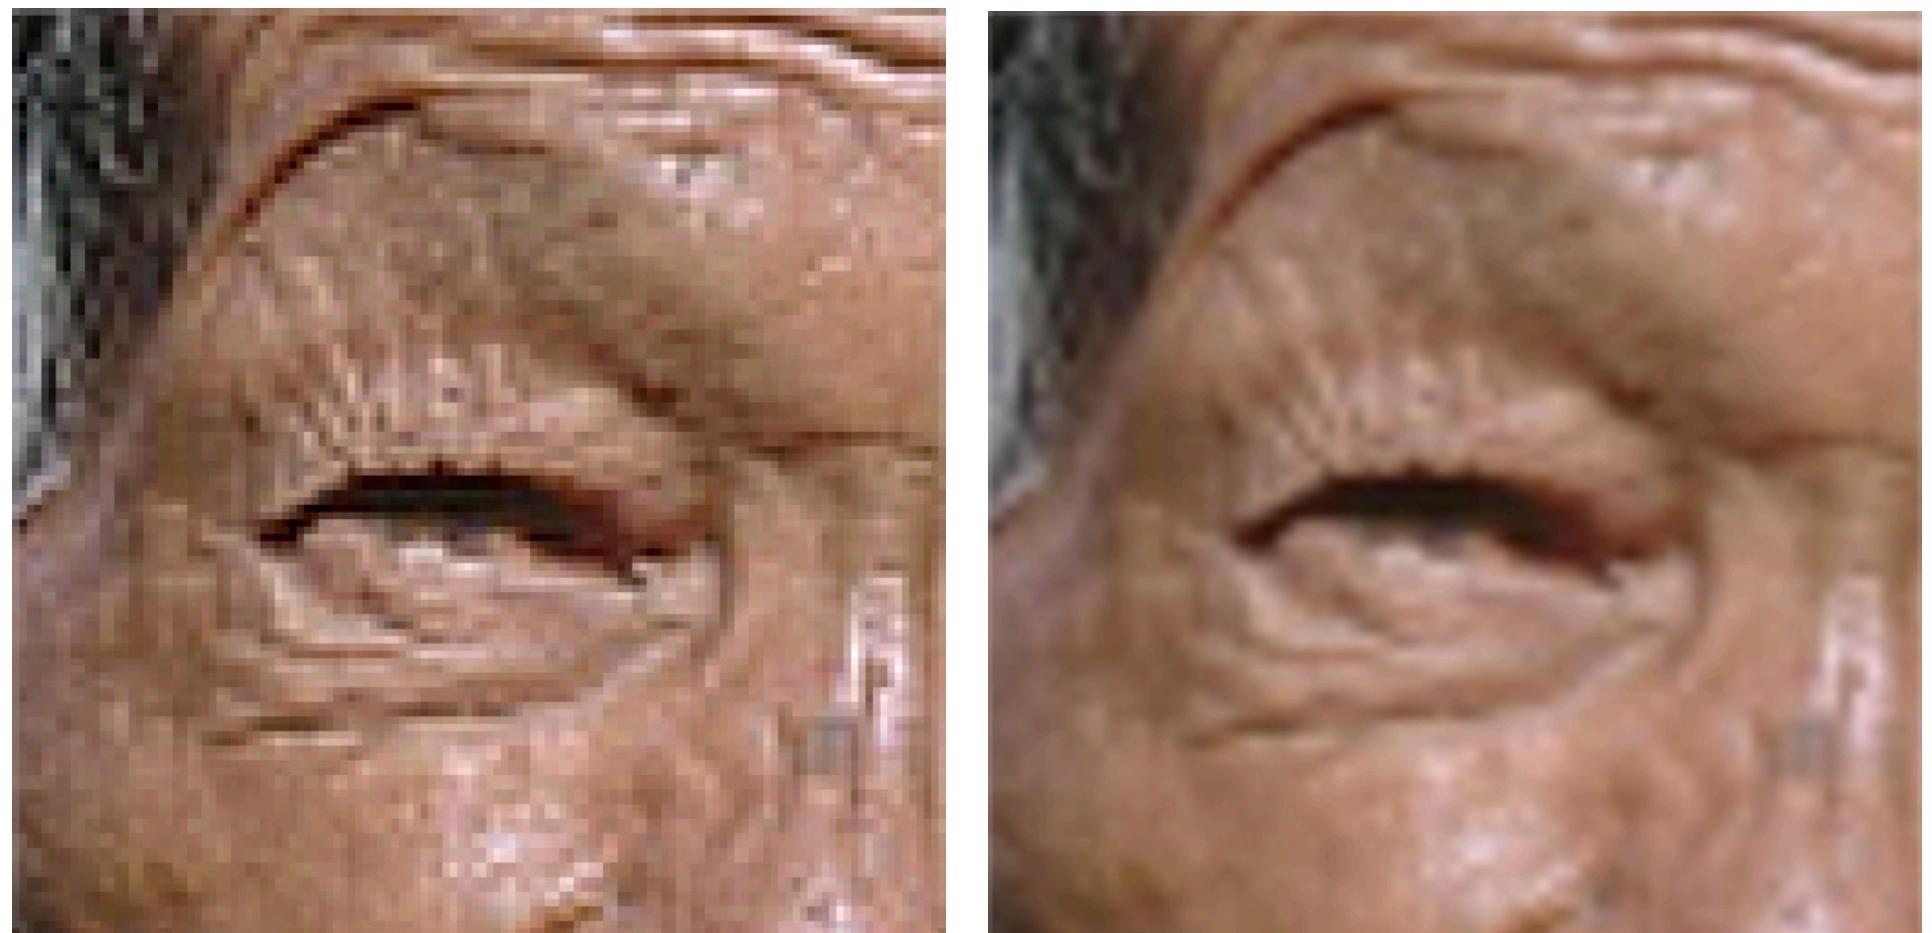
Geometry Processing: Reconstruction

- Given samples of geometry, reconstruct continuous surface
- What are “samples”? Many possibilities:
 - points, points & normals, ...
 - image pairs / sets (multi-view stereo)
 - line density integrals (MRI/CT scans)
- How do you get a surface? Many techniques:
 - PDE-based (e.g., Poisson reconstruction)
 - Voronoi-based (e.g., power crust)
 - silhouette-based (visual hull)
 - Radon transform / isosurfacing (marching cubes)



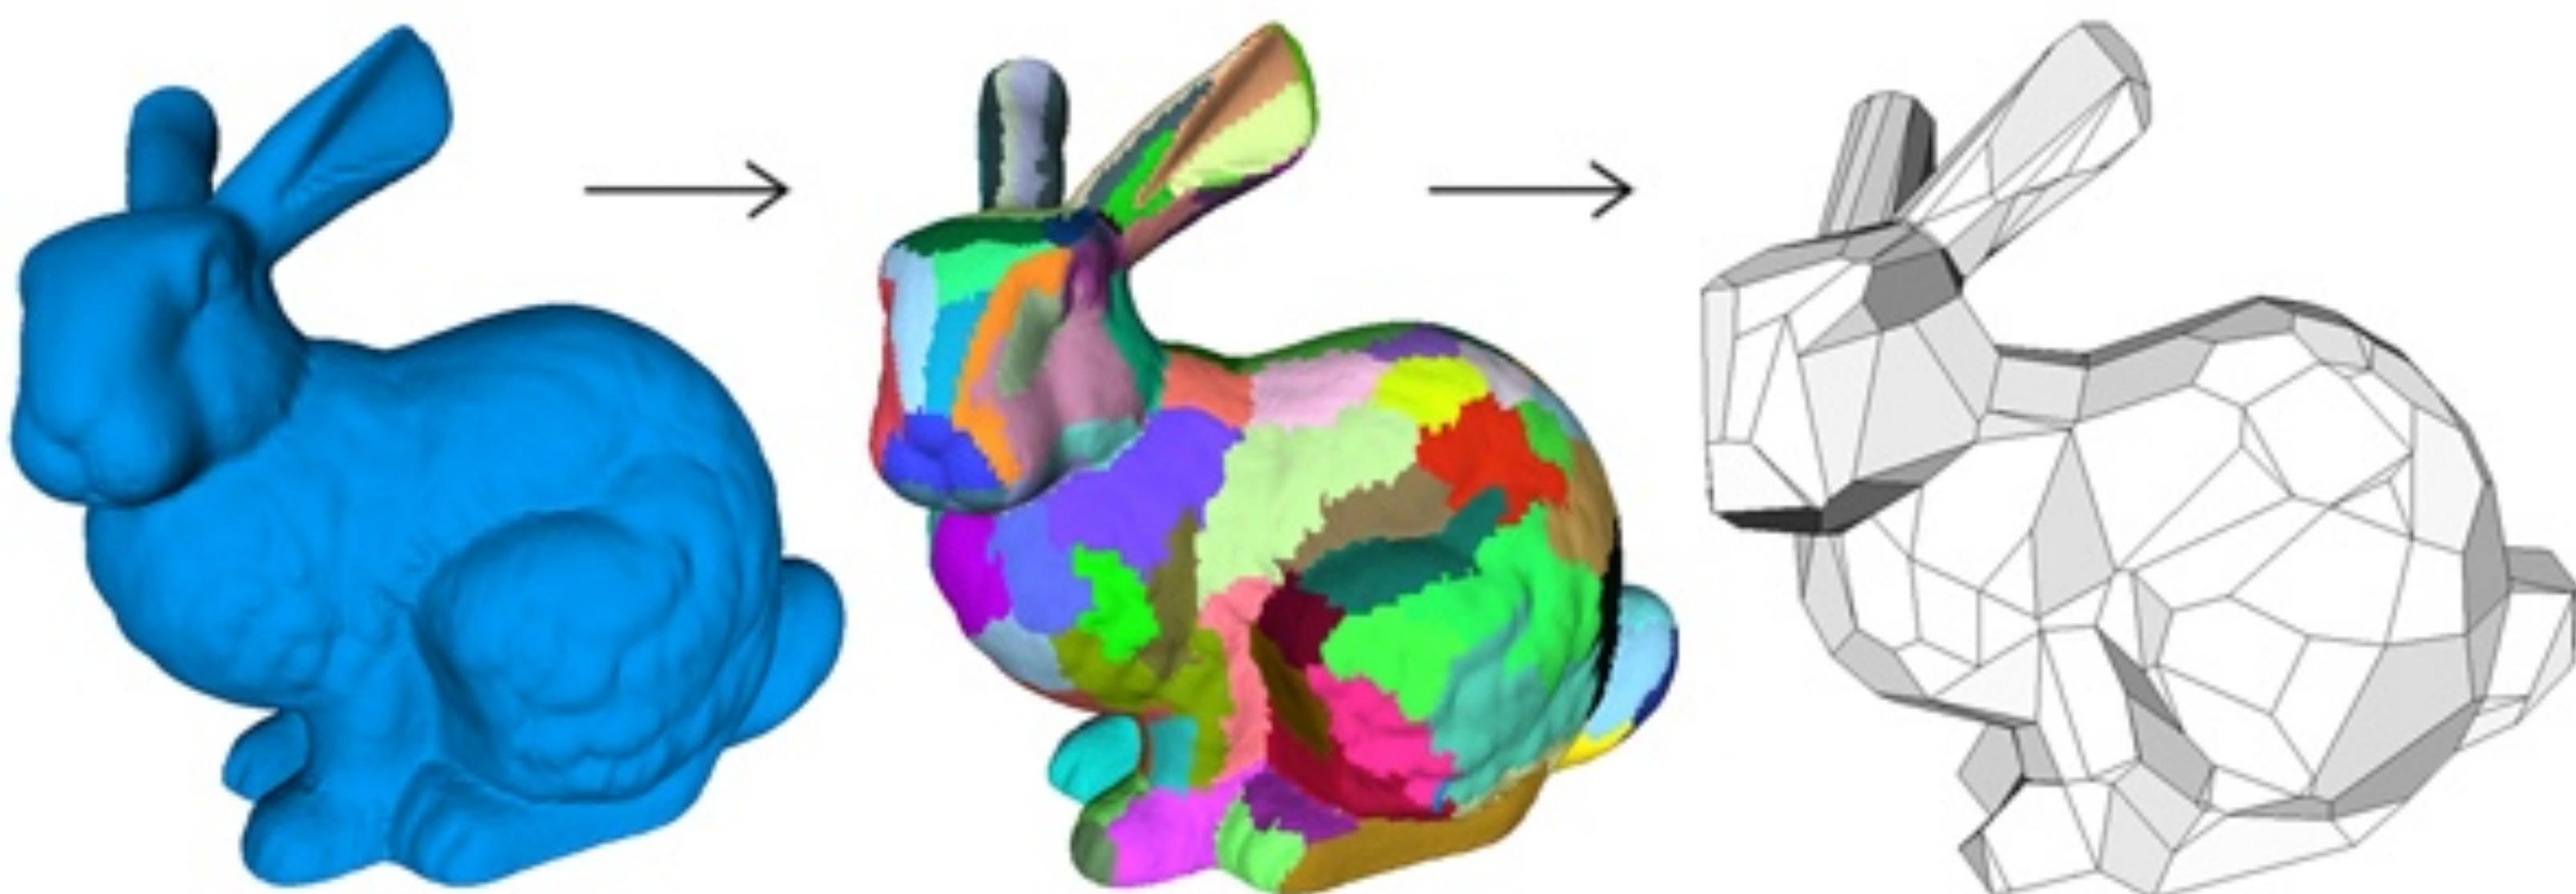
Geometry Processing: Upsampling

- Increase resolution via interpolation
- Images: e.g., bilinear, bicubic interpolation
- Polygon meshes:
 - subdivision
 - bilateral upsampling
 - ...



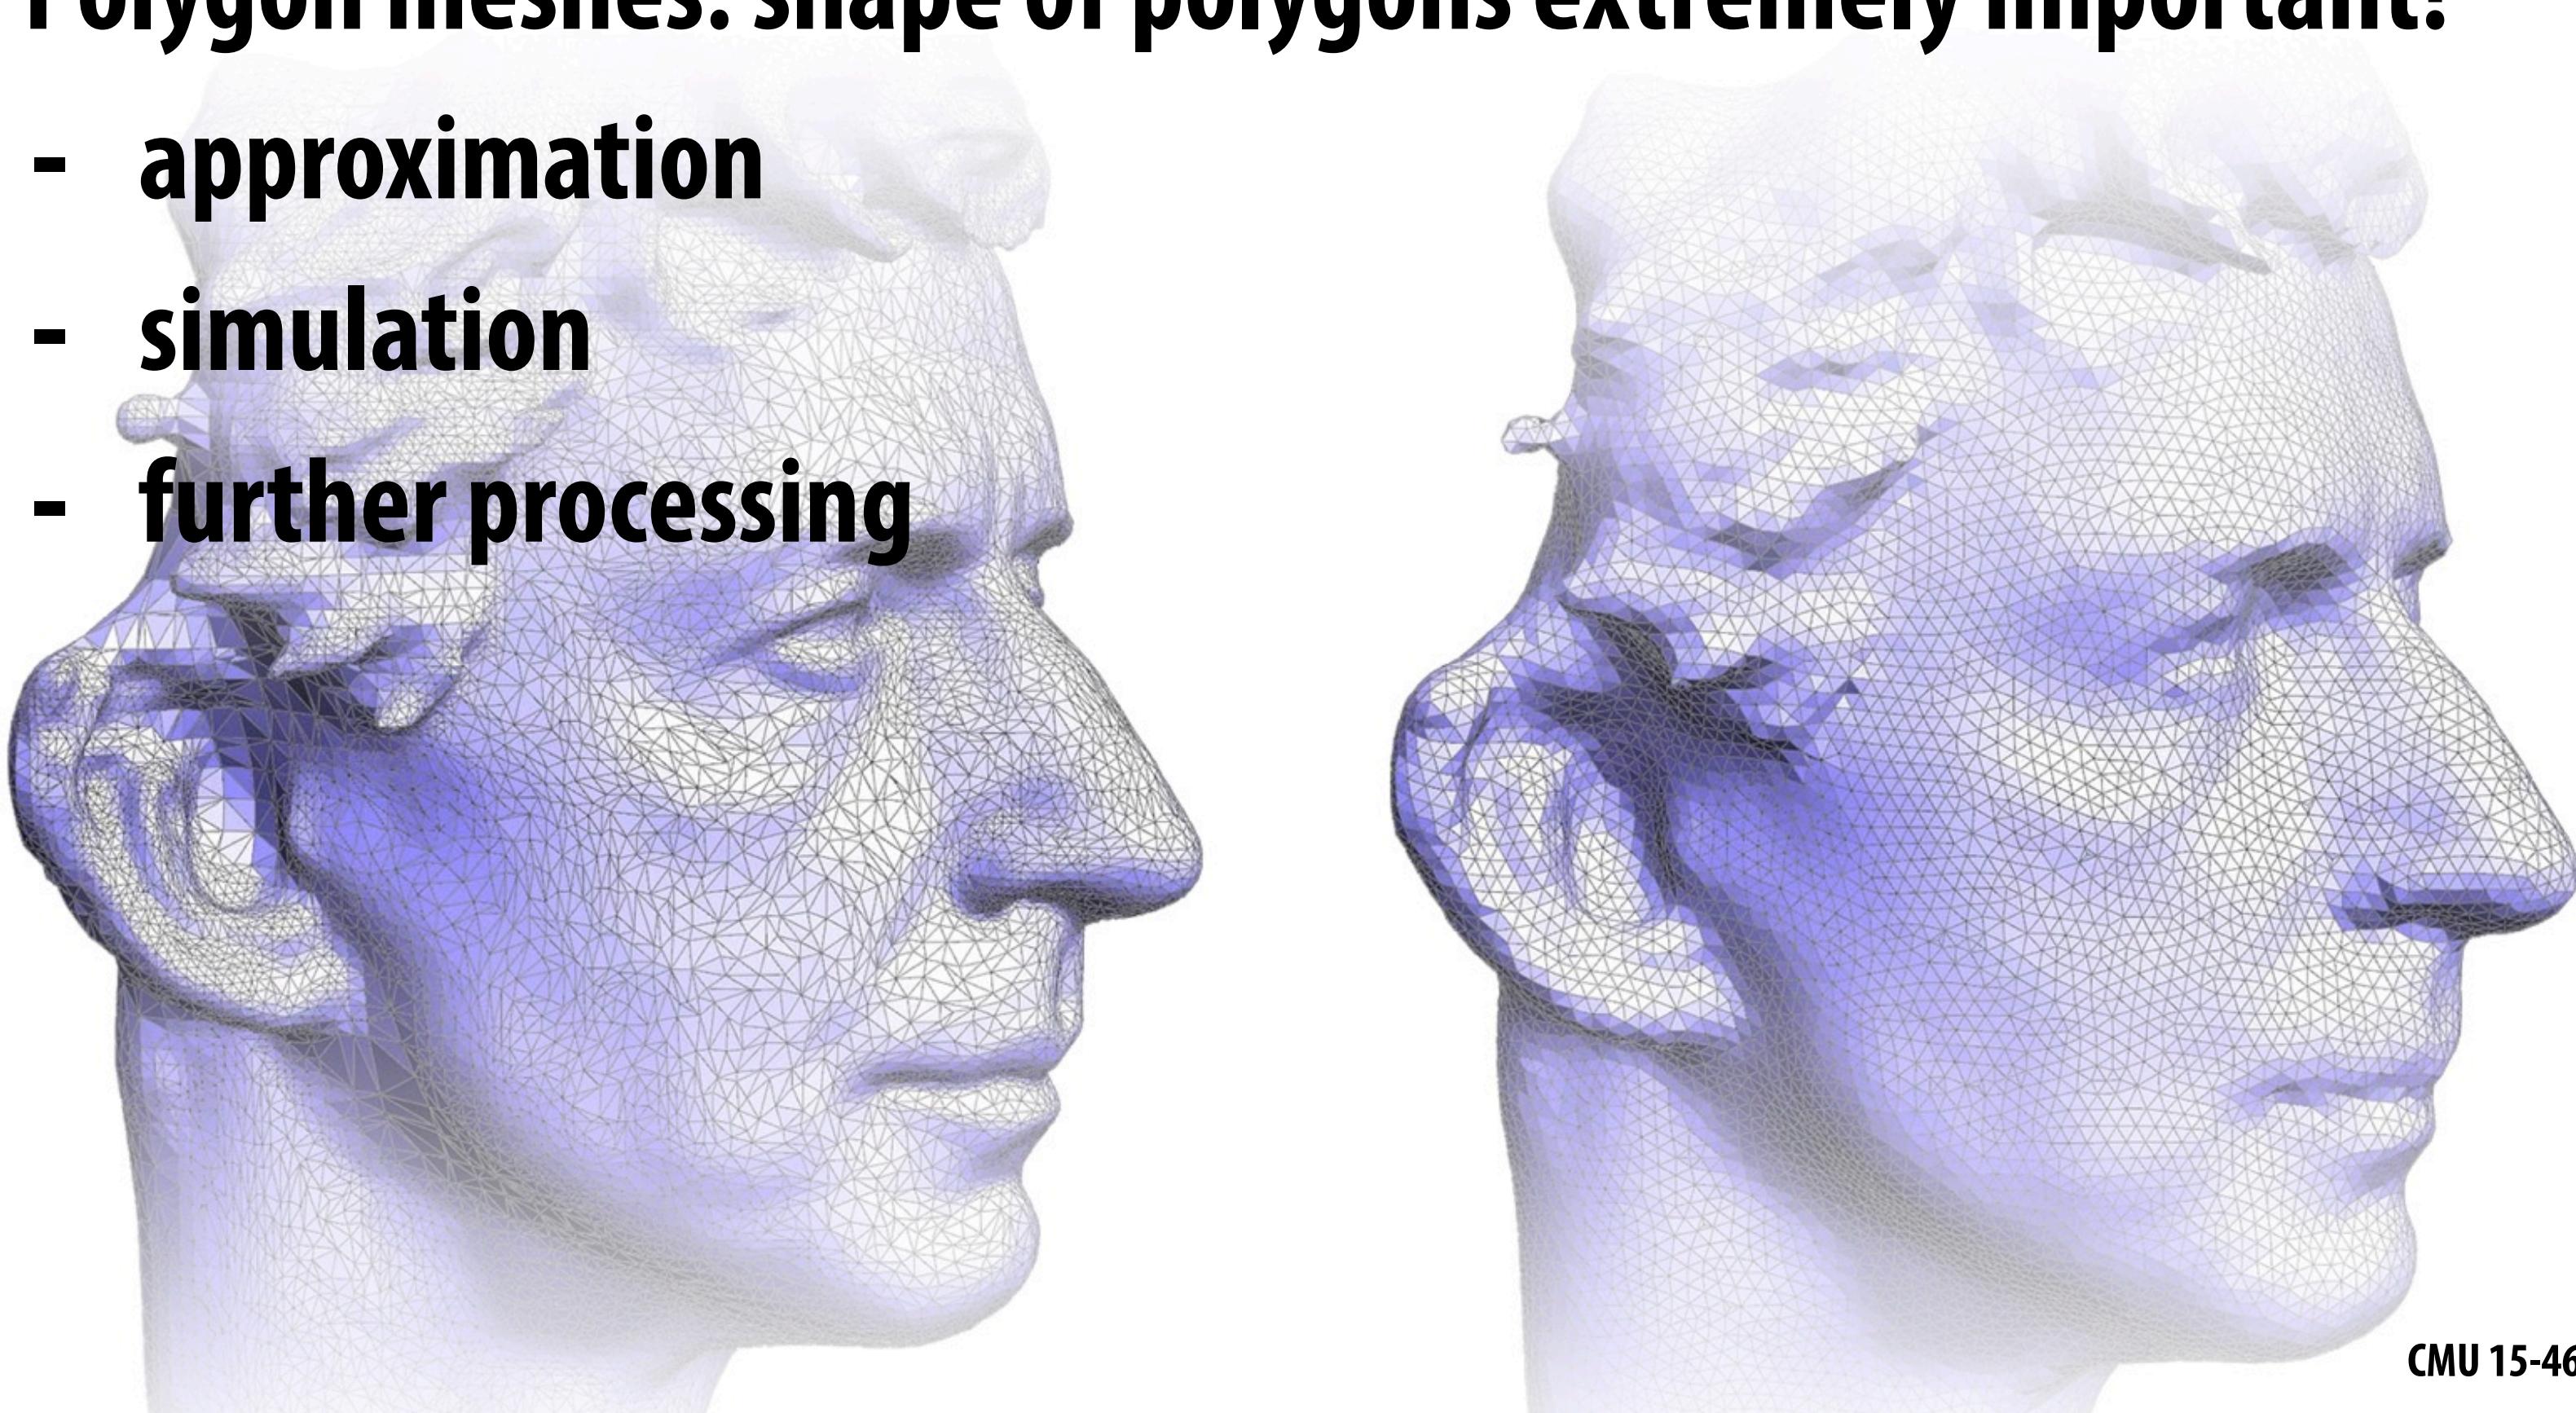
Geometry Processing: Downsampling

- Decrease resolution; try to preserve shape/appearance
- Images: again, bilinear, bicubic interpolation (again)
- Polygon meshes:
 - iterative decimation
 - variational shape approximation
 - ...



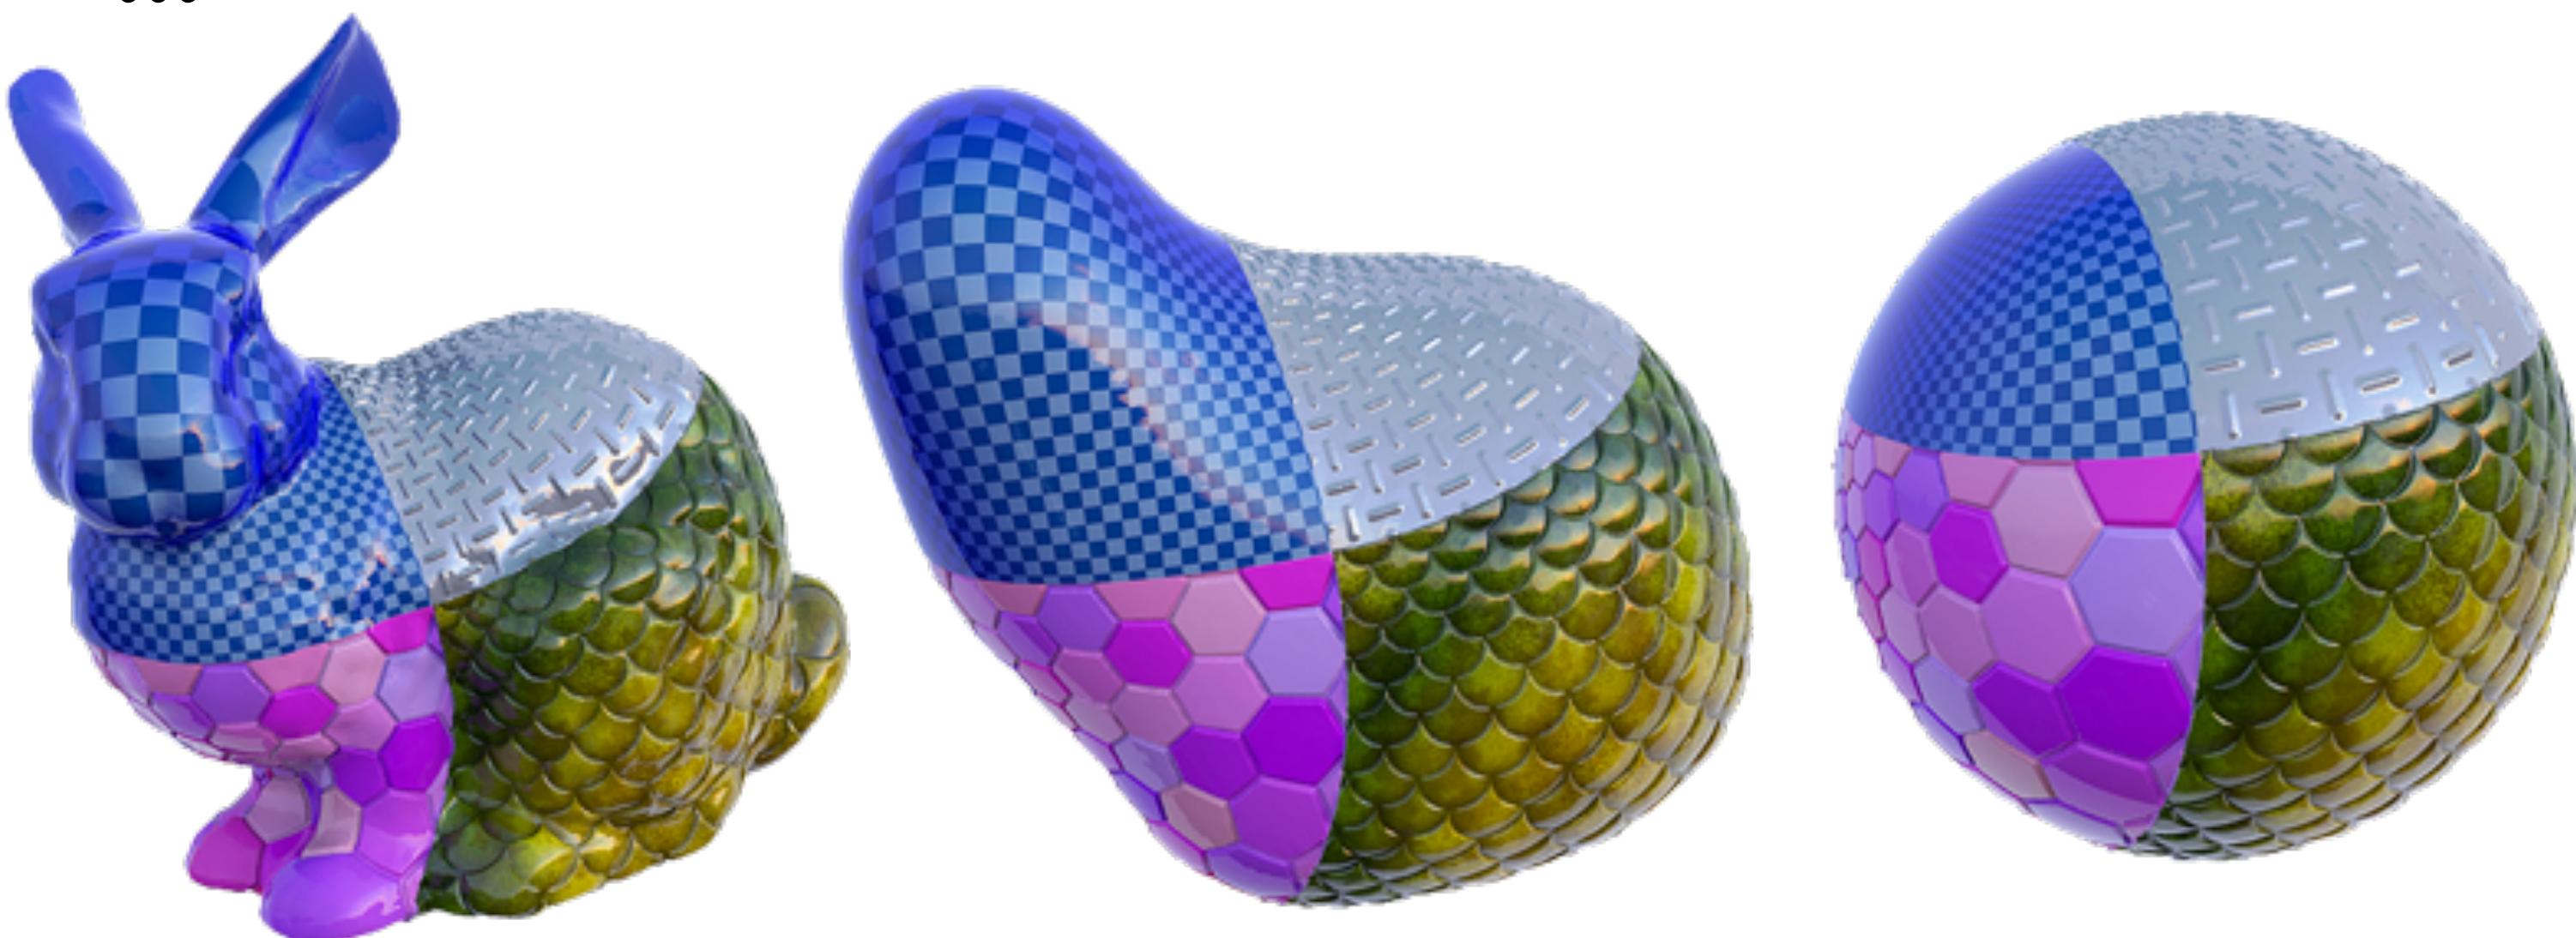
Geometry Processing: Resampling

- **Modify sample distribution to improve quality**
- **Images: ...not usually an issue!**
 - pixels are always stored on a regular grid
- **Polygon meshes: shape of polygons extremely important!**
 - approximation
 - simulation
 - further processing



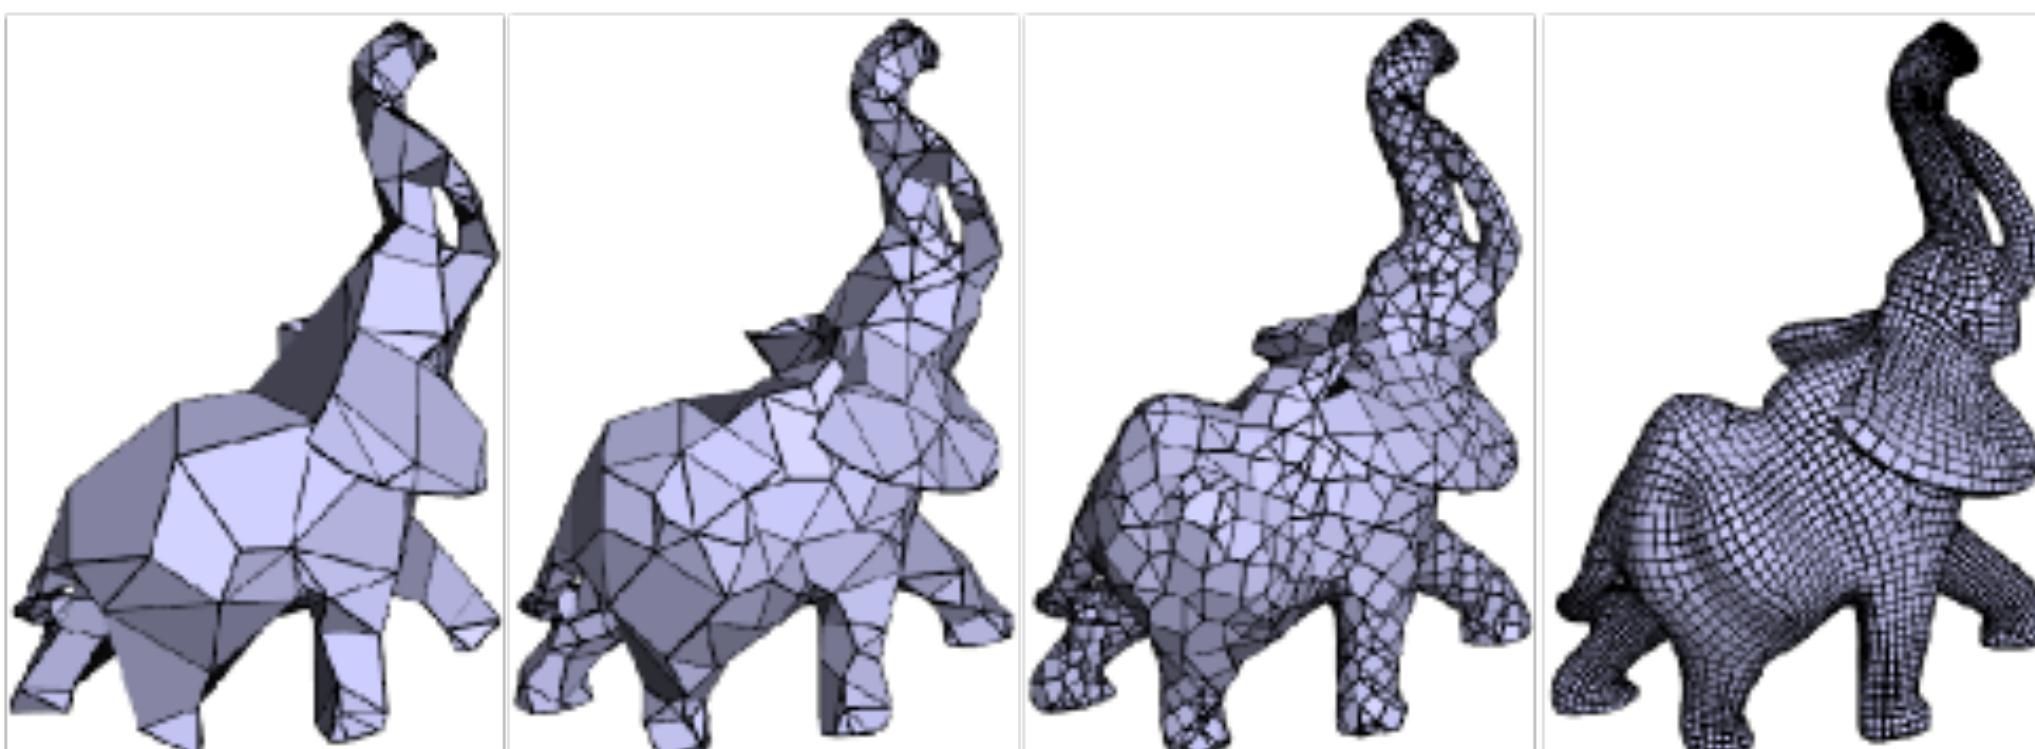
Geometry Processing: Filtering

- Remove noise, or emphasize important features (e.g., edges)
- Images: blurring, bilateral filter, compressed sensing, ...
- Polygon meshes:
 - curvature flow
 - bilateral filter
 - ...



Geometry Processing: Compression

- Reduce storage size by eliminating redundant data/approximating unimportant data
- Images:
 - run-length encoding (RLE) - no loss of information
 - spectral/wavelet encoding (e.g., JPEG/MPEG) - lossy
- Polygon meshes:
 - have to compress geometry *and* connectivity
 - many techniques (spectral, diffusion, ...)



Geometry Processing: Shape Analysis

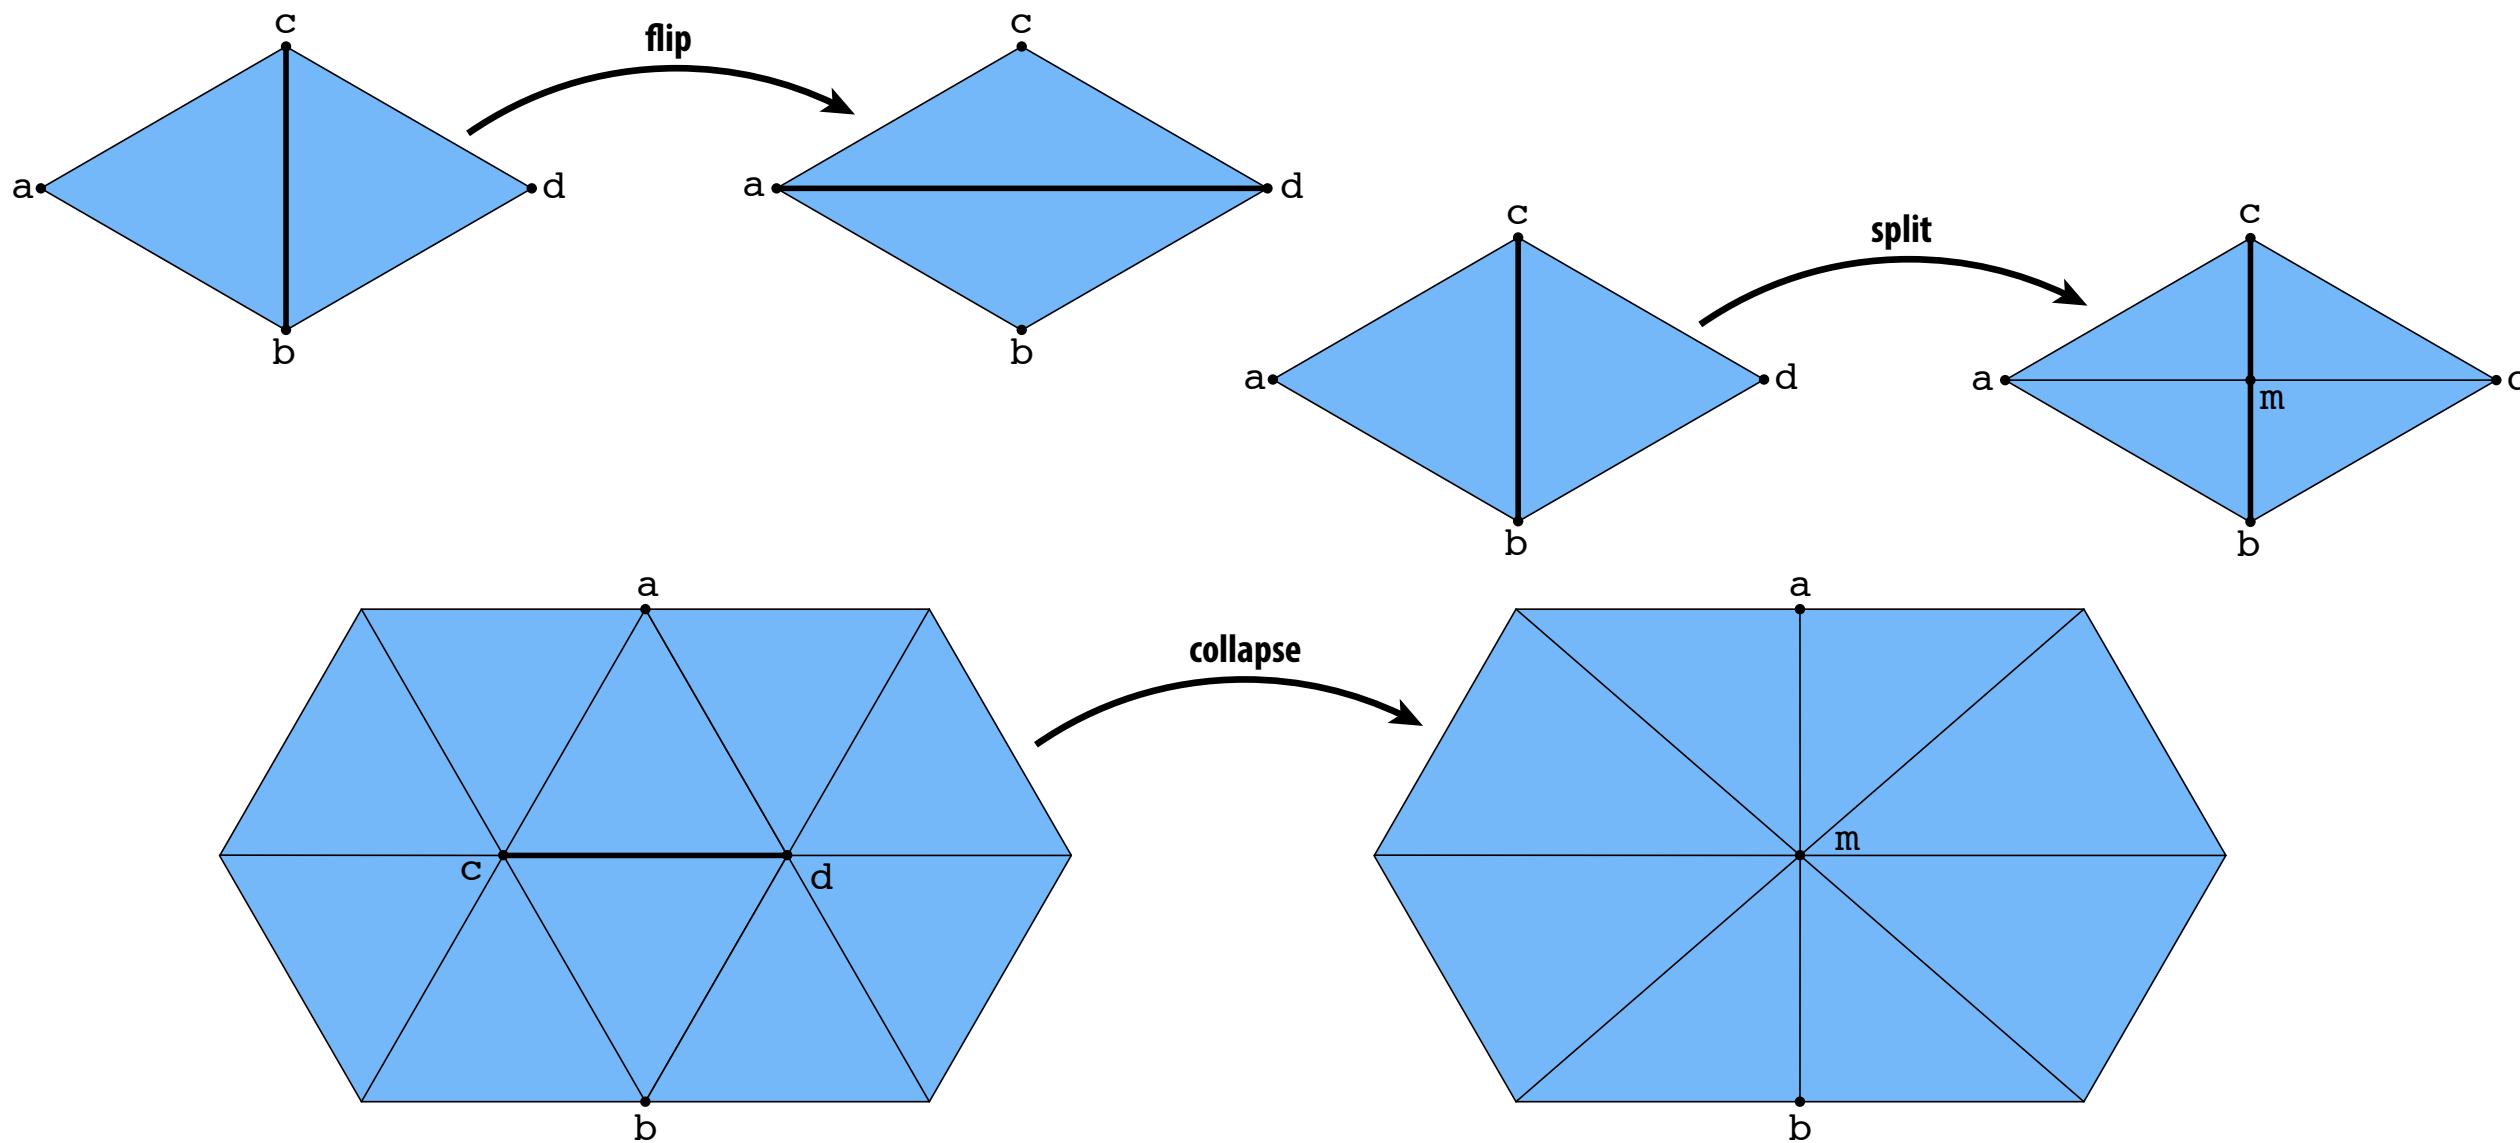
- Identify/understand important semantic features
- Images: computer vision, segmentation, face detection, ...
- Polygon meshes:
 - segmentation
 - correspondence
 - symmetry detection
 - ...



0k, enough! Let's process some geometry!

Processing geometry with halfedges

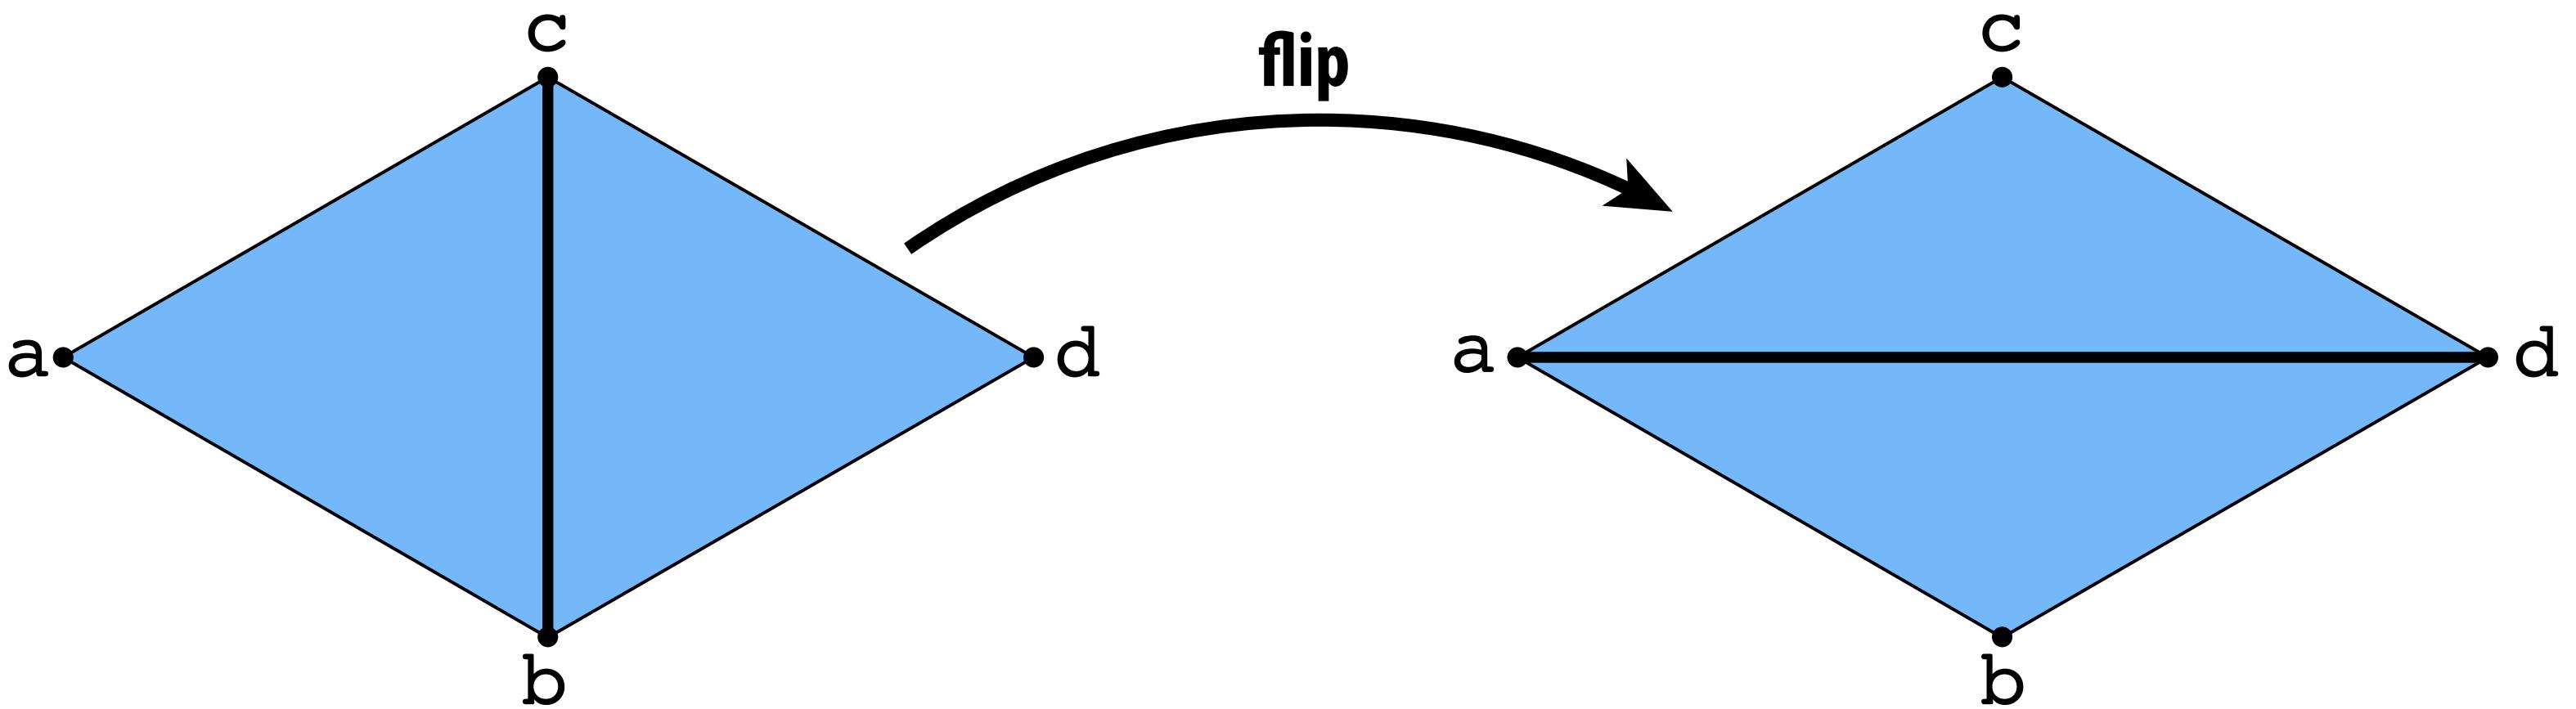
- Remember key feature of linked list: insert/delete elements
- Same story with halfedge mesh (“linked list on steroids”)
- Several atomic operations for triangle meshes:



- How? Allocate/delete elements; reassigning pointers.
- (Should be careful to preserve manifoldness!)

Edge Flip

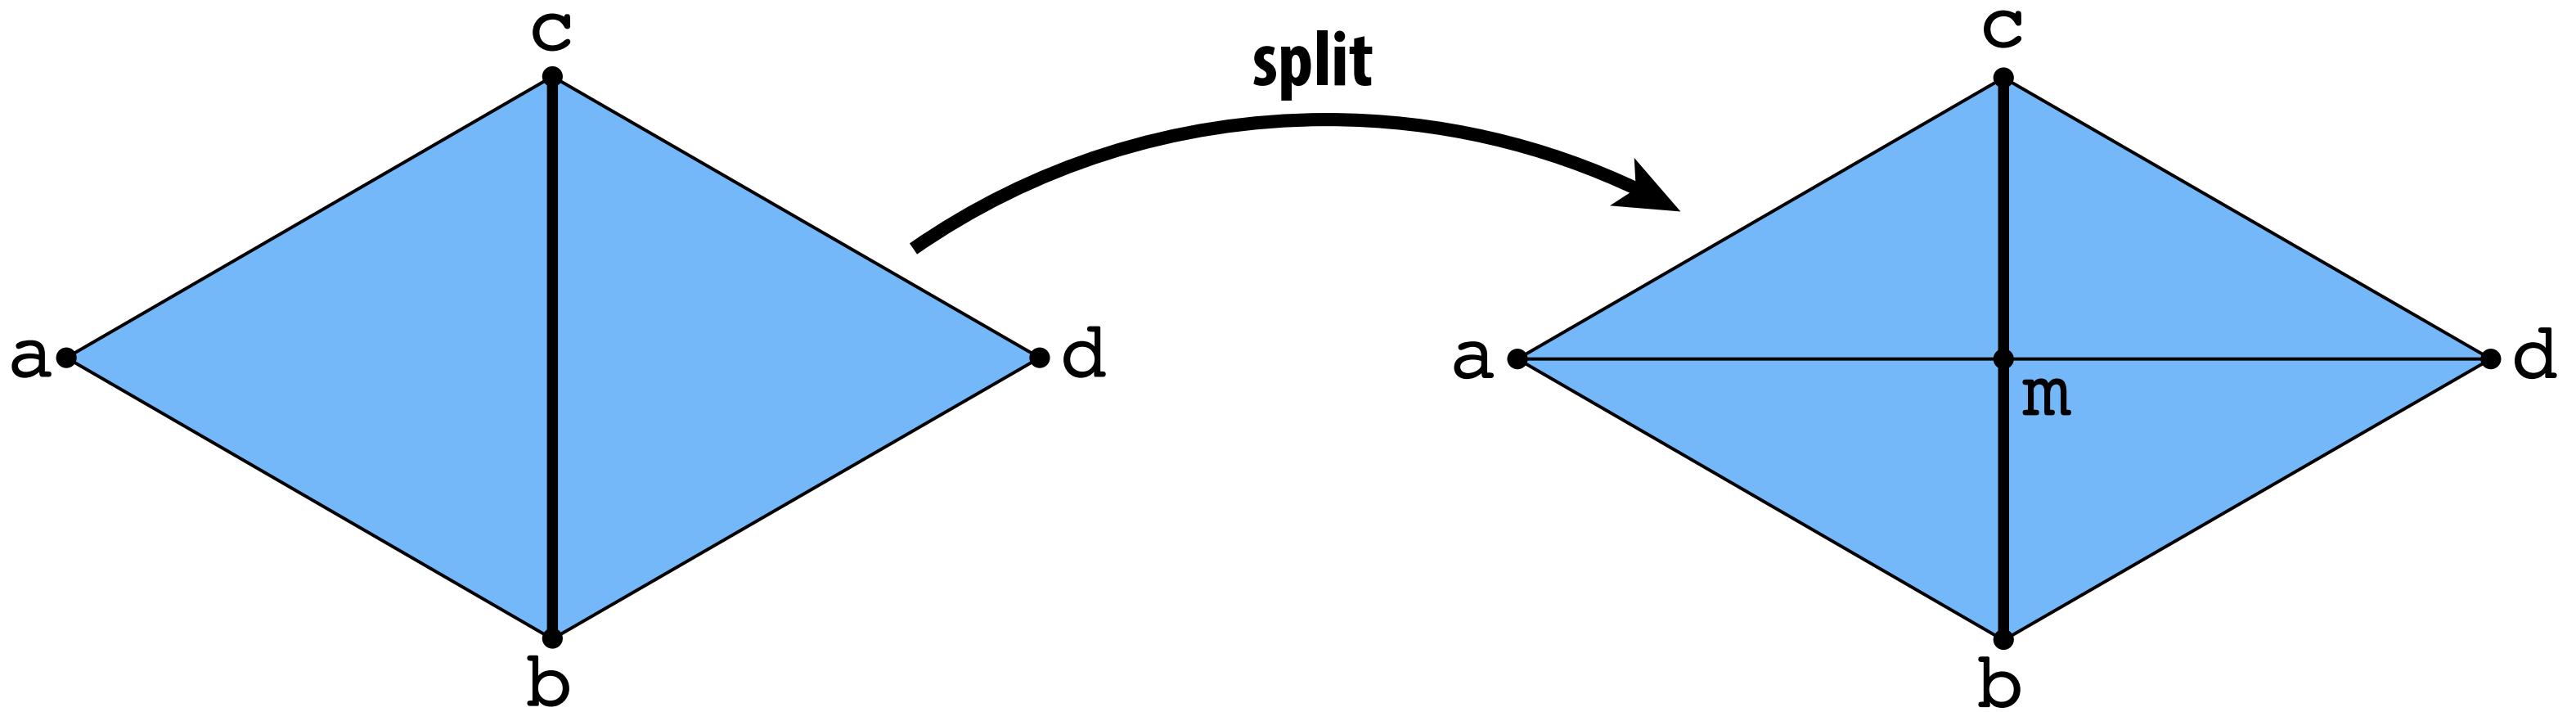
- Triangles $(a,b,c), (b,d,c)$ become $(a,d,c), (a,b,d)$:



- Long list of pointer reassessments (`edge->halfedge = ...`)
- However, no elements created/destroyed.
- Q: What happens if we flip twice?
- (Challenge: can you implement edge flip such that pointers are unchanged after two flips?)

Edge Split

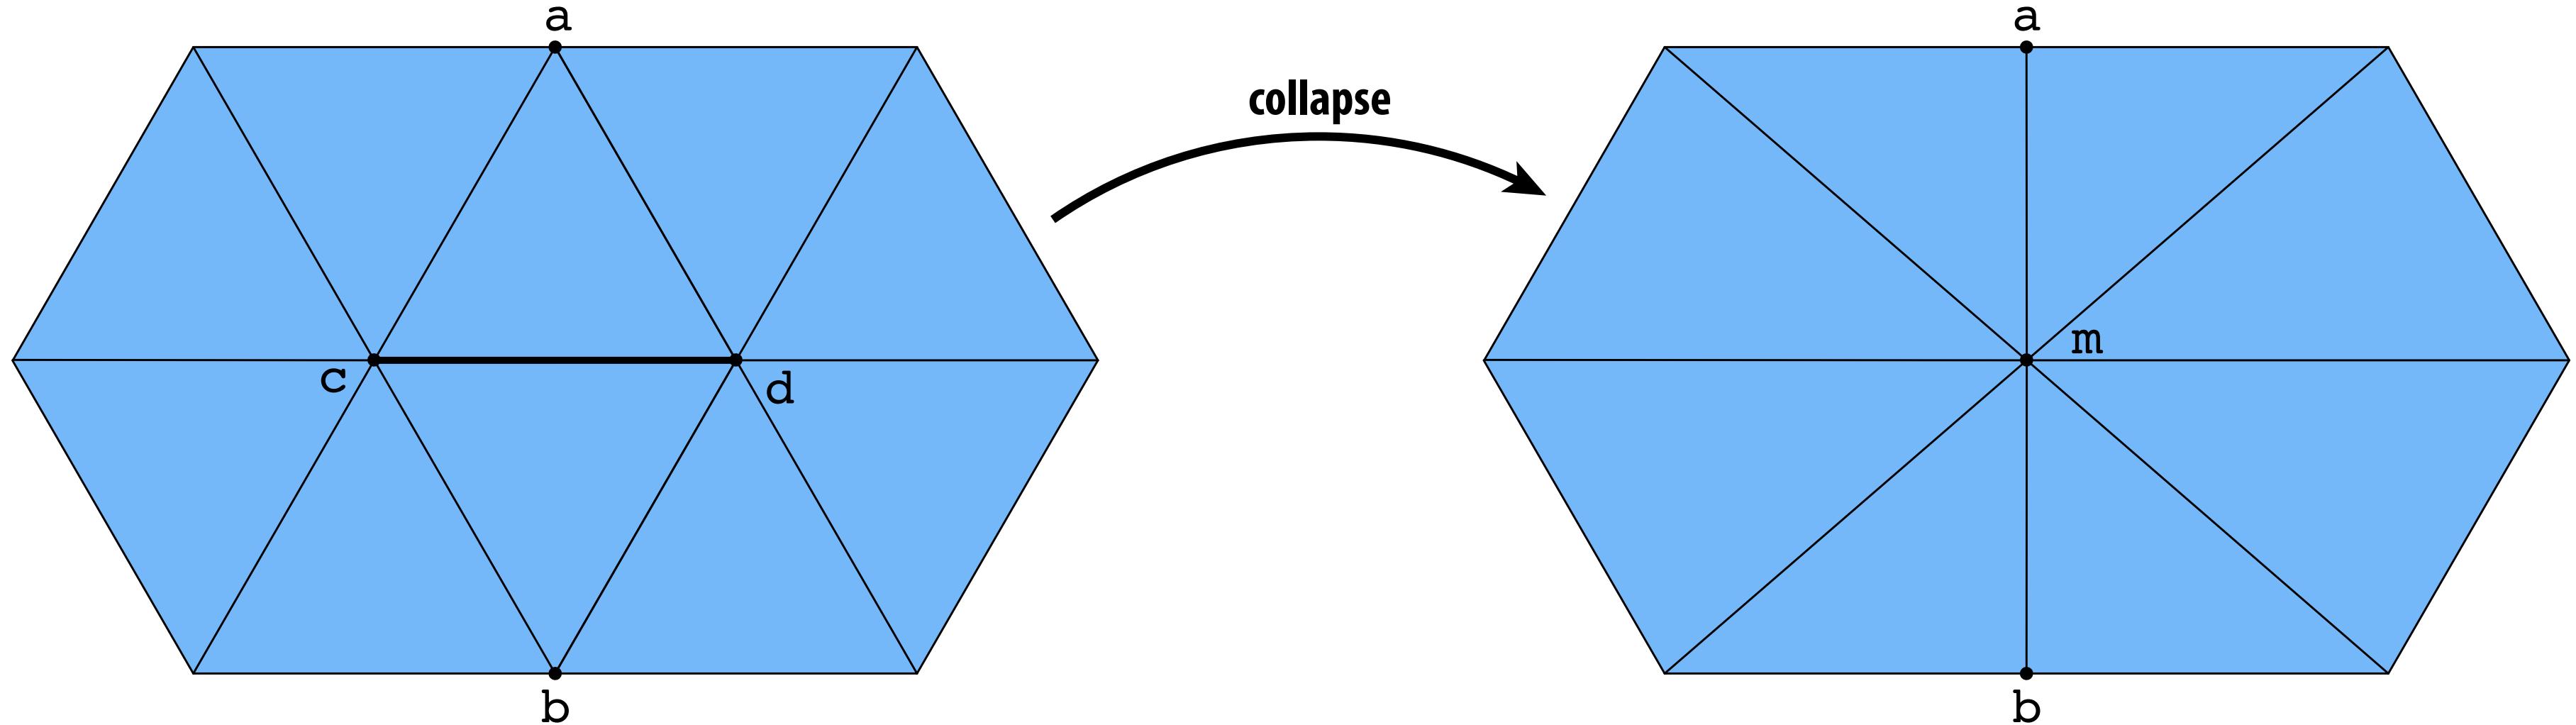
- Insert midpoint m of edge (c,b) , connect to get four triangles:



- This time, have to *add* new elements.
- Lots of pointer reassessments.
- Q: Can we “reverse” this operation?

Edge Collapse

- Replace edge (c,d) with a single vertex m:

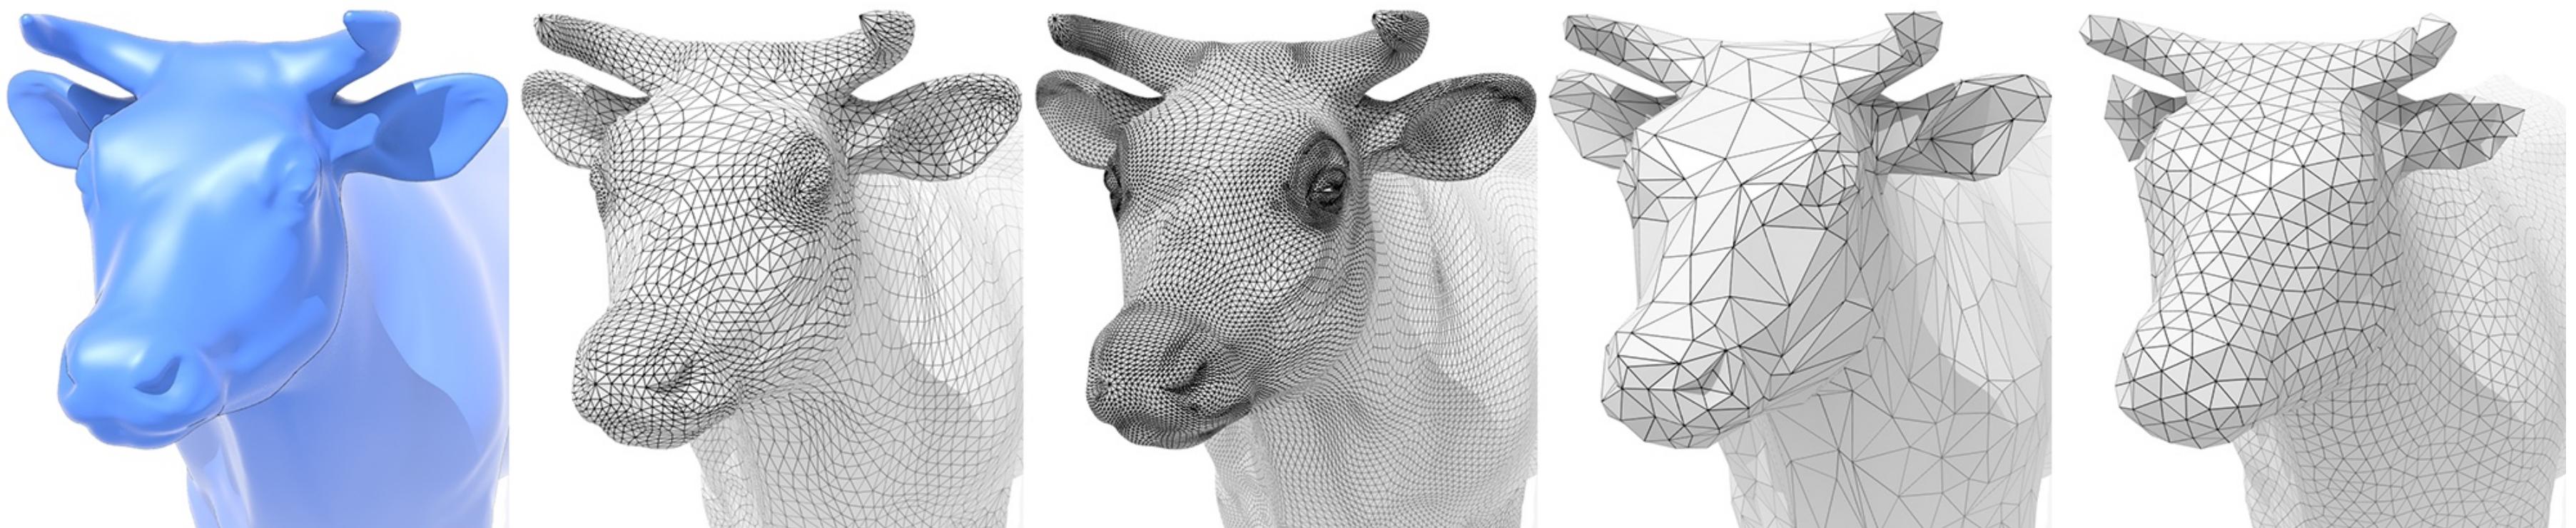
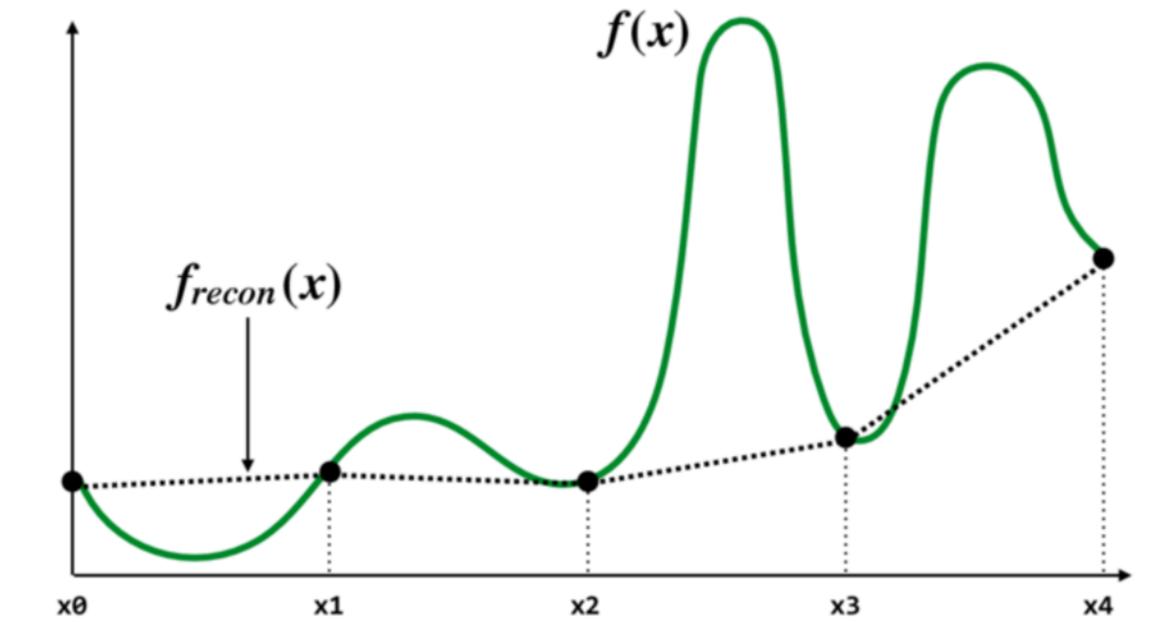


- Now have to *delete* elements.
- Still lots of pointer assignments!
- Q: How would we implement this with a polygon soup?
- Any other good way to do it? (E.g., different data structure?)

**Ok, but what can we actually *do*
with these operations?**

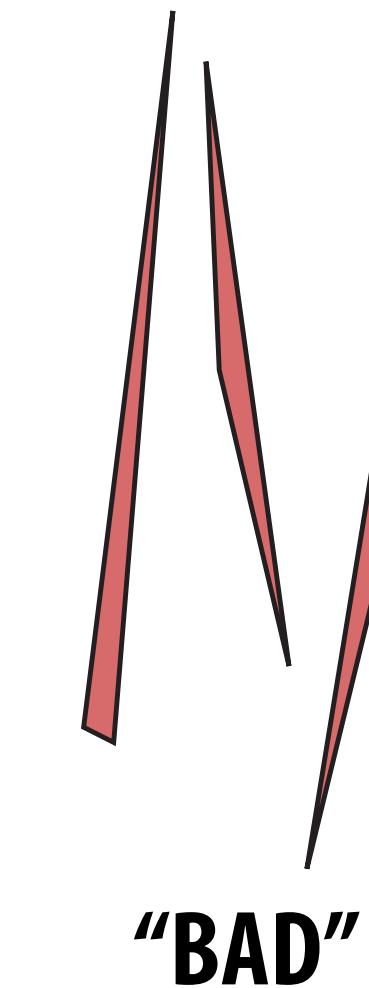
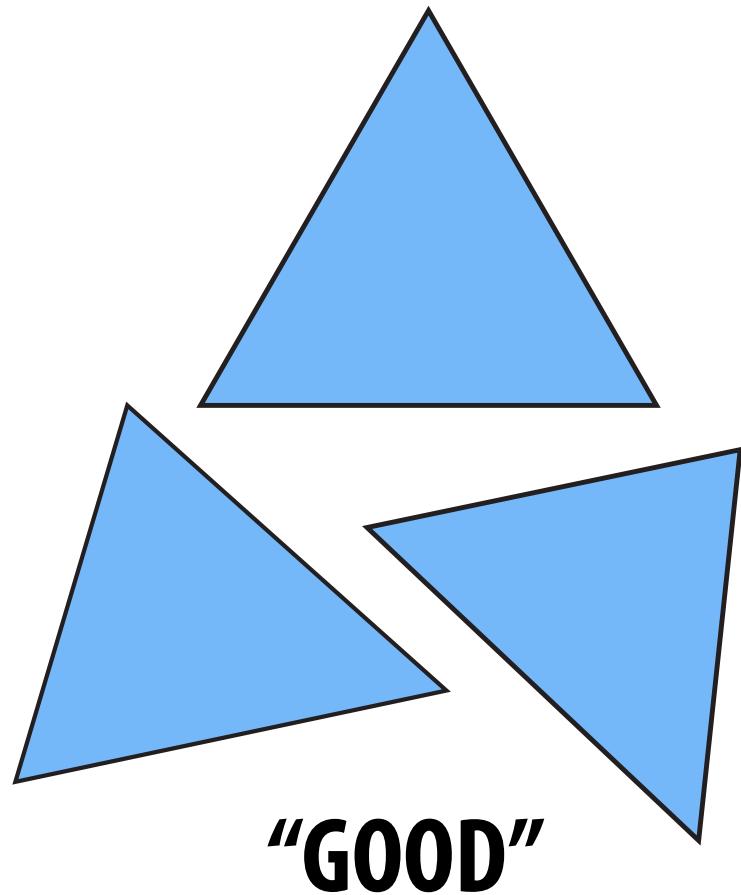
Remeshing as resampling

- Remember our discussion of *aliasing*
- Bad sampling makes signal appear different than it really is
- E.g., undersampled curve looks flat
- Geometry is no different!
 - undersampling destroys features
 - oversampling destroys performance

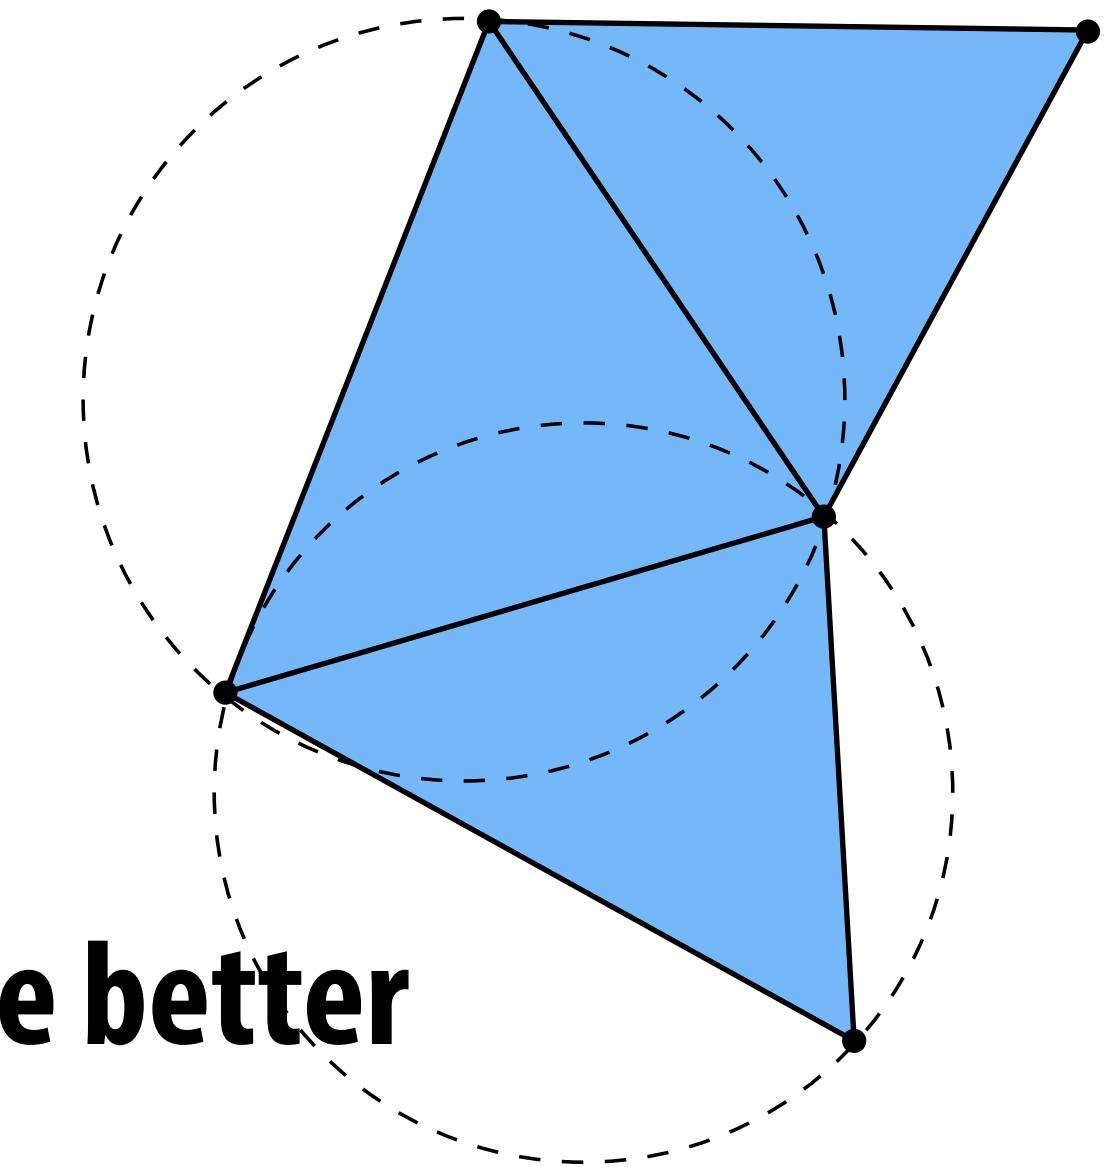


What makes a “good” triangle mesh?

- One rule of thumb: *triangle shape*



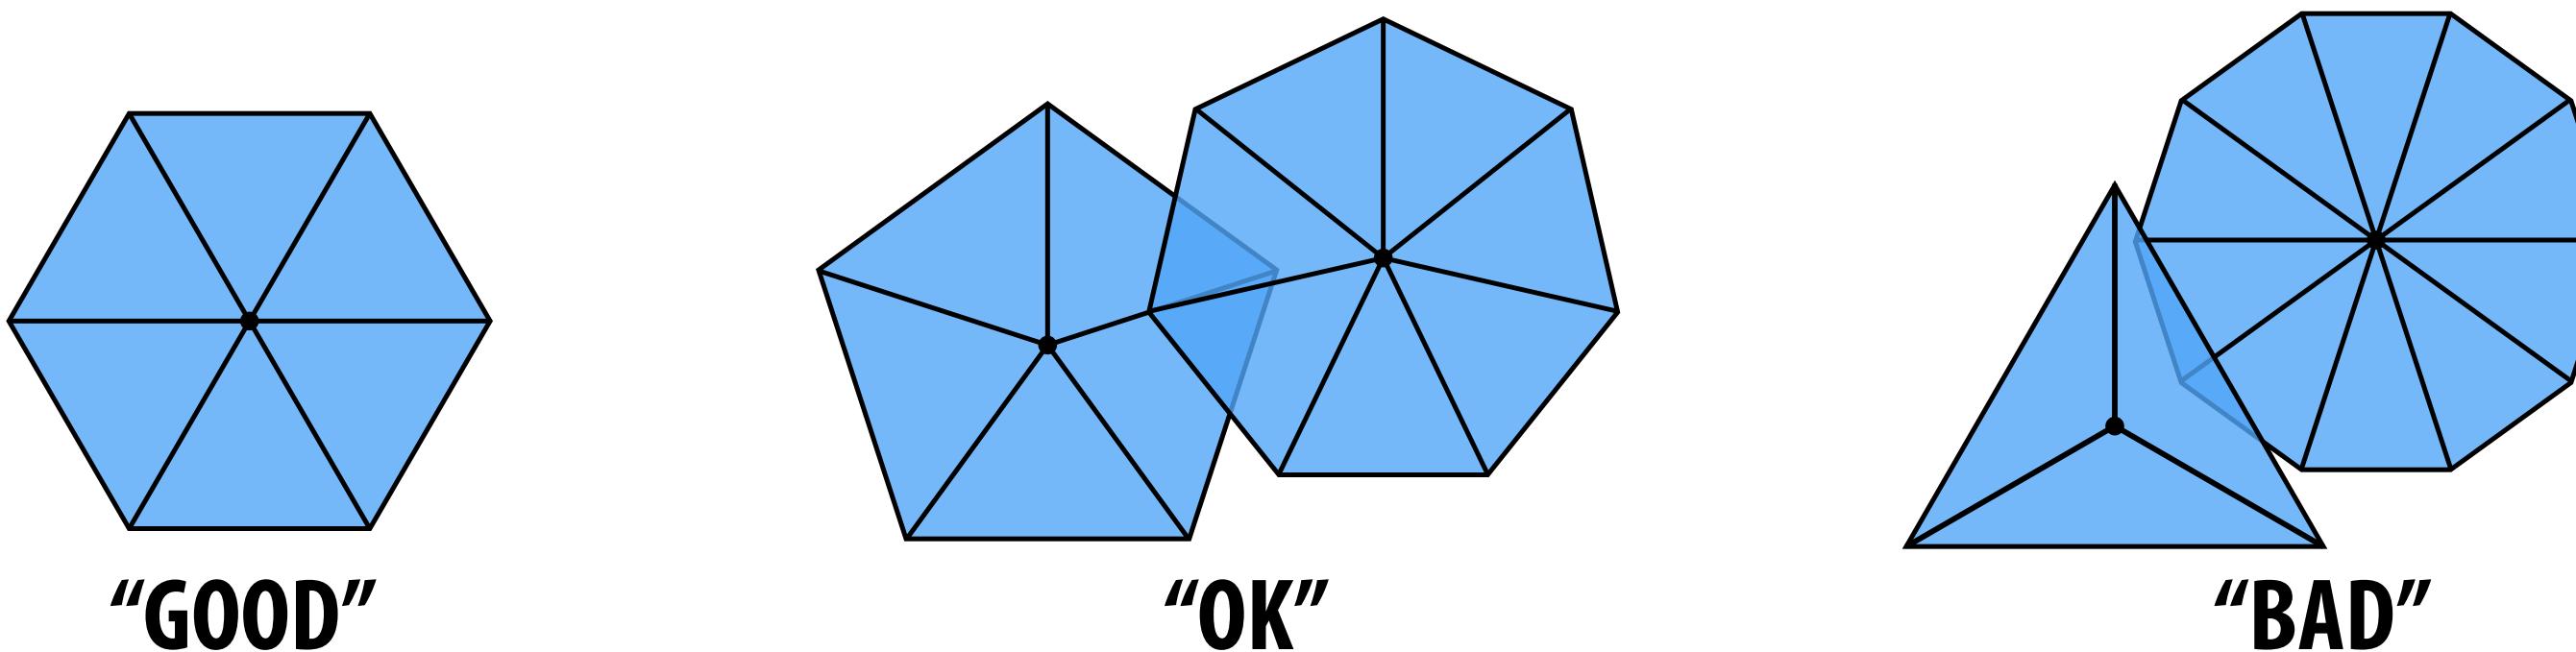
- More specific condition: *Delaunay*
- “Circumcircle interiors contain no vertices.”
- Not *always* a good condition, but often*
 - especially important for simulation
 - for approximation, long triangles may be better



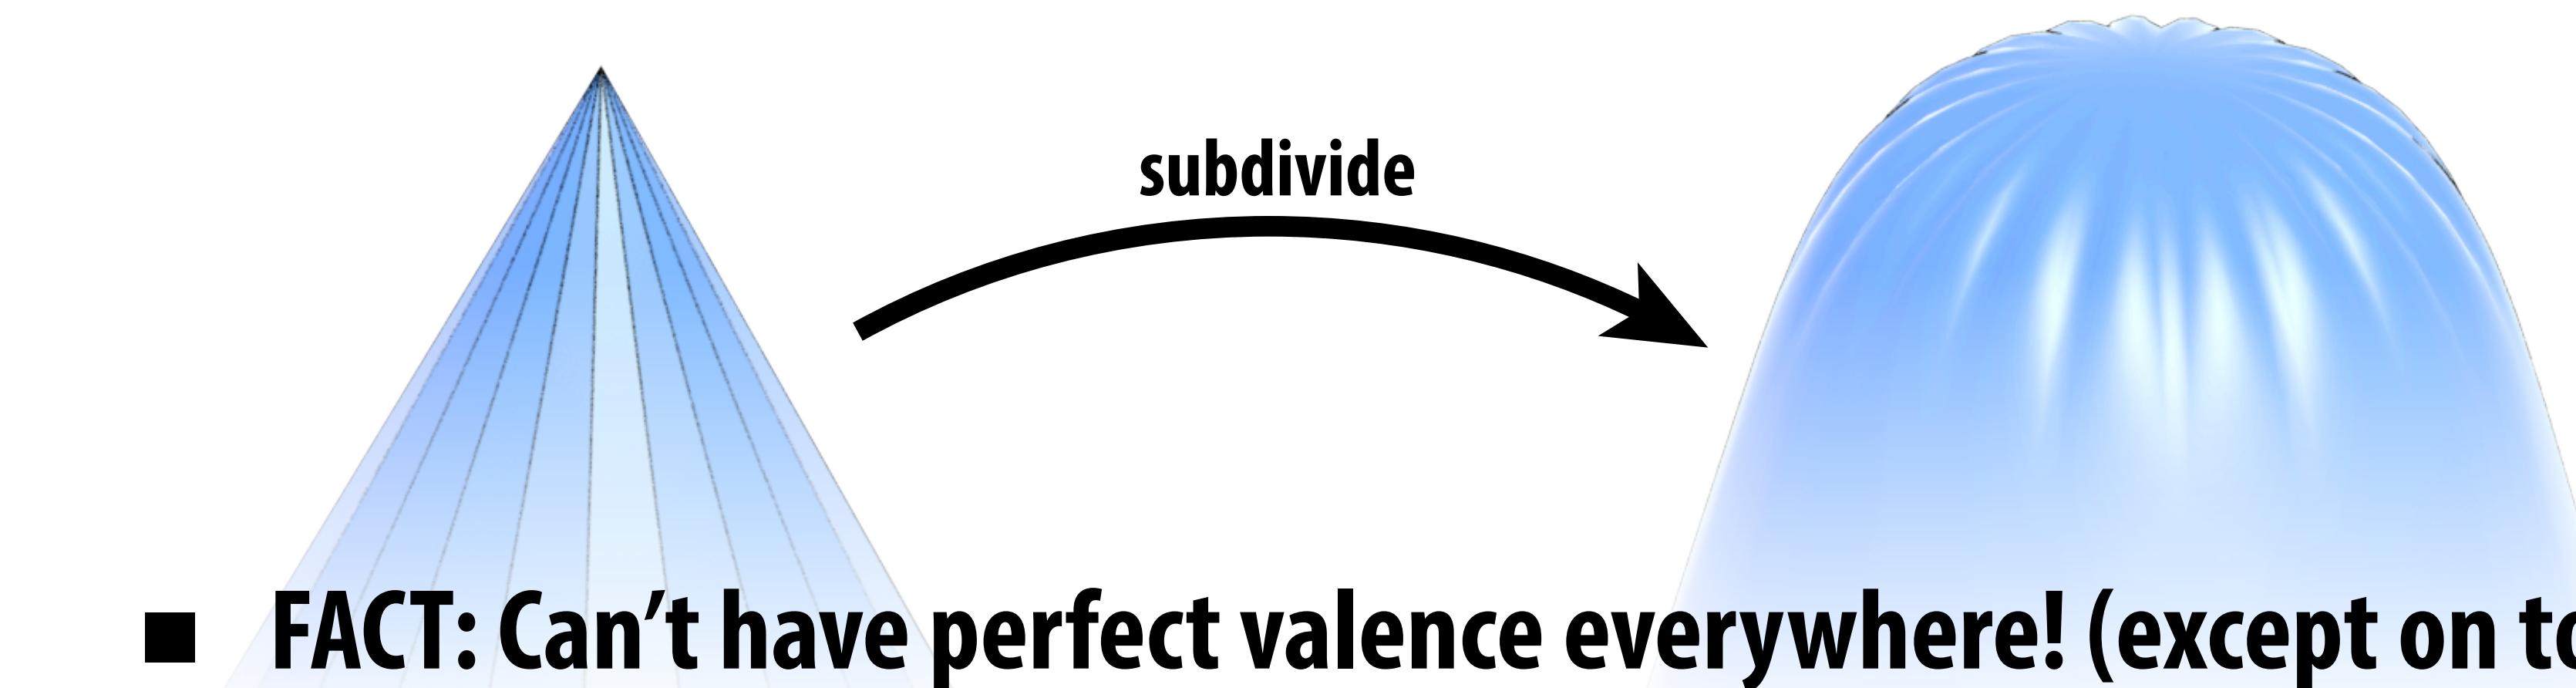
*See Shewchuk, “What is a Good Linear Element”

What else constitutes a good mesh?

- Another rule of thumb: *regular vertex degree*
- Ideal for triangle meshes: make every vertex valence 6:



- Why? Better triangle shape, important for (e.g.) subdivision:

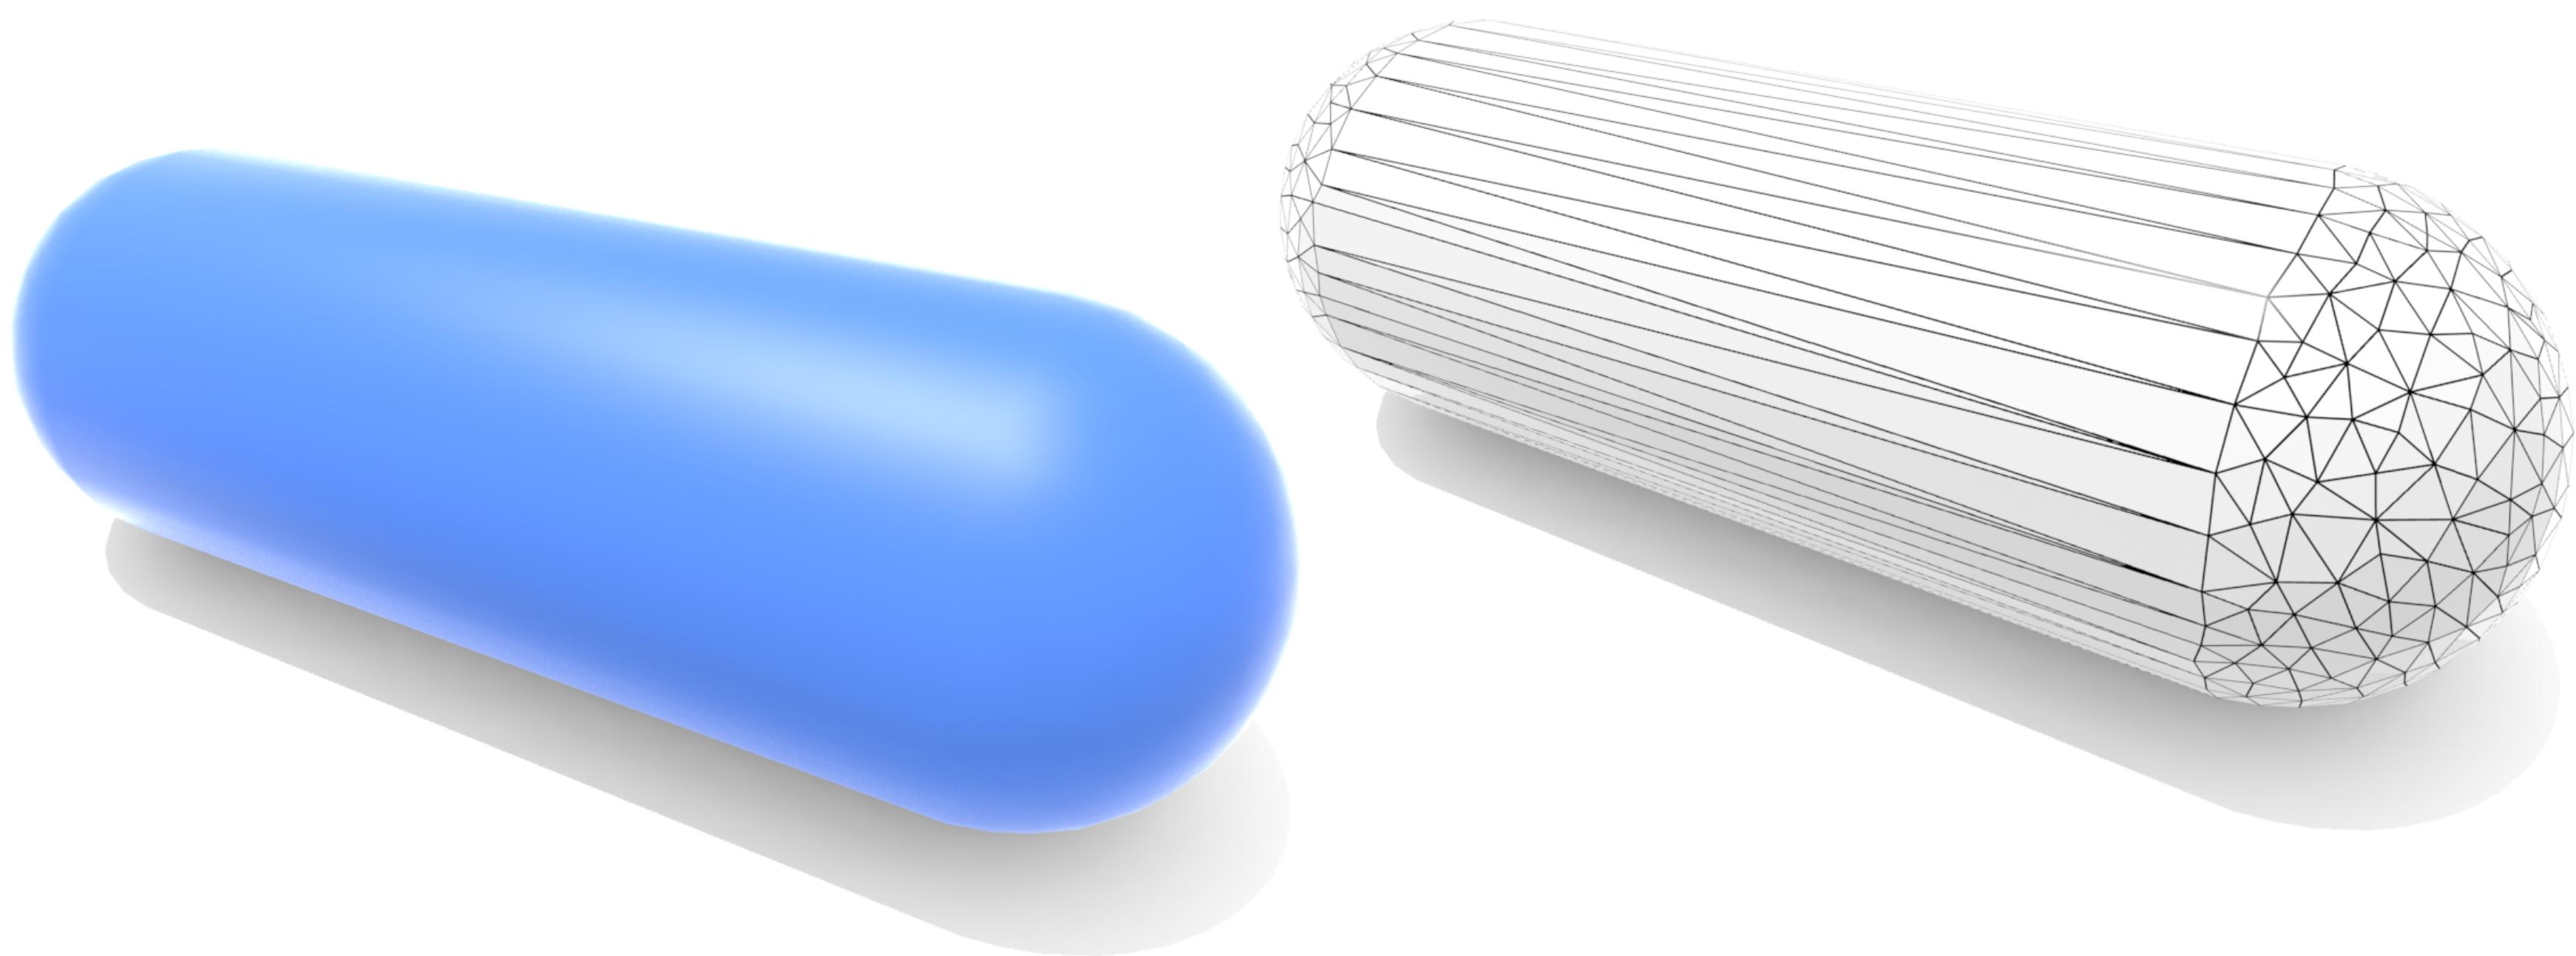


- FACT: Can't have perfect valence everywhere! (except on torus)

*See Shewchuk, "What is a Good Linear Element"

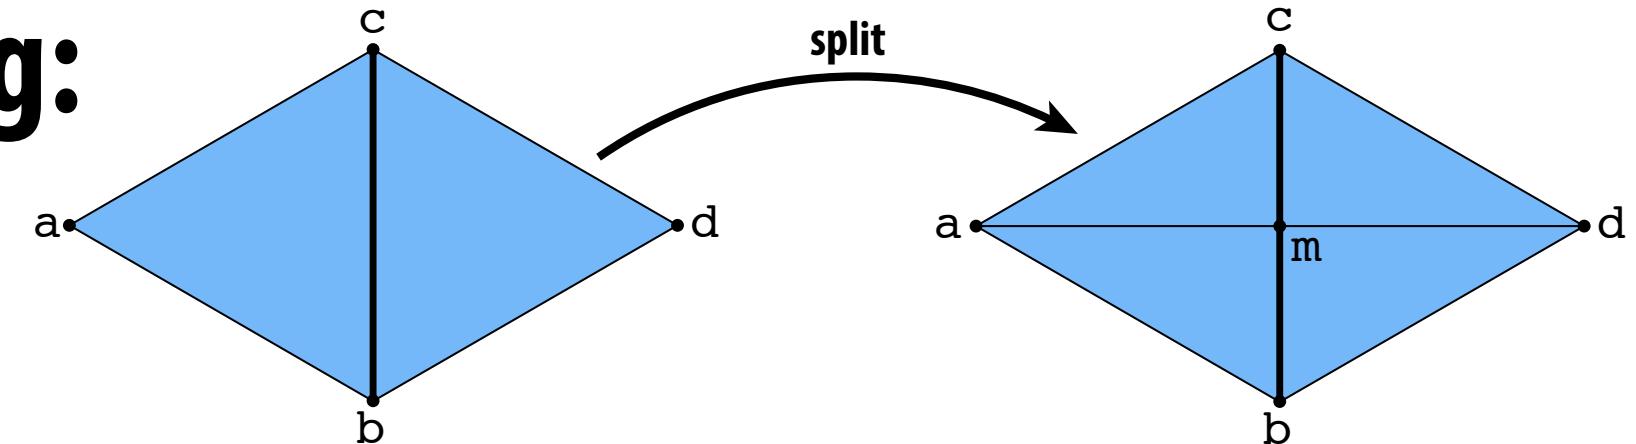
What else makes a “good” mesh?

- Good approximation of original shape!
- Keep only elements that contribute *information* about shape
- Add additional information where, e.g., *curvature* is large
- Balance with element quality
 - Delaunay, “round” triangles, regular degree, etc.

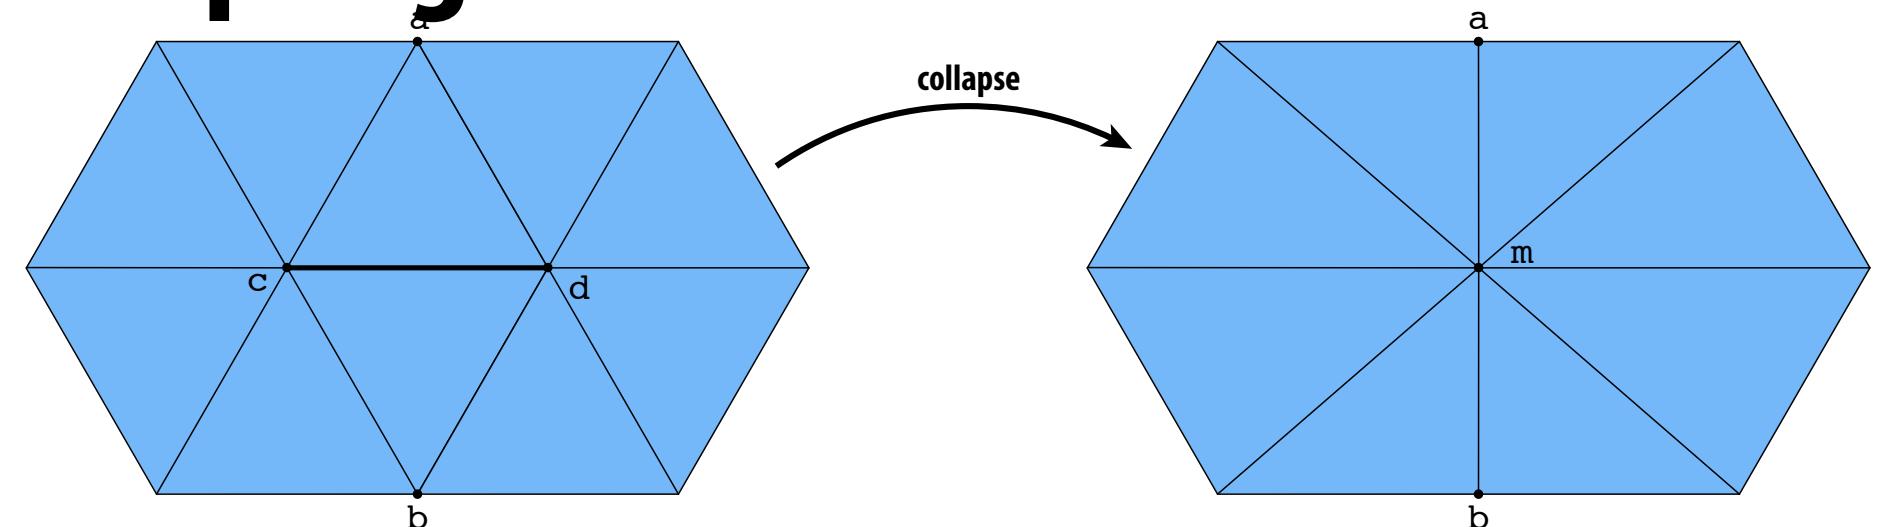


How do we resample? Already know how!

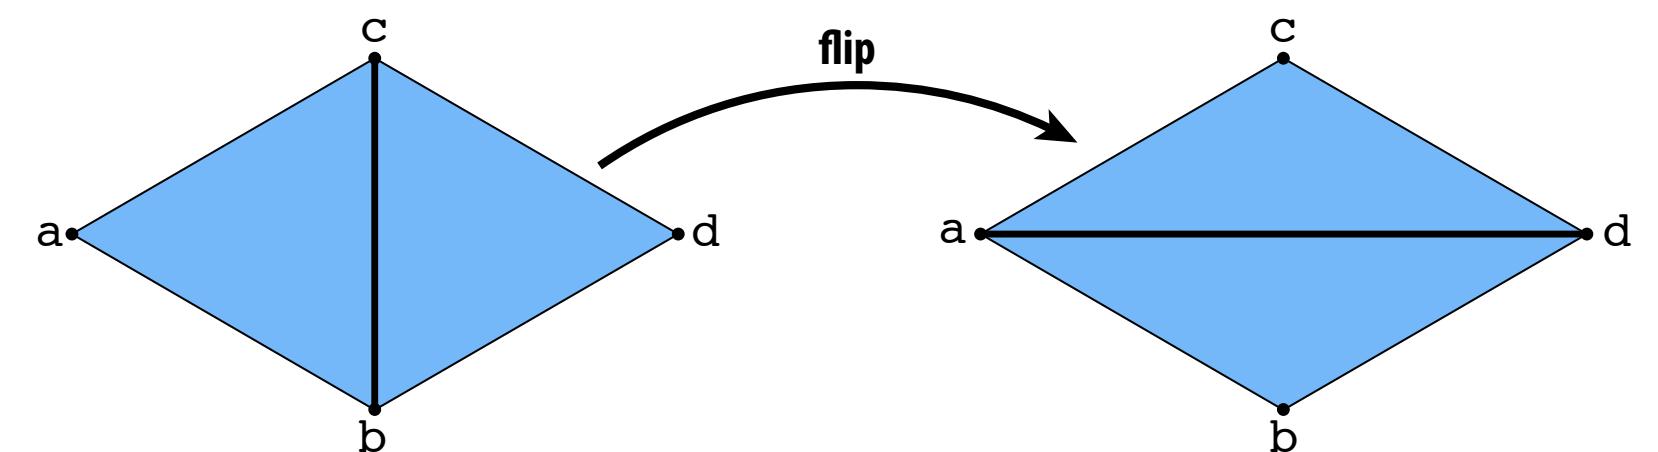
- Edge split is (local) upsampling:



- Edge collapse is (local) downsampling:



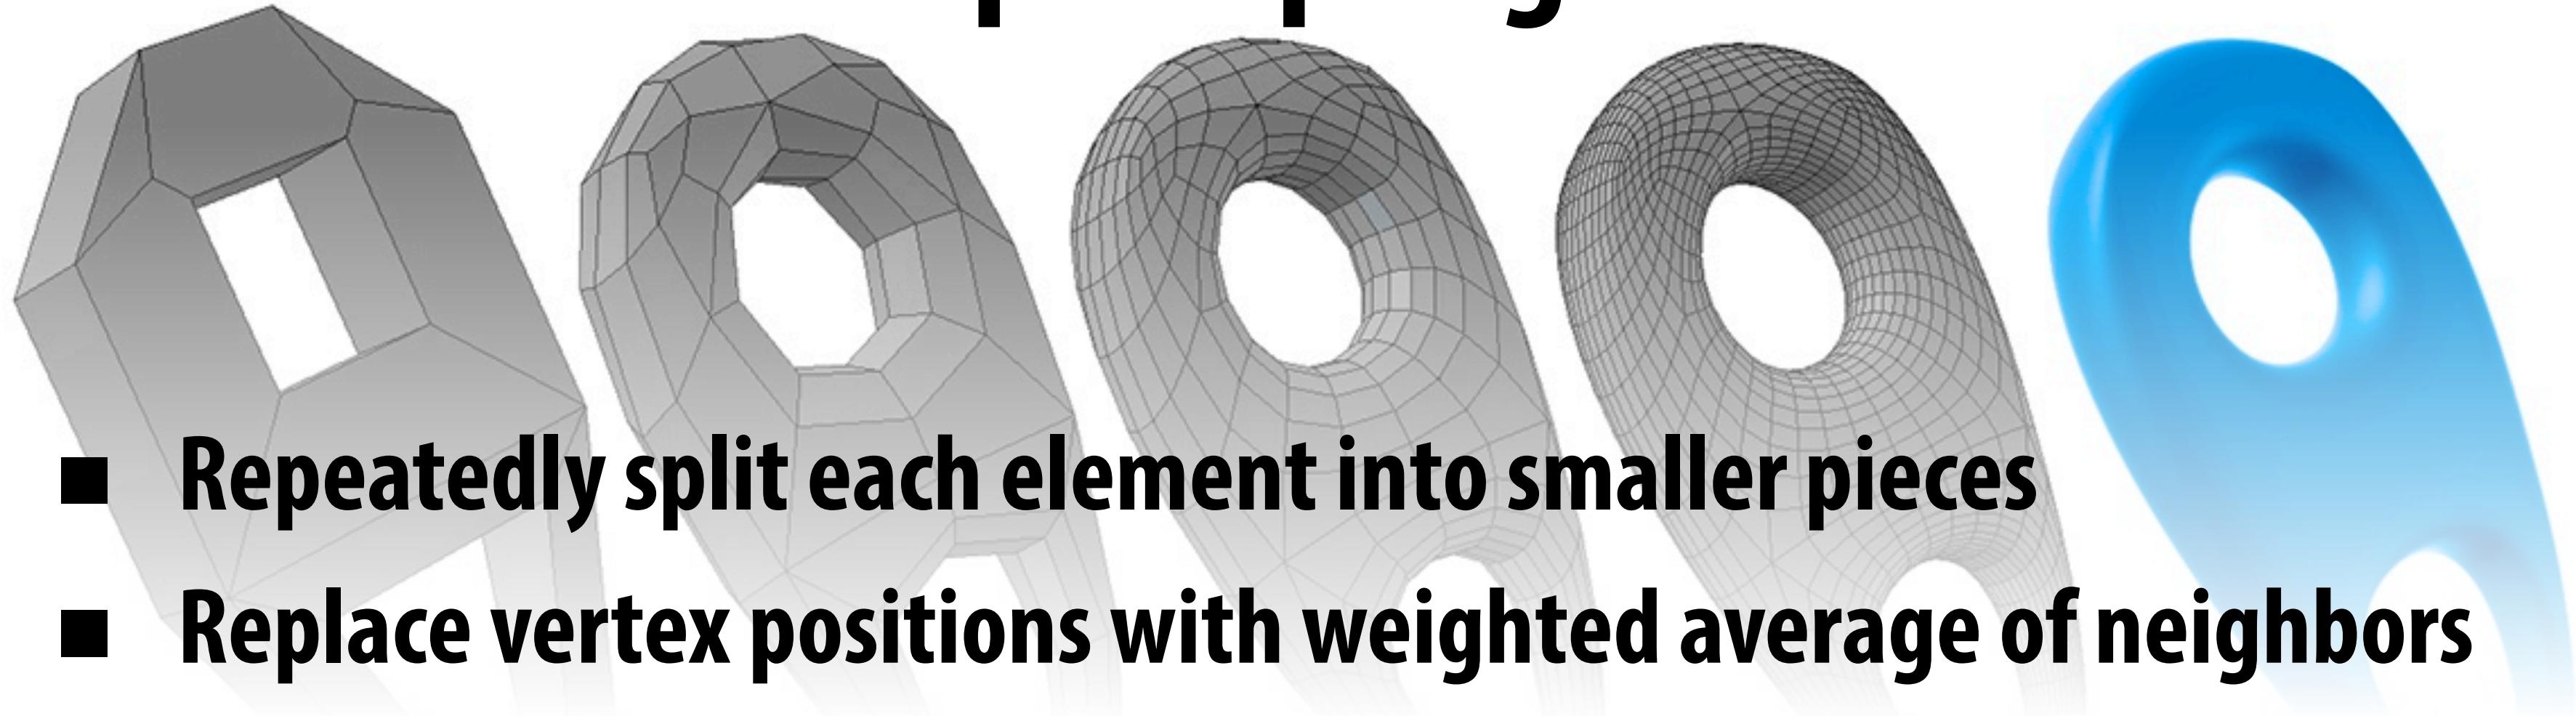
- Edge flip is (local) resampling:



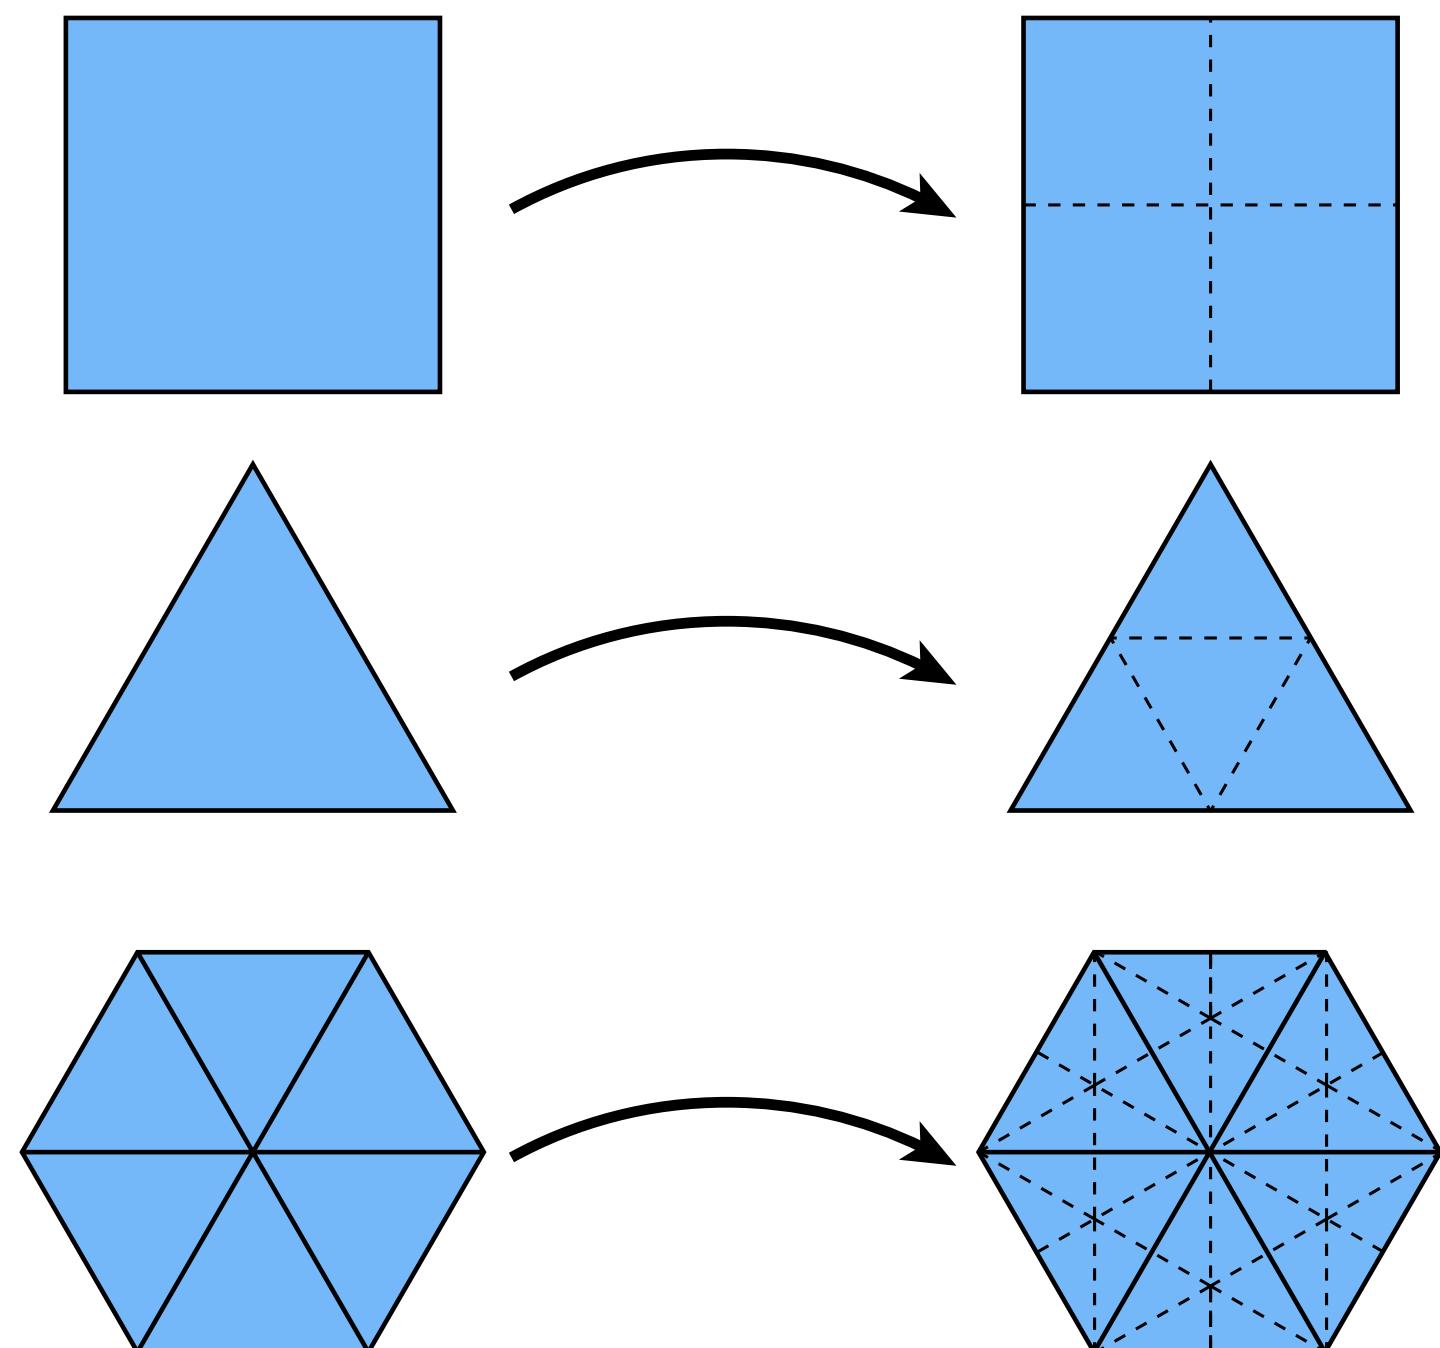
- Still need to intelligently decide *which* edges to modify!

***Which edges should we split to
upsample the whole mesh?***

Subdivision as Upsampling

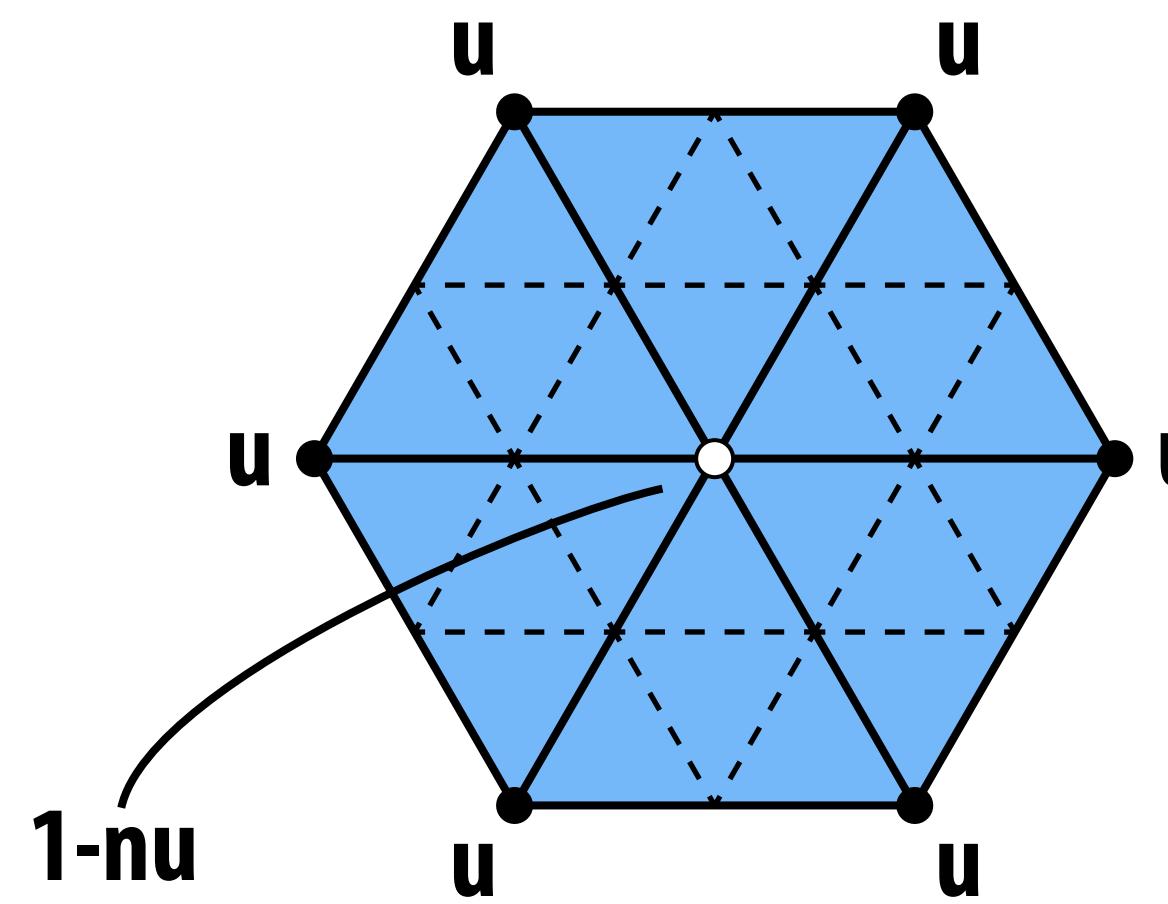
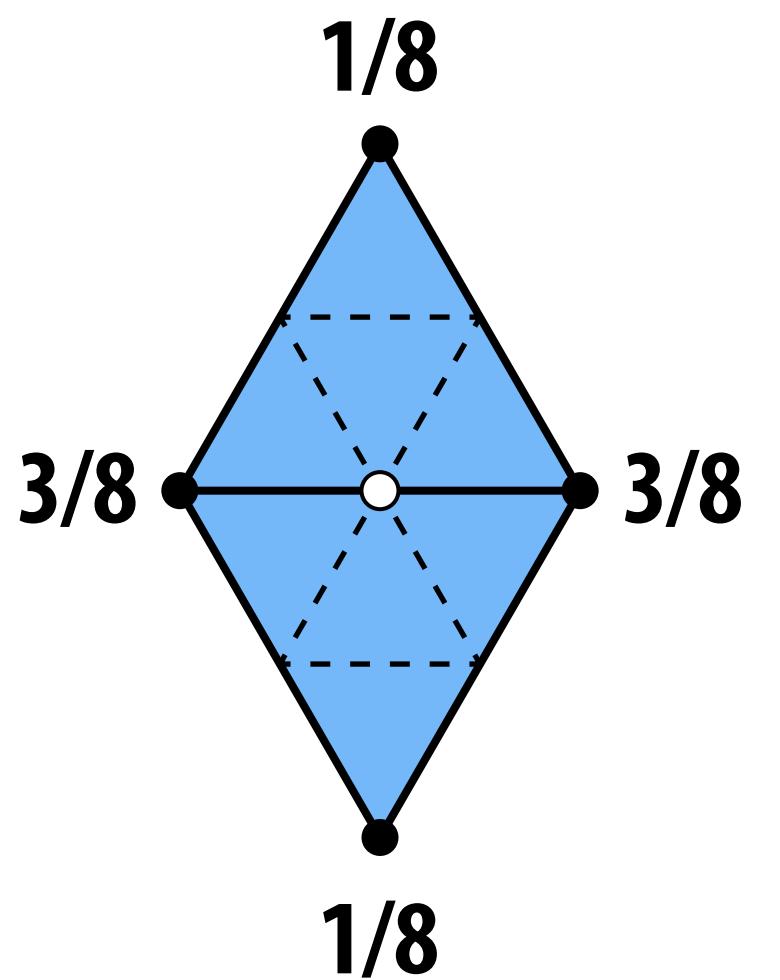
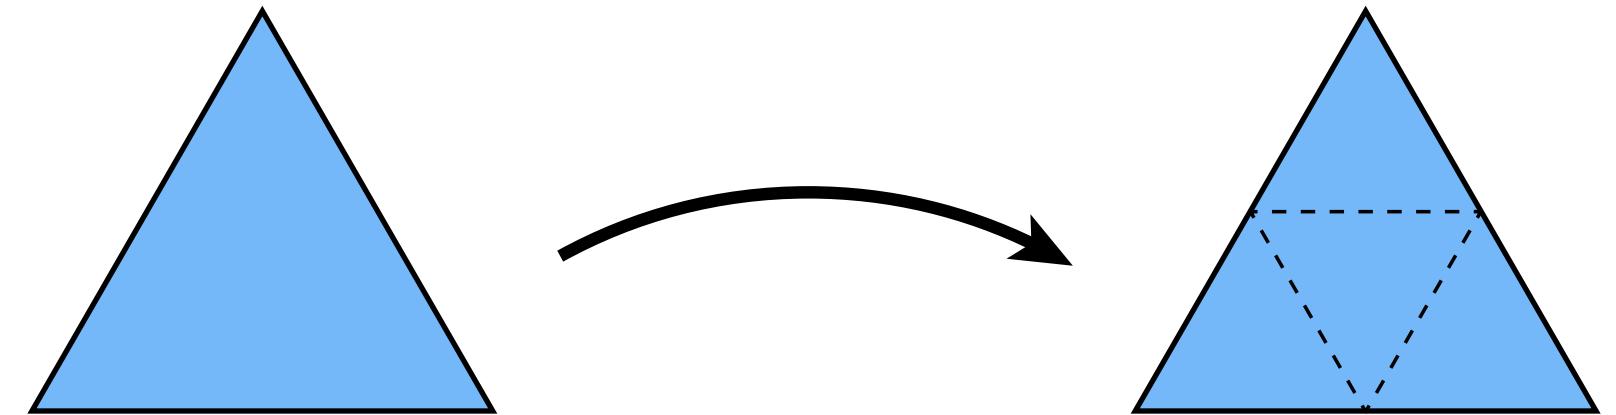


- Repeatedly split each element into smaller pieces
- Replace vertex positions with weighted average of neighbors
- Main considerations:
 - interpolating vs. approximating
 - limit surface continuity (C^1, C^2, \dots)
 - behavior at irregular vertices
- Many options:
 - Quad: Catmull-Clark
 - Triangle: Loop, Butterfly, $\text{Sqrt}(3)$



Loop Subdivision

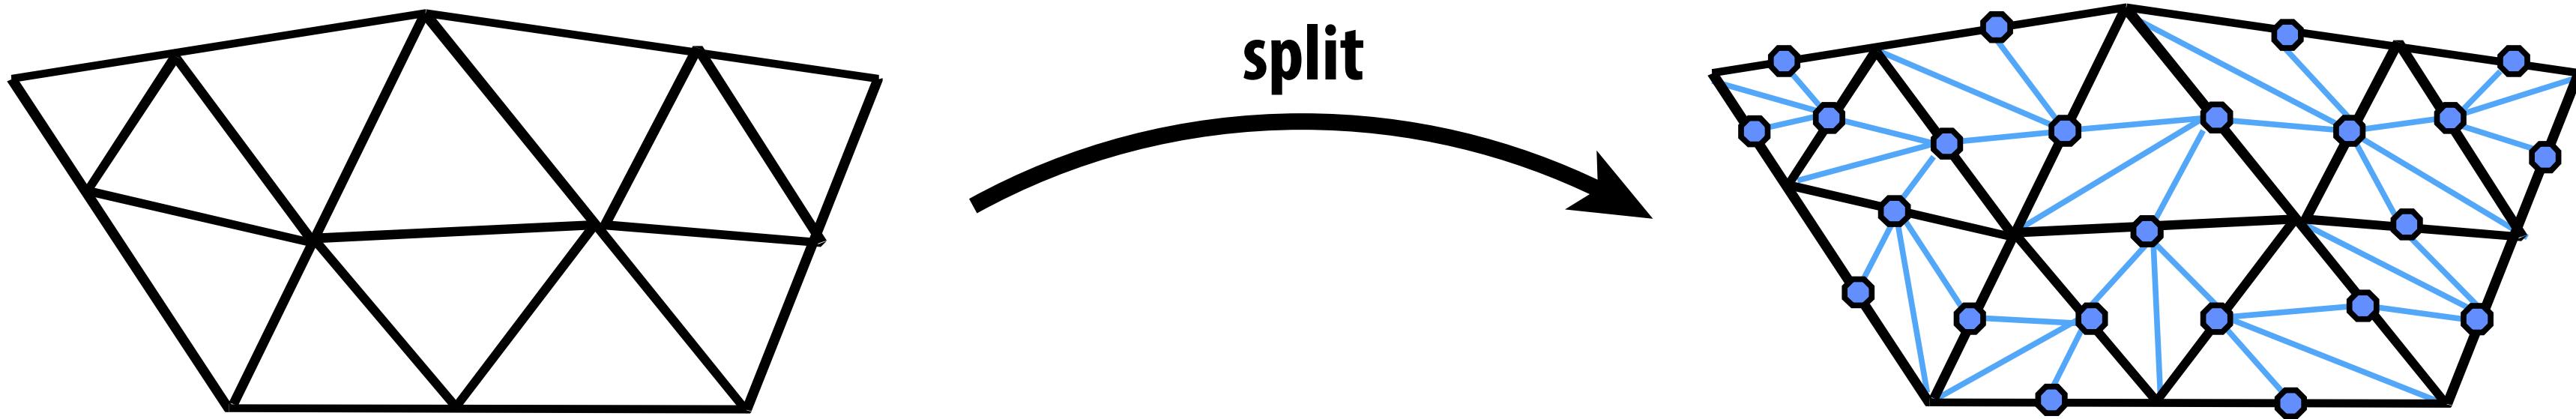
- Fairly common subdivision rule for triangles
- Curvature is continuous away from irregular vertices ("C²")
- *Approximating*, not interpolating
- Algorithm:
 - Split each triangle into four
 - Assign new vertex positions according to weights:



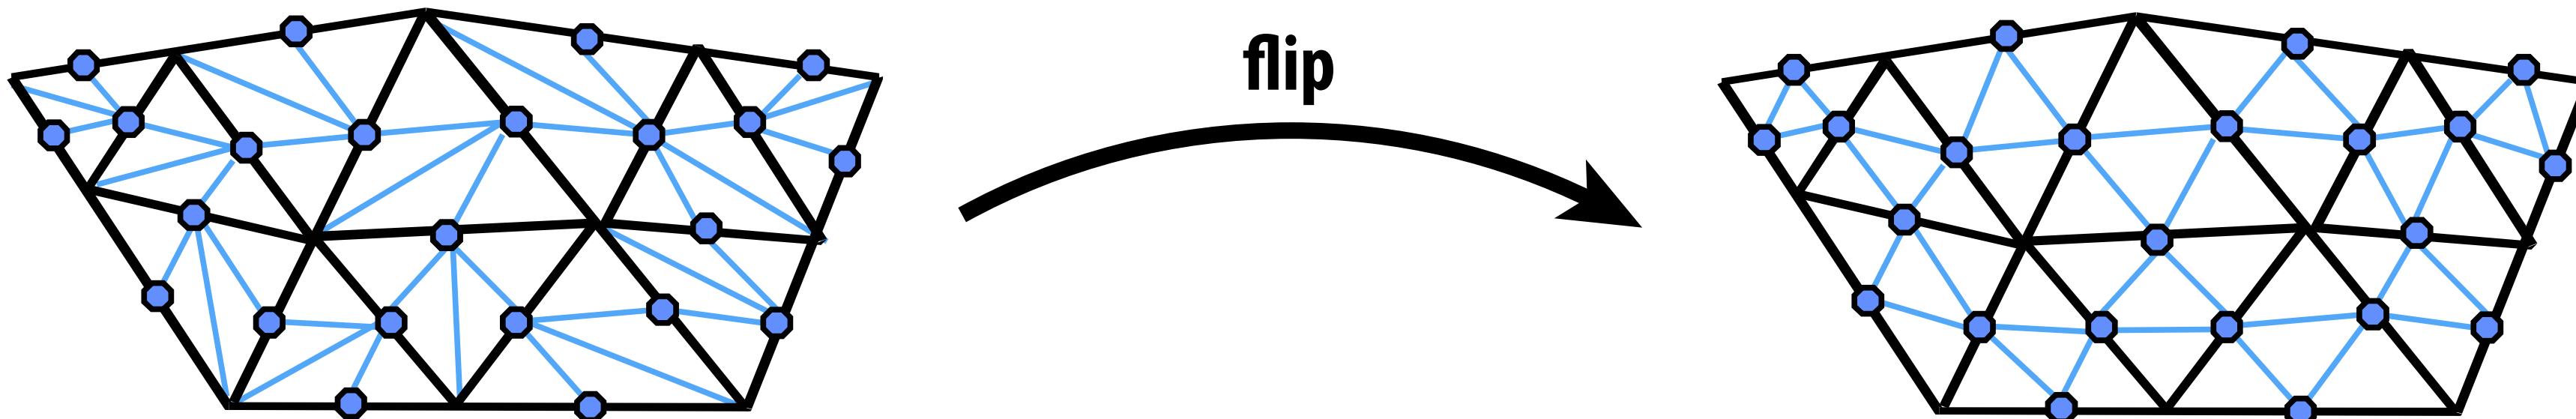
n : vertex degree
 u : $3/16$ if $n=3$, $3/(8n)$ otherwise

Loop Subdivision via Edge Operations

- First, split edges of original mesh in *any* order:



- Next, flip new edges that touch a new & old vertex:

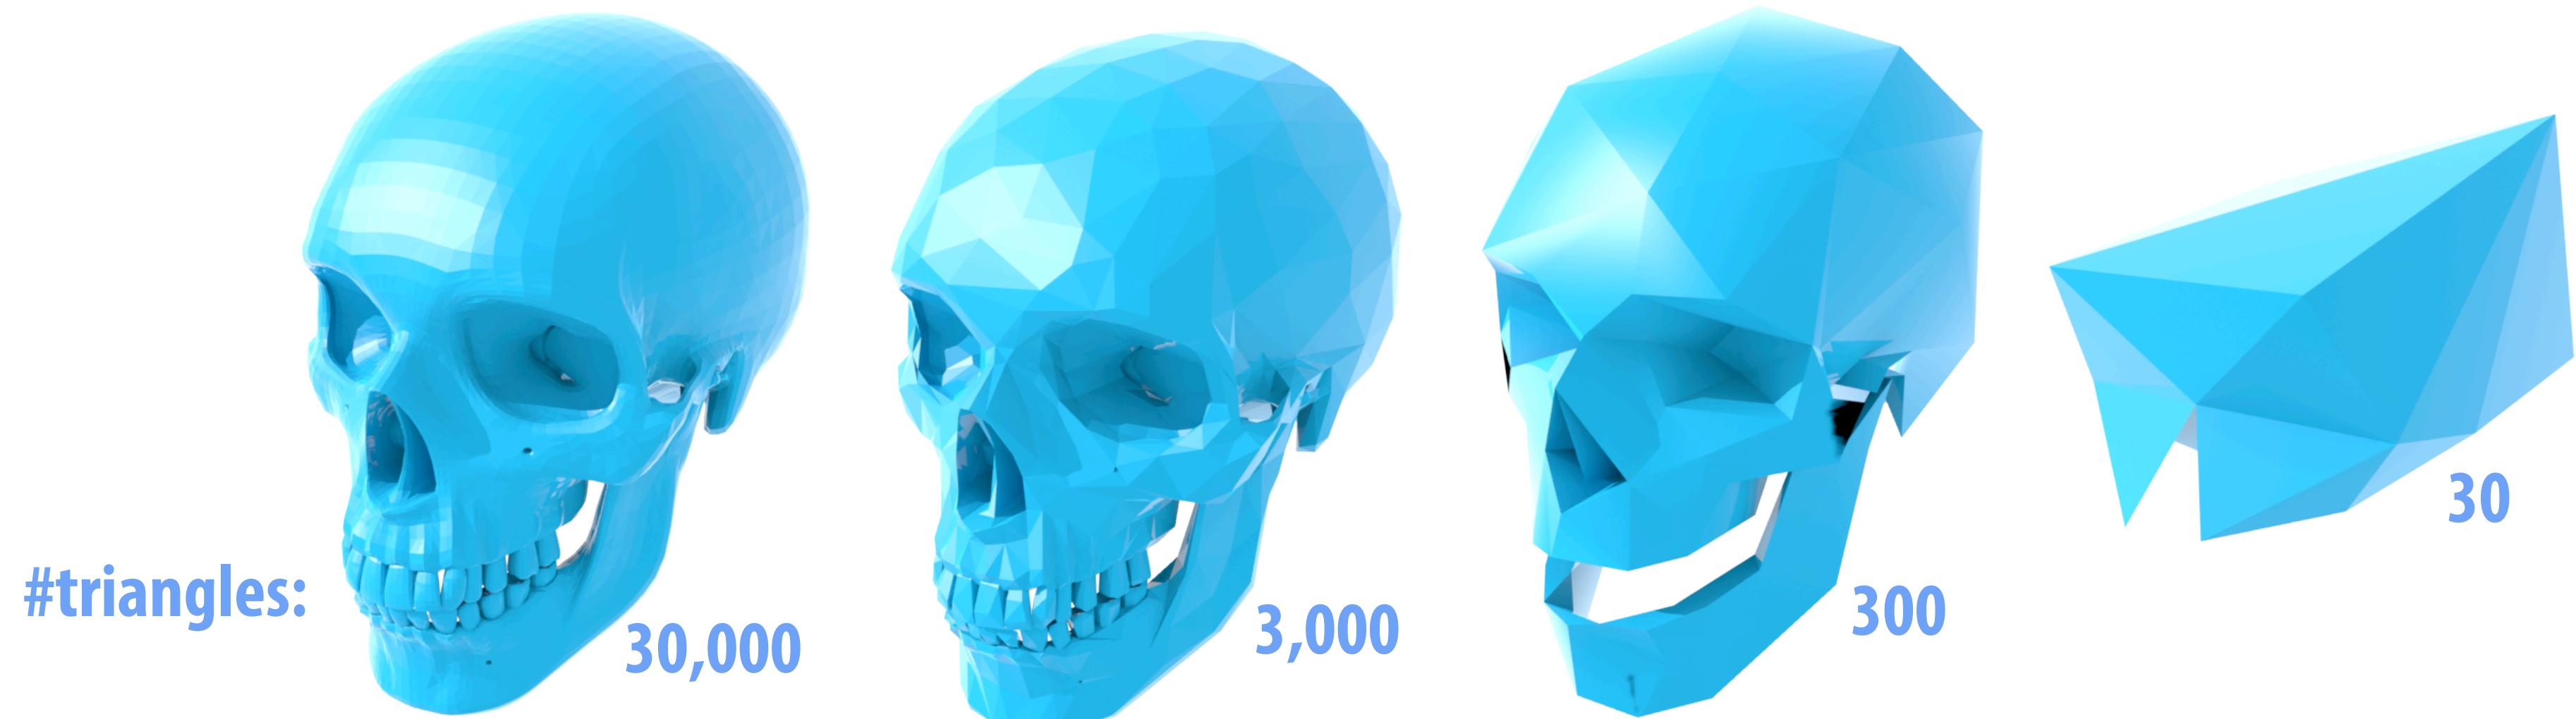


(Don't forget to update vertex positions!)

What if we want *fewer* triangles?

Simplification via Quadric Error Metric

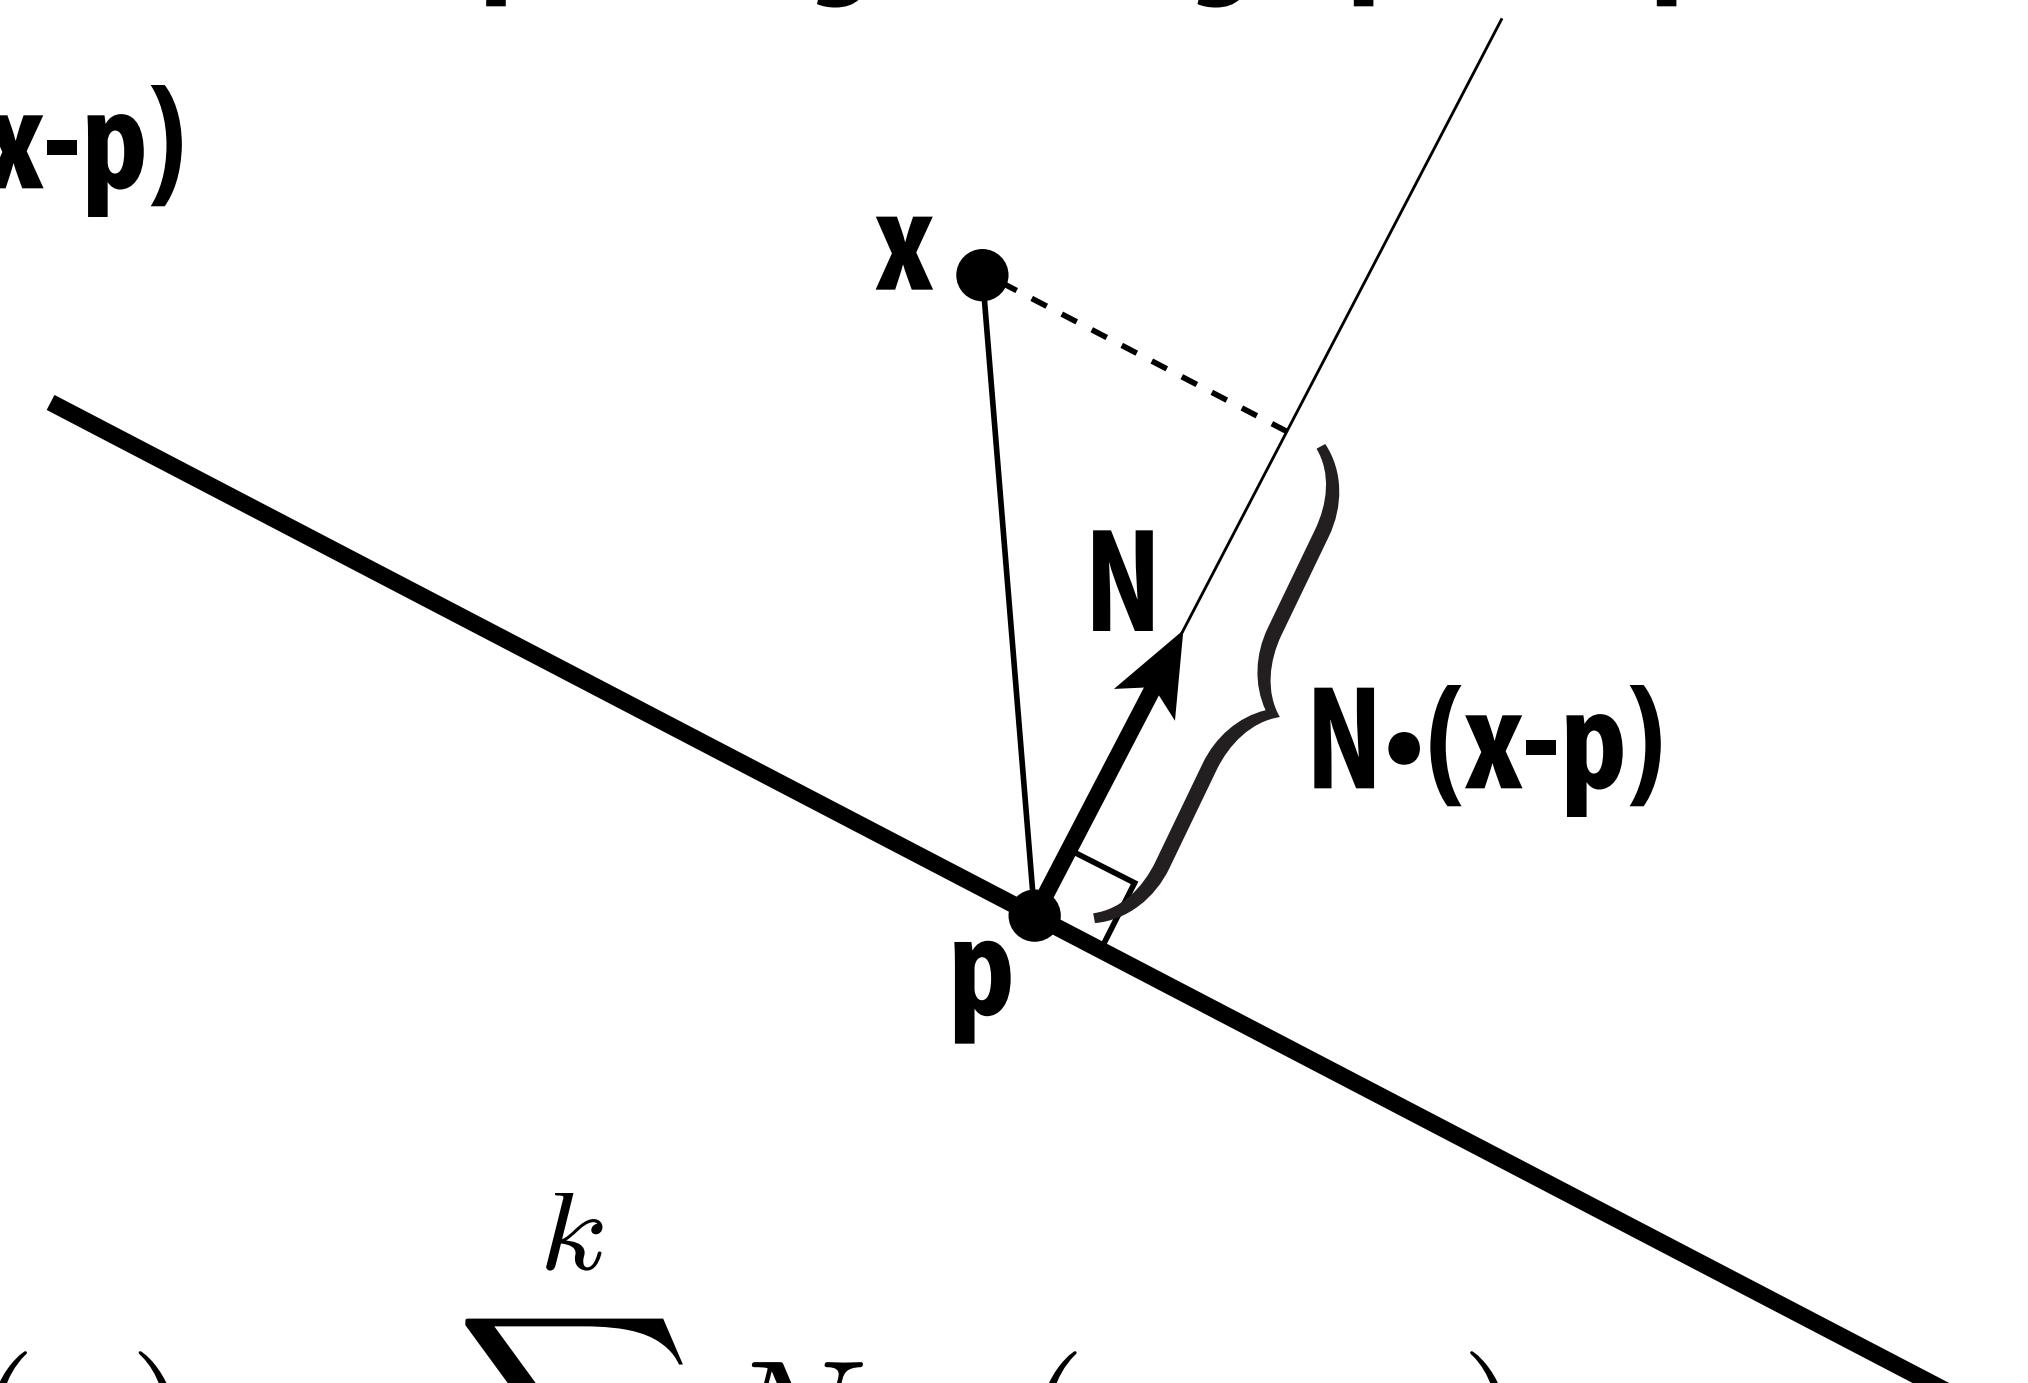
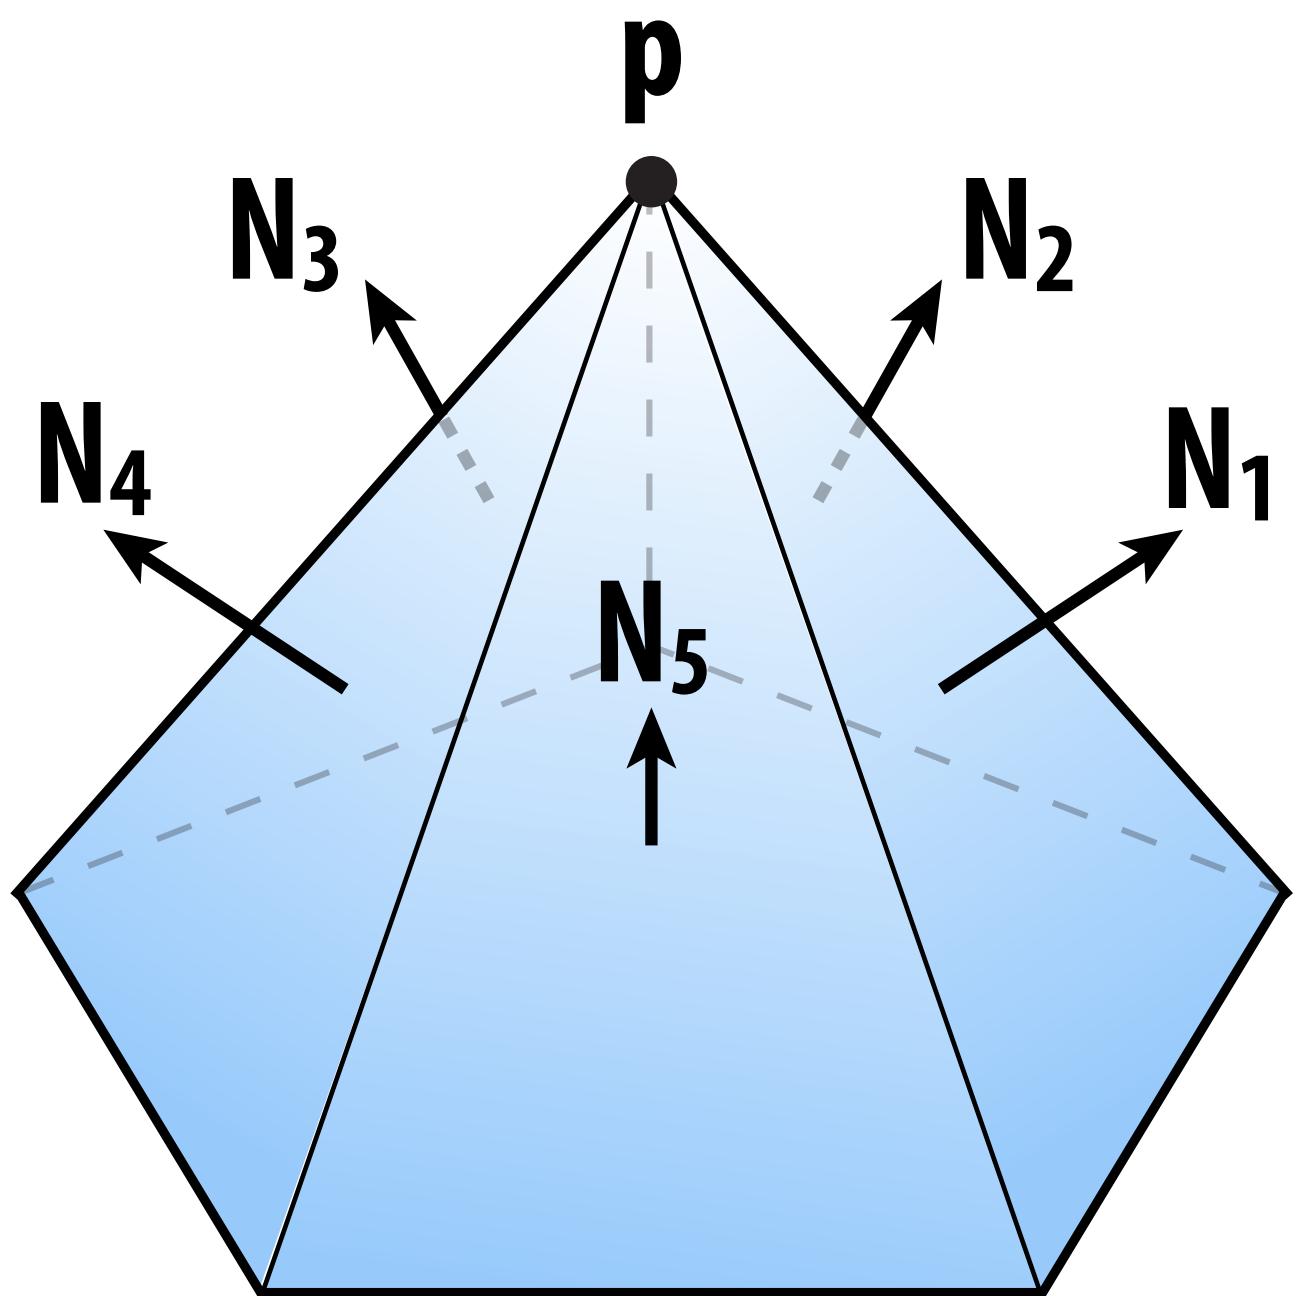
- One popular scheme: iteratively collapse edges
- Which edges? Assign score with *quadric error metric**
 - approximate distance to surface as sum of distance to aggregated triangles
 - iteratively collapse edge with smallest score
 - greedy algorithm... great results!



*invented here at CMU! (Garland & Heckbert 1997)

Quadric Error Metric

- Approximate distance to a collection of triangles
- Distance is sum of point-to-plane distances
 - Q: Distance to plane w/ normal N passing through point p ?
 - A: $d(x) = N \cdot x - N \cdot p = N \cdot (x-p)$
- Sum of distances:



$$d(x) := \sum_{i=1}^k N_i \cdot (x - p)$$

Quadric Error - Homogeneous Coordinates

- Suppose in coordinates we have

- a query point (x, y, z)
- a normal (a, b, c)
- an offset $d := -(x, y, z) \cdot (a, b, c)$

$$Q = \begin{bmatrix} a^2 & ab & ac & ad \\ ab & b^2 & bc & bd \\ ac & bc & c^2 & cd \\ ad & bd & cd & d^2 \end{bmatrix}$$

- Then in homogeneous coordinates, let

- $u := (x, y, z, 1)$
- $v := (a, b, c, d)$

- *Signed distance to plane is then just $u \cdot v = ax + by + cz + d$*

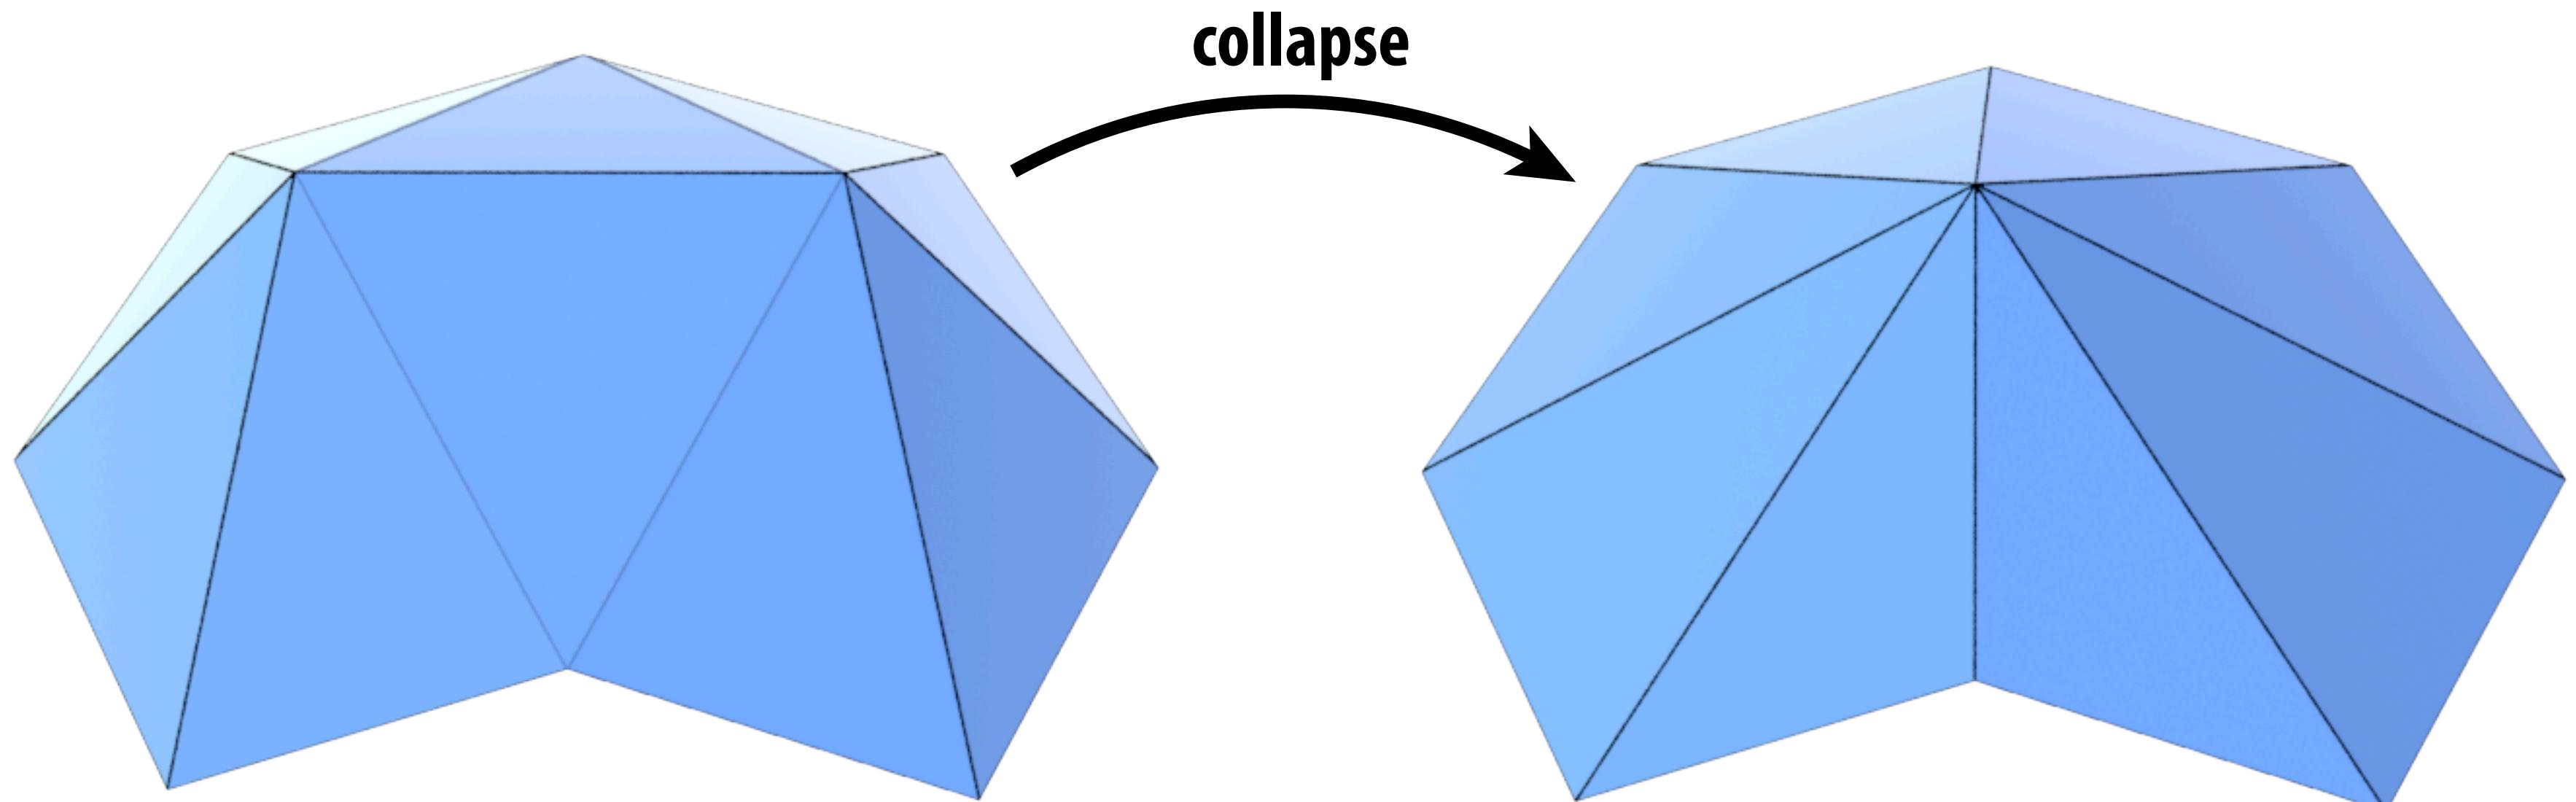
- *Squared distance is $(u^T v)^2 = u^T (v v^T) u =: u^T Q u$*

- Key idea: *matrix Q encodes distance to plane*

- Q is symmetric, contains 10 unique coefficients (small storage)

Quadric Error of Edge Collapse

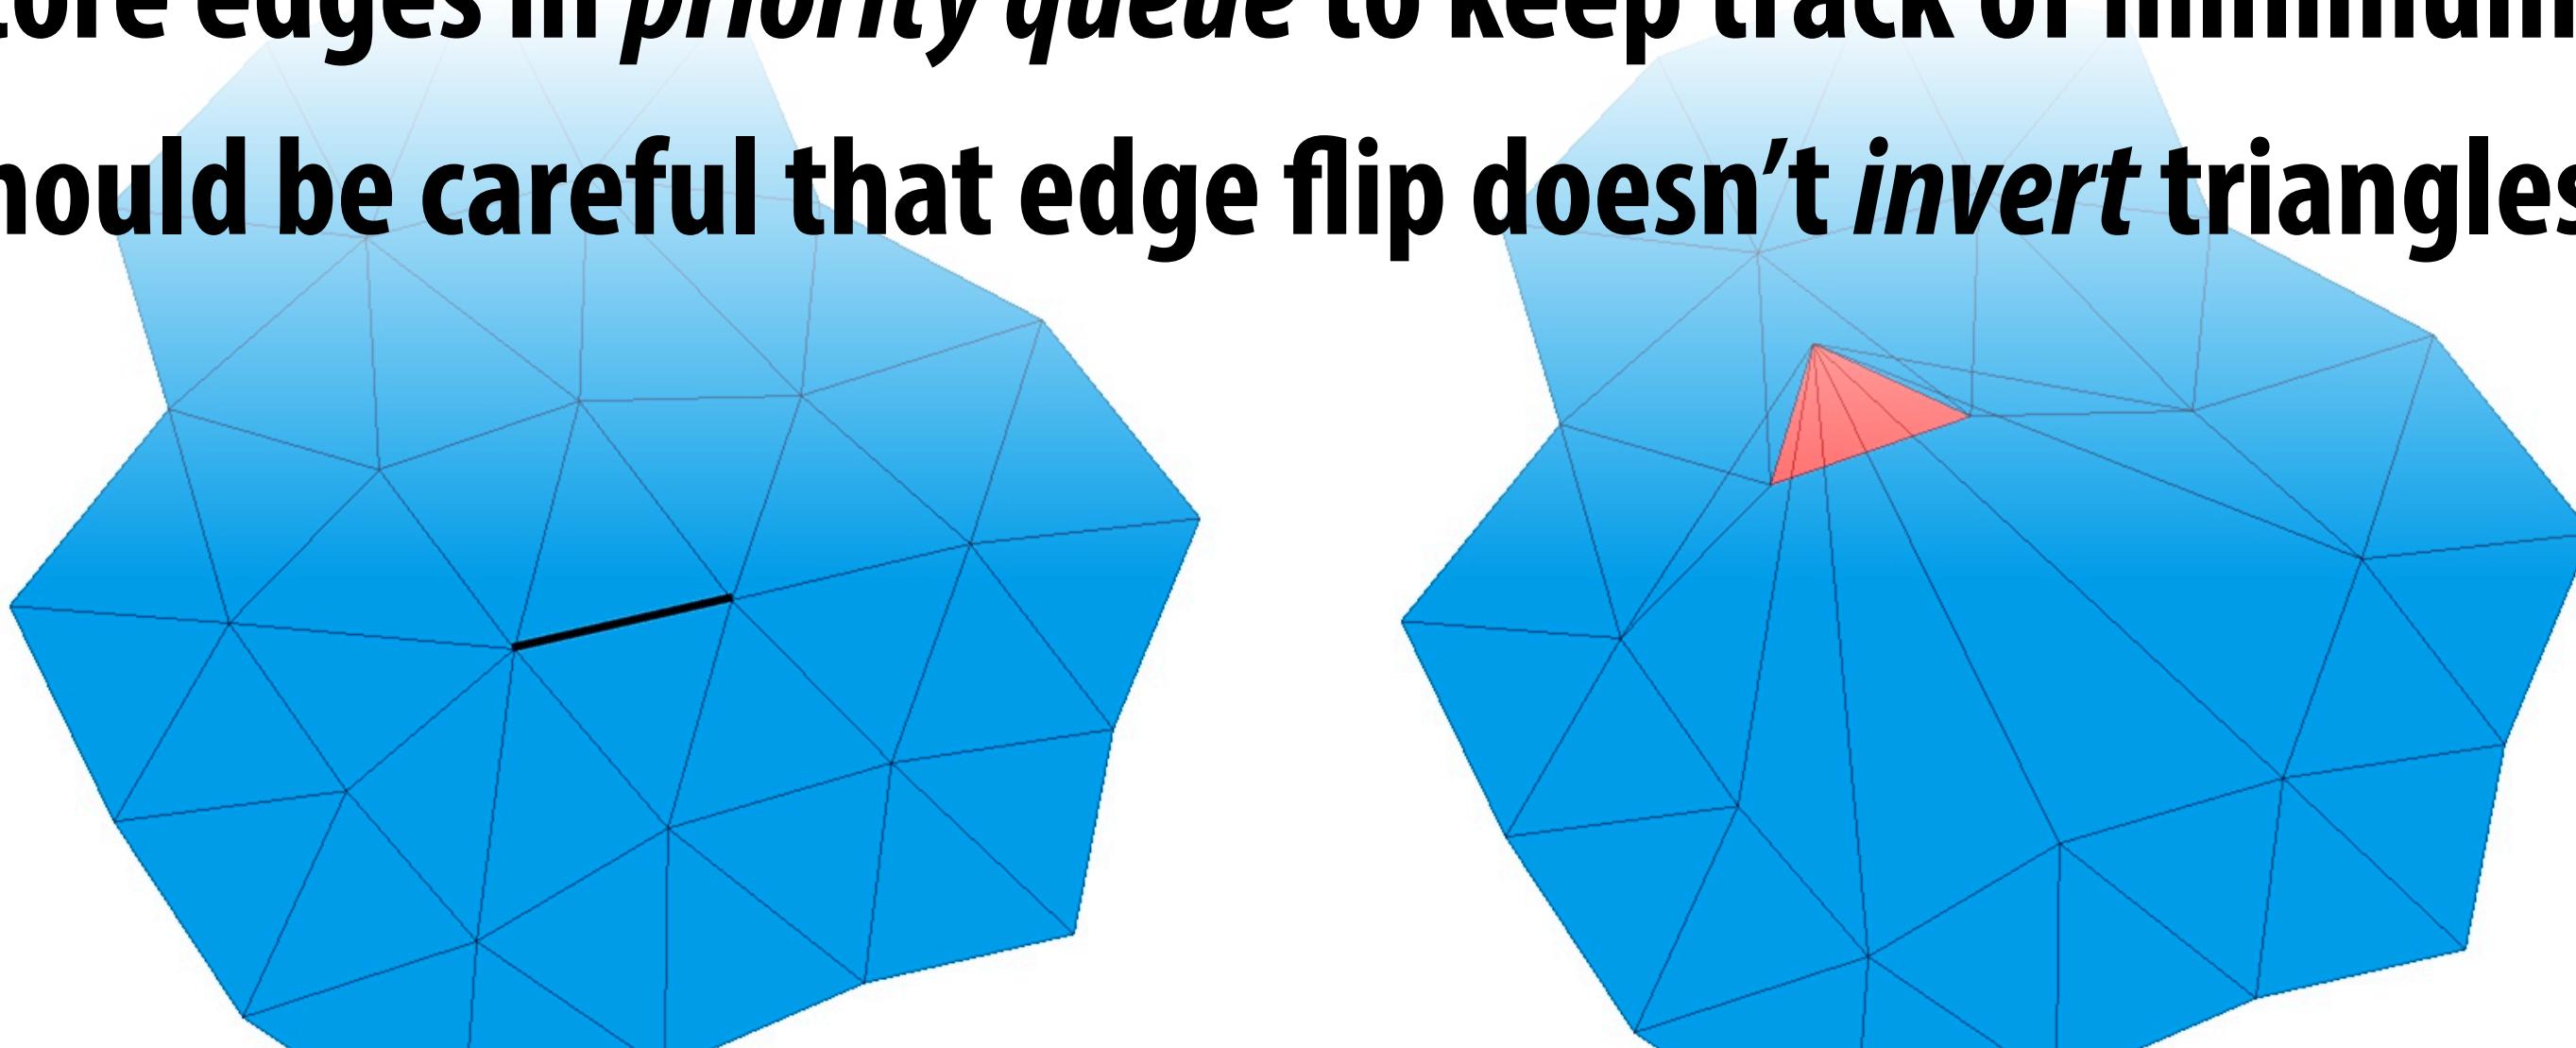
- How much does it cost to collapse an edge?
- Idea: compute edge midpoint, measure quadric error



- Better idea: use point that minimizes quadric error as new point!
- (More details in assignment; see also Garland & Heckbert 1997.)

Quadric Error Simplification

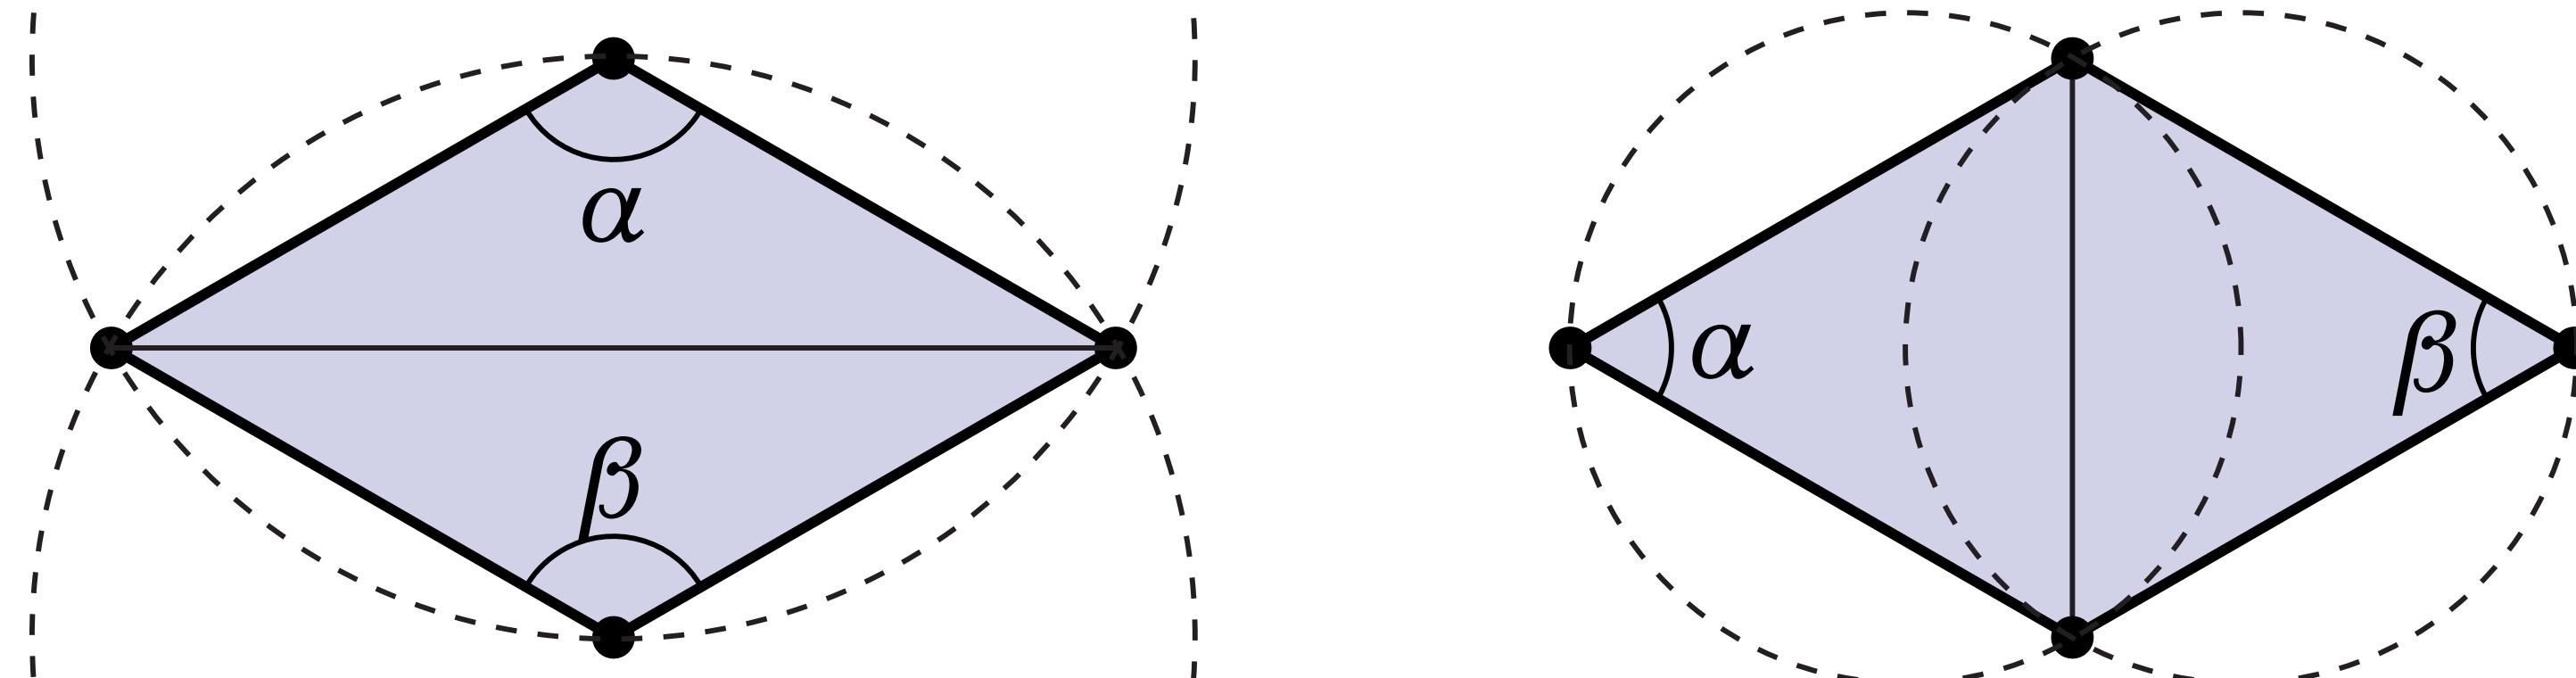
- Compute Q for each triangle
- Set Q at each vertex to sum of Q s from incident triangles
- Until we reach target # of triangles:
 - collapse edge (i,j) with smallest cost to get new vertex k
 - add Q_i and Q_j to get new quadric Q_k
 - update cost of any edge touching new vertex k
- Store edges in *priority queue* to keep track of minimum cost
- Should be careful that edge flip doesn't *invert* triangles:



What if we're happy with the *number* of triangles, but want to improve *quality*?

How do we make a mesh “more Delaunay”?

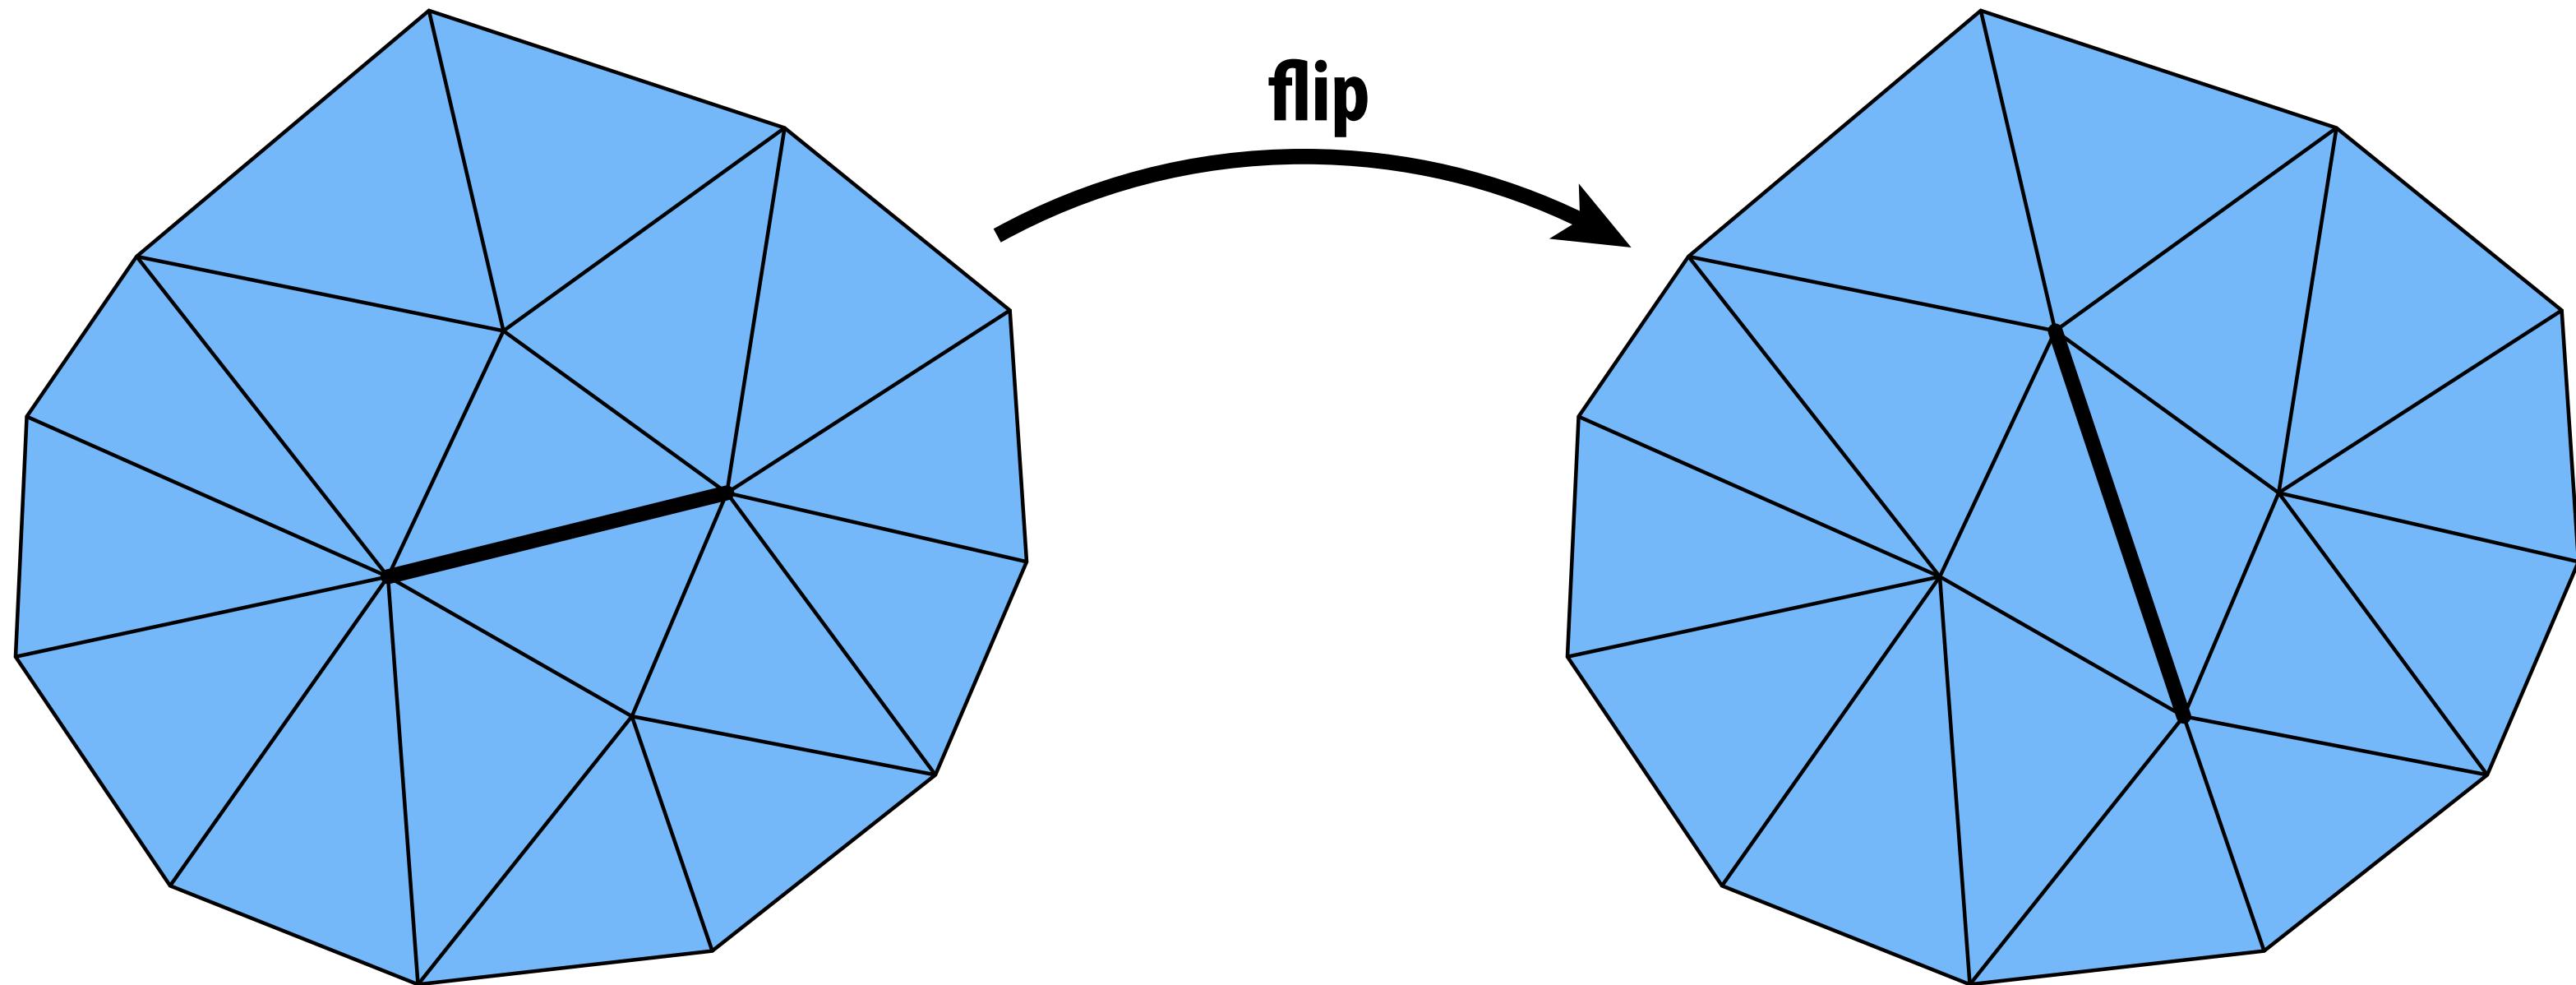
- Already have a good tool: edge flips!
- If $\alpha + \beta > \pi$, flip it!



- FACT: in 2D, flipping edges eventually yields Delaunay mesh
- Theory: worst case $O(n^2)$; no longer true for surfaces in 3D.
- Practice: simple, effective way to improve mesh quality

Alternatively: how do we improve degree?

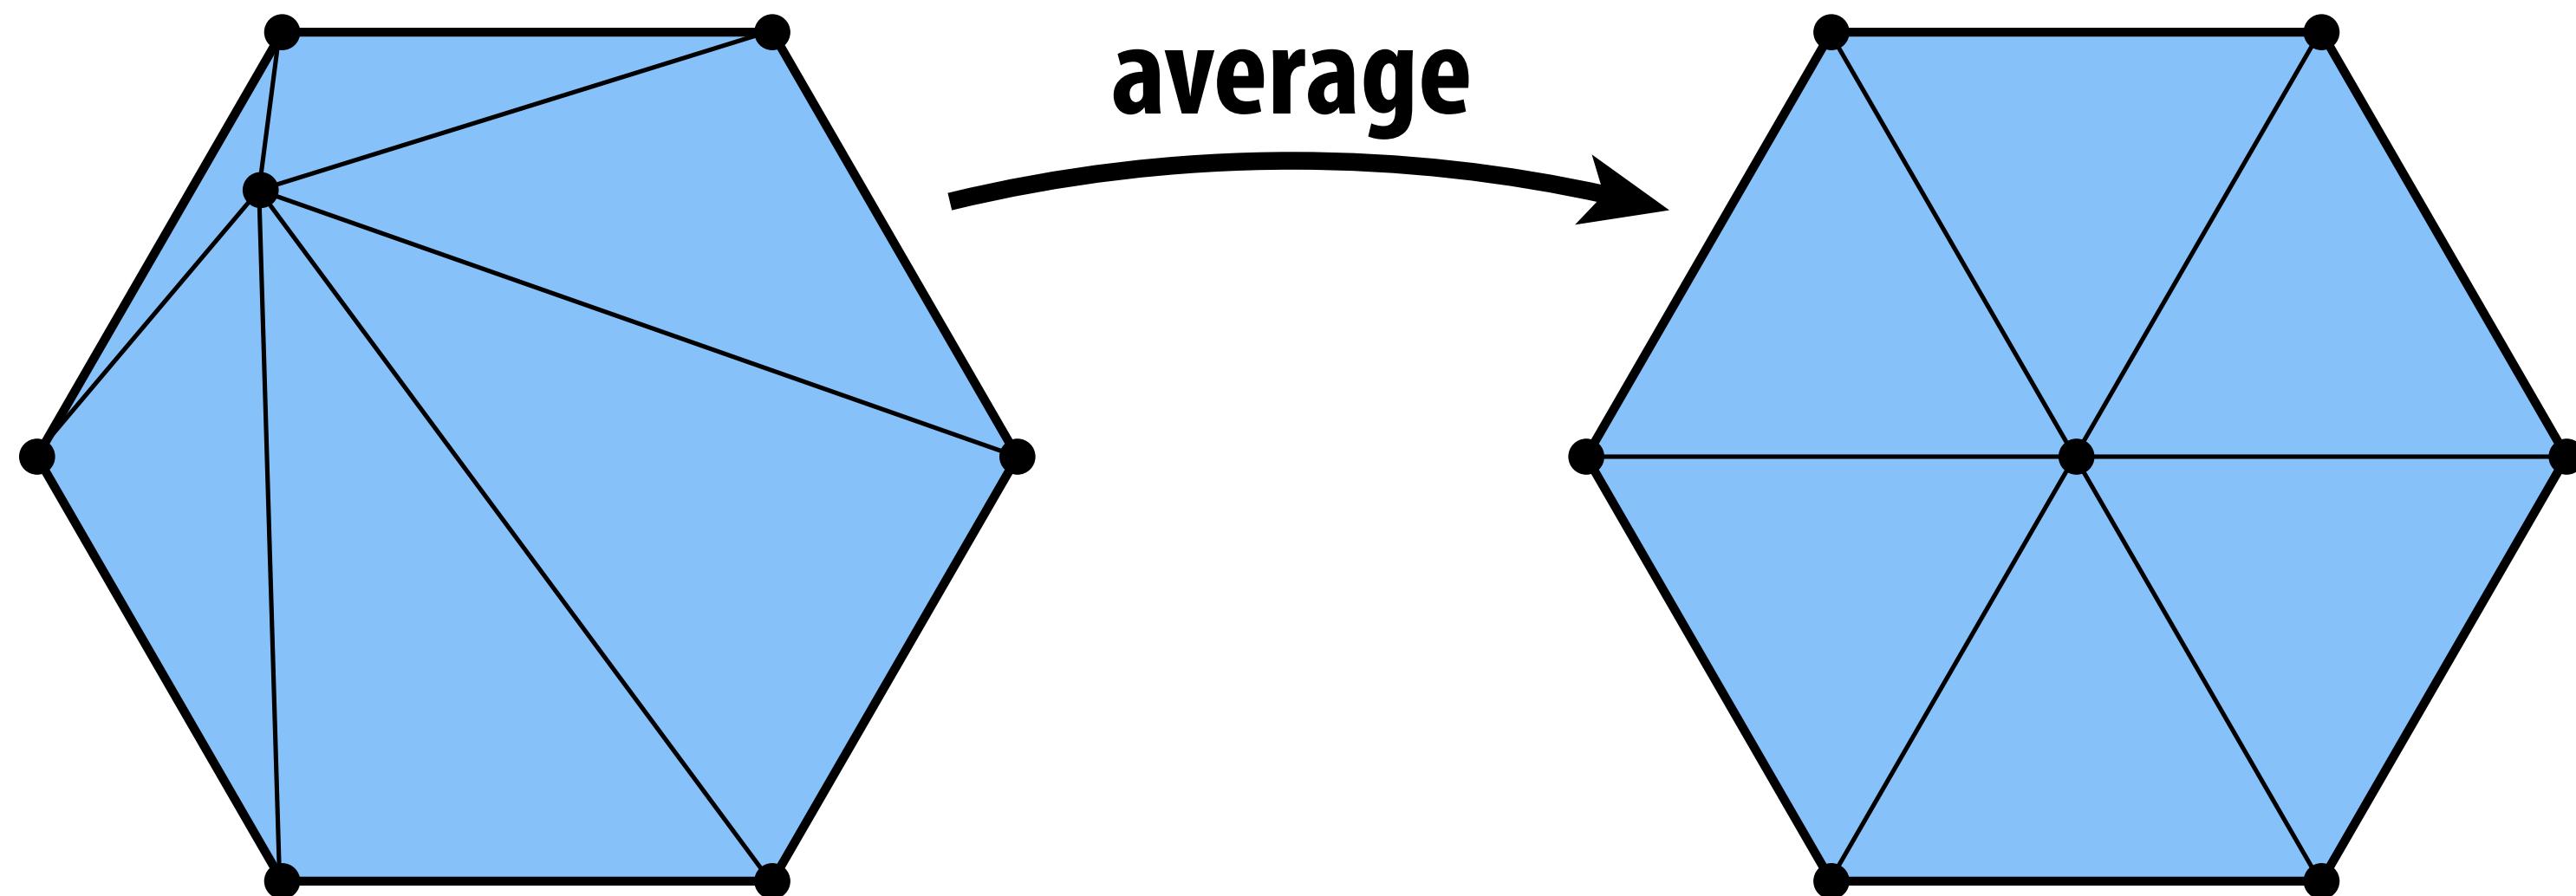
- Same tool: edge flips!
- If total deviation from degree-6 gets smaller, flip it!



- FACT: average valence of any triangle mesh is 6
- Iterative edge flipping acts like “discrete diffusion” of degree
- Again, no (known) guarantees; works well in practice

How do we make a triangles “more round”?

- Delaunay doesn’t mean triangles are “round” (angles near 60°)
- Can often improve shape by centering vertices:

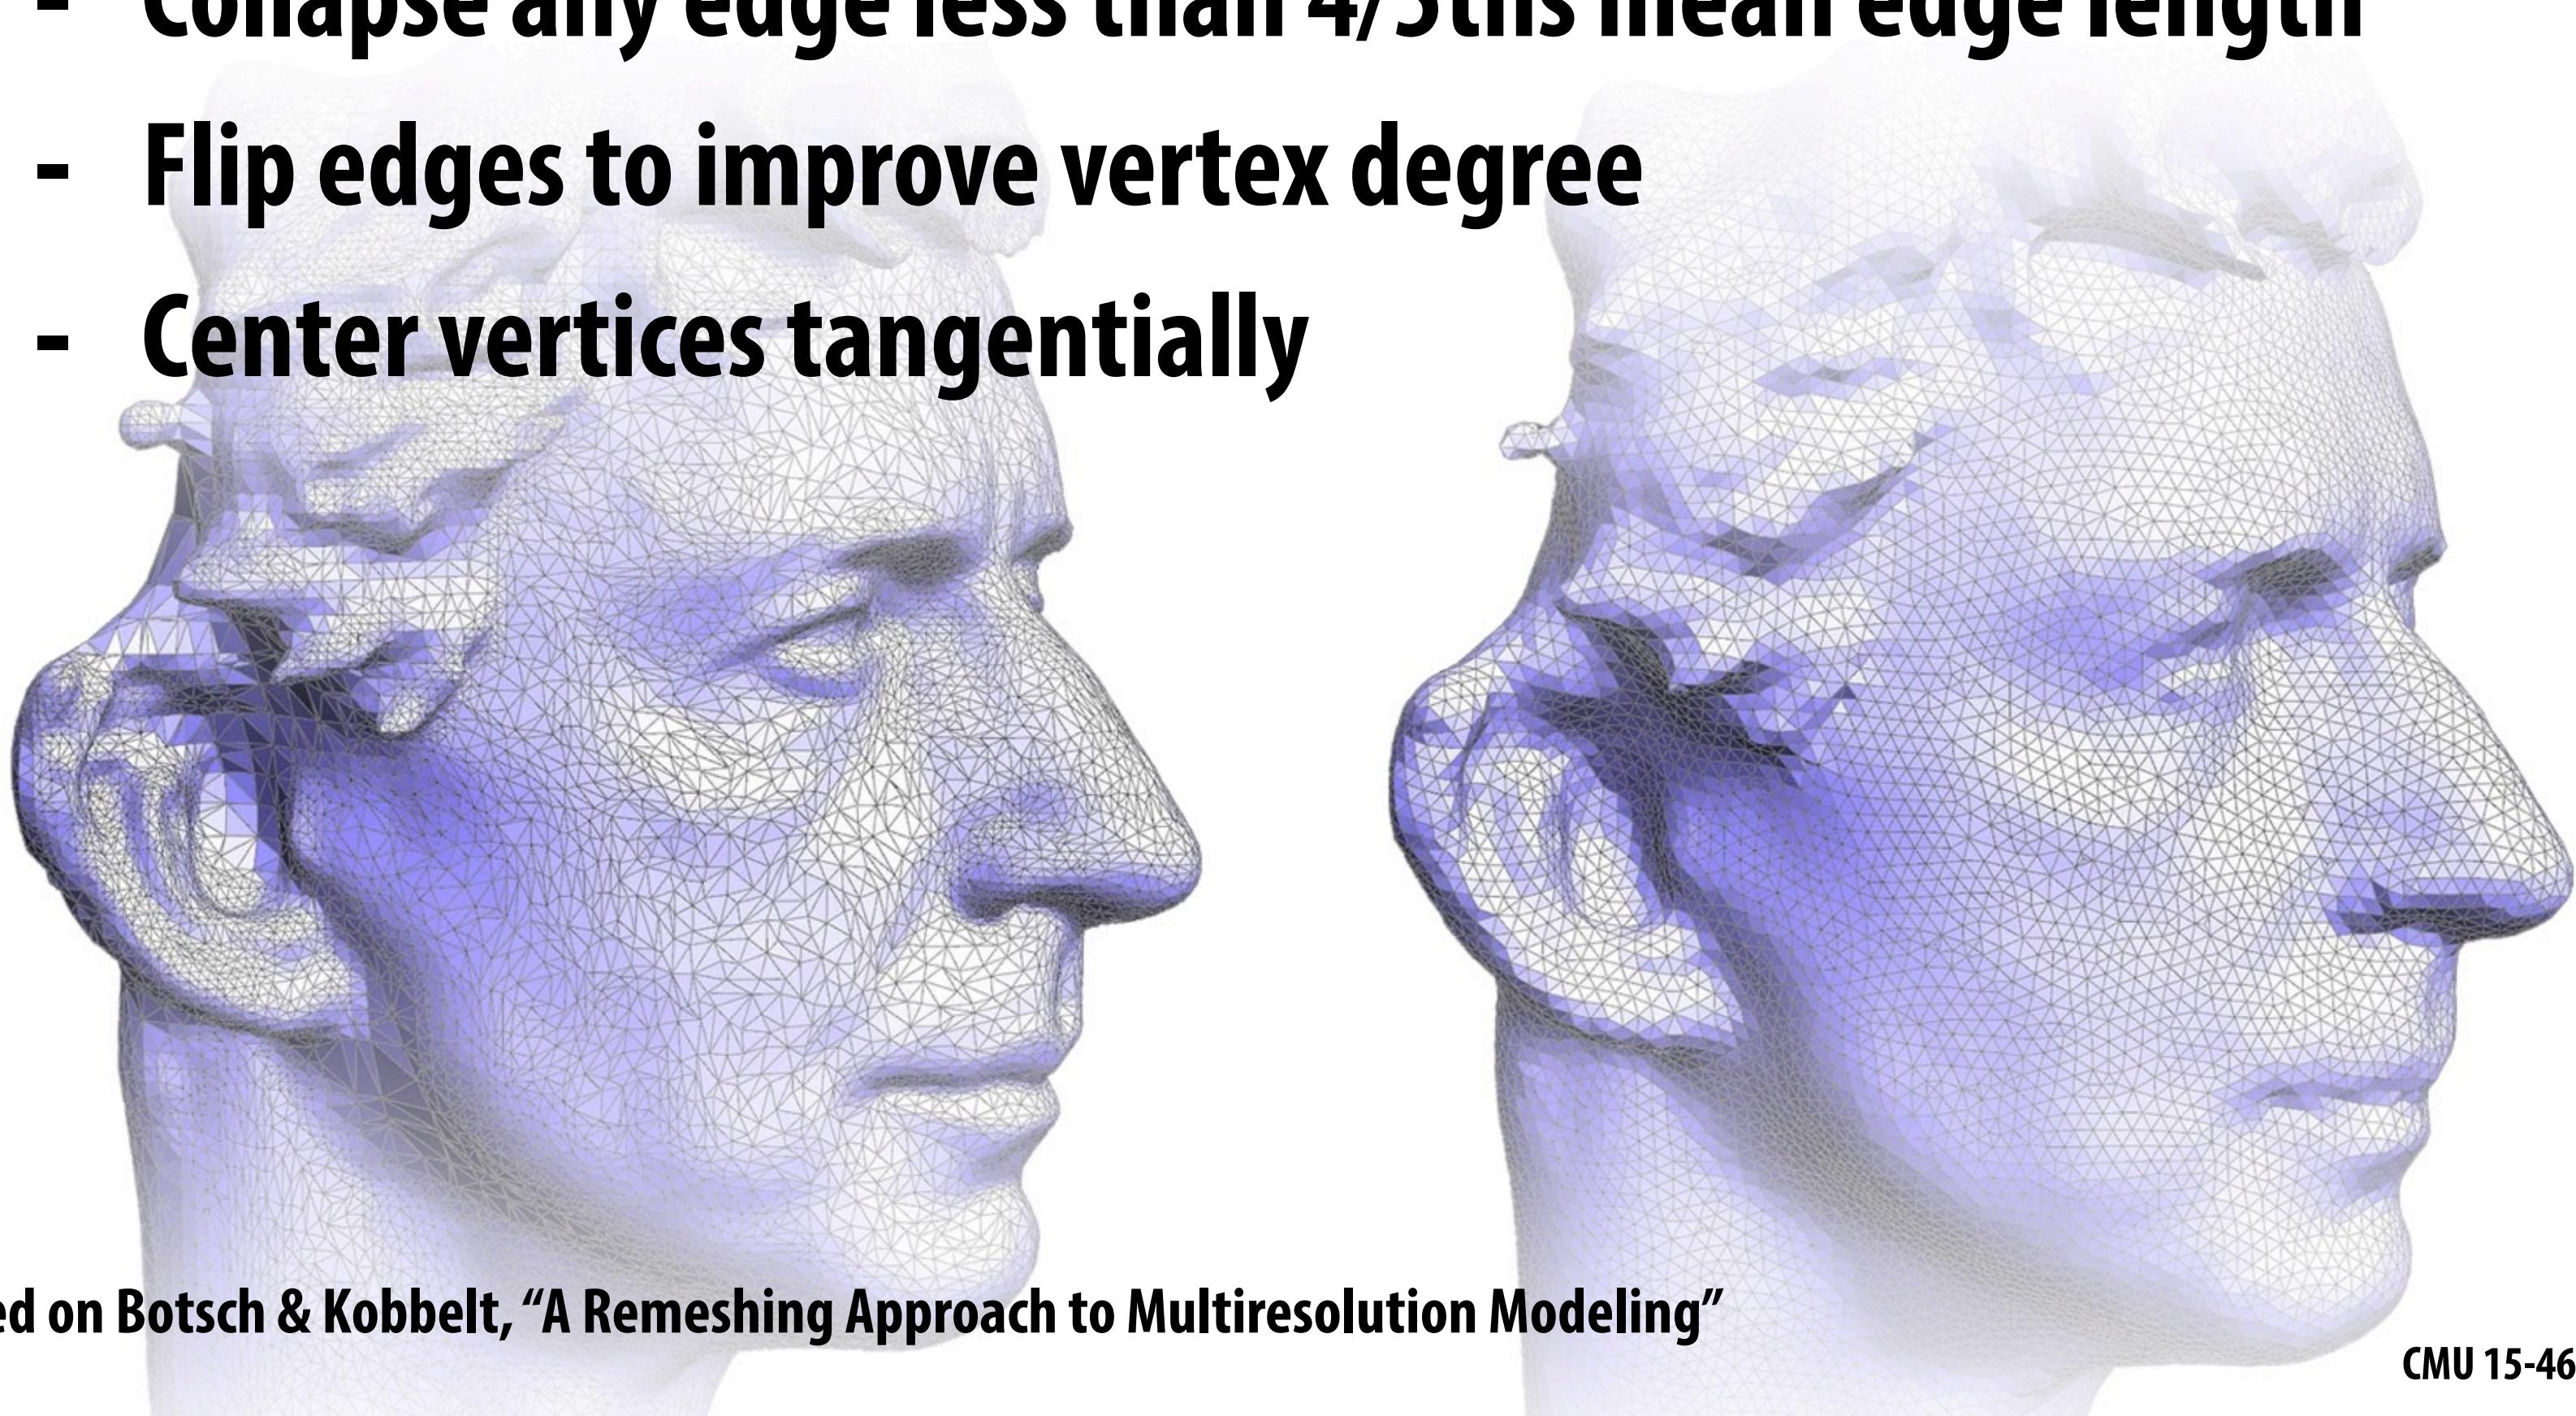


- Simple version of technique called “Laplacian smoothing”.*
- On surface: move only in *tangent* direction
- How? Remove normal component from update vector.

*See Crane, “Digital Geometry Processing with Discrete Exterior Calculus” <http://keenan.is/dgpdec>

Isotropic Remeshing Algorithm*

- Try to make triangles uniform shape & size
- Repeat four steps:
 - Split any edge over 4/3rds mean edge length
 - Collapse any edge less than 4/5ths mean edge length
 - Flip edges to improve vertex degree
 - Center vertices tangentially



*Based on Botsch & Kobbelt, "A Remeshing Approach to Multiresolution Modeling"

Next time: Geometric Queries

- Not all queries can be answered “on the surface”
- E.g.,
 - distance from point to surface?
 - intersection between two surfaces?
 - where does a ray pierce the surface?
- Need totally different data structures/algorithms

